



Lecture 14: Scikit Learn Pipelines

Recap

- Kmeans algorithm
- Expectation Maximization
- Centroid initialization with kmeans++
- Kmeans decision boundary
- Kmeans limitations with oblong data clusters
- Kmeans limitations with outliers
- Using Kmeans and silhouette analysis for outlier detection

**Data
Preprocessing
(duplicate
removal,
imputation,
encoding)**

2) Data understanding

Goal: Understand the data



3) Data preparation

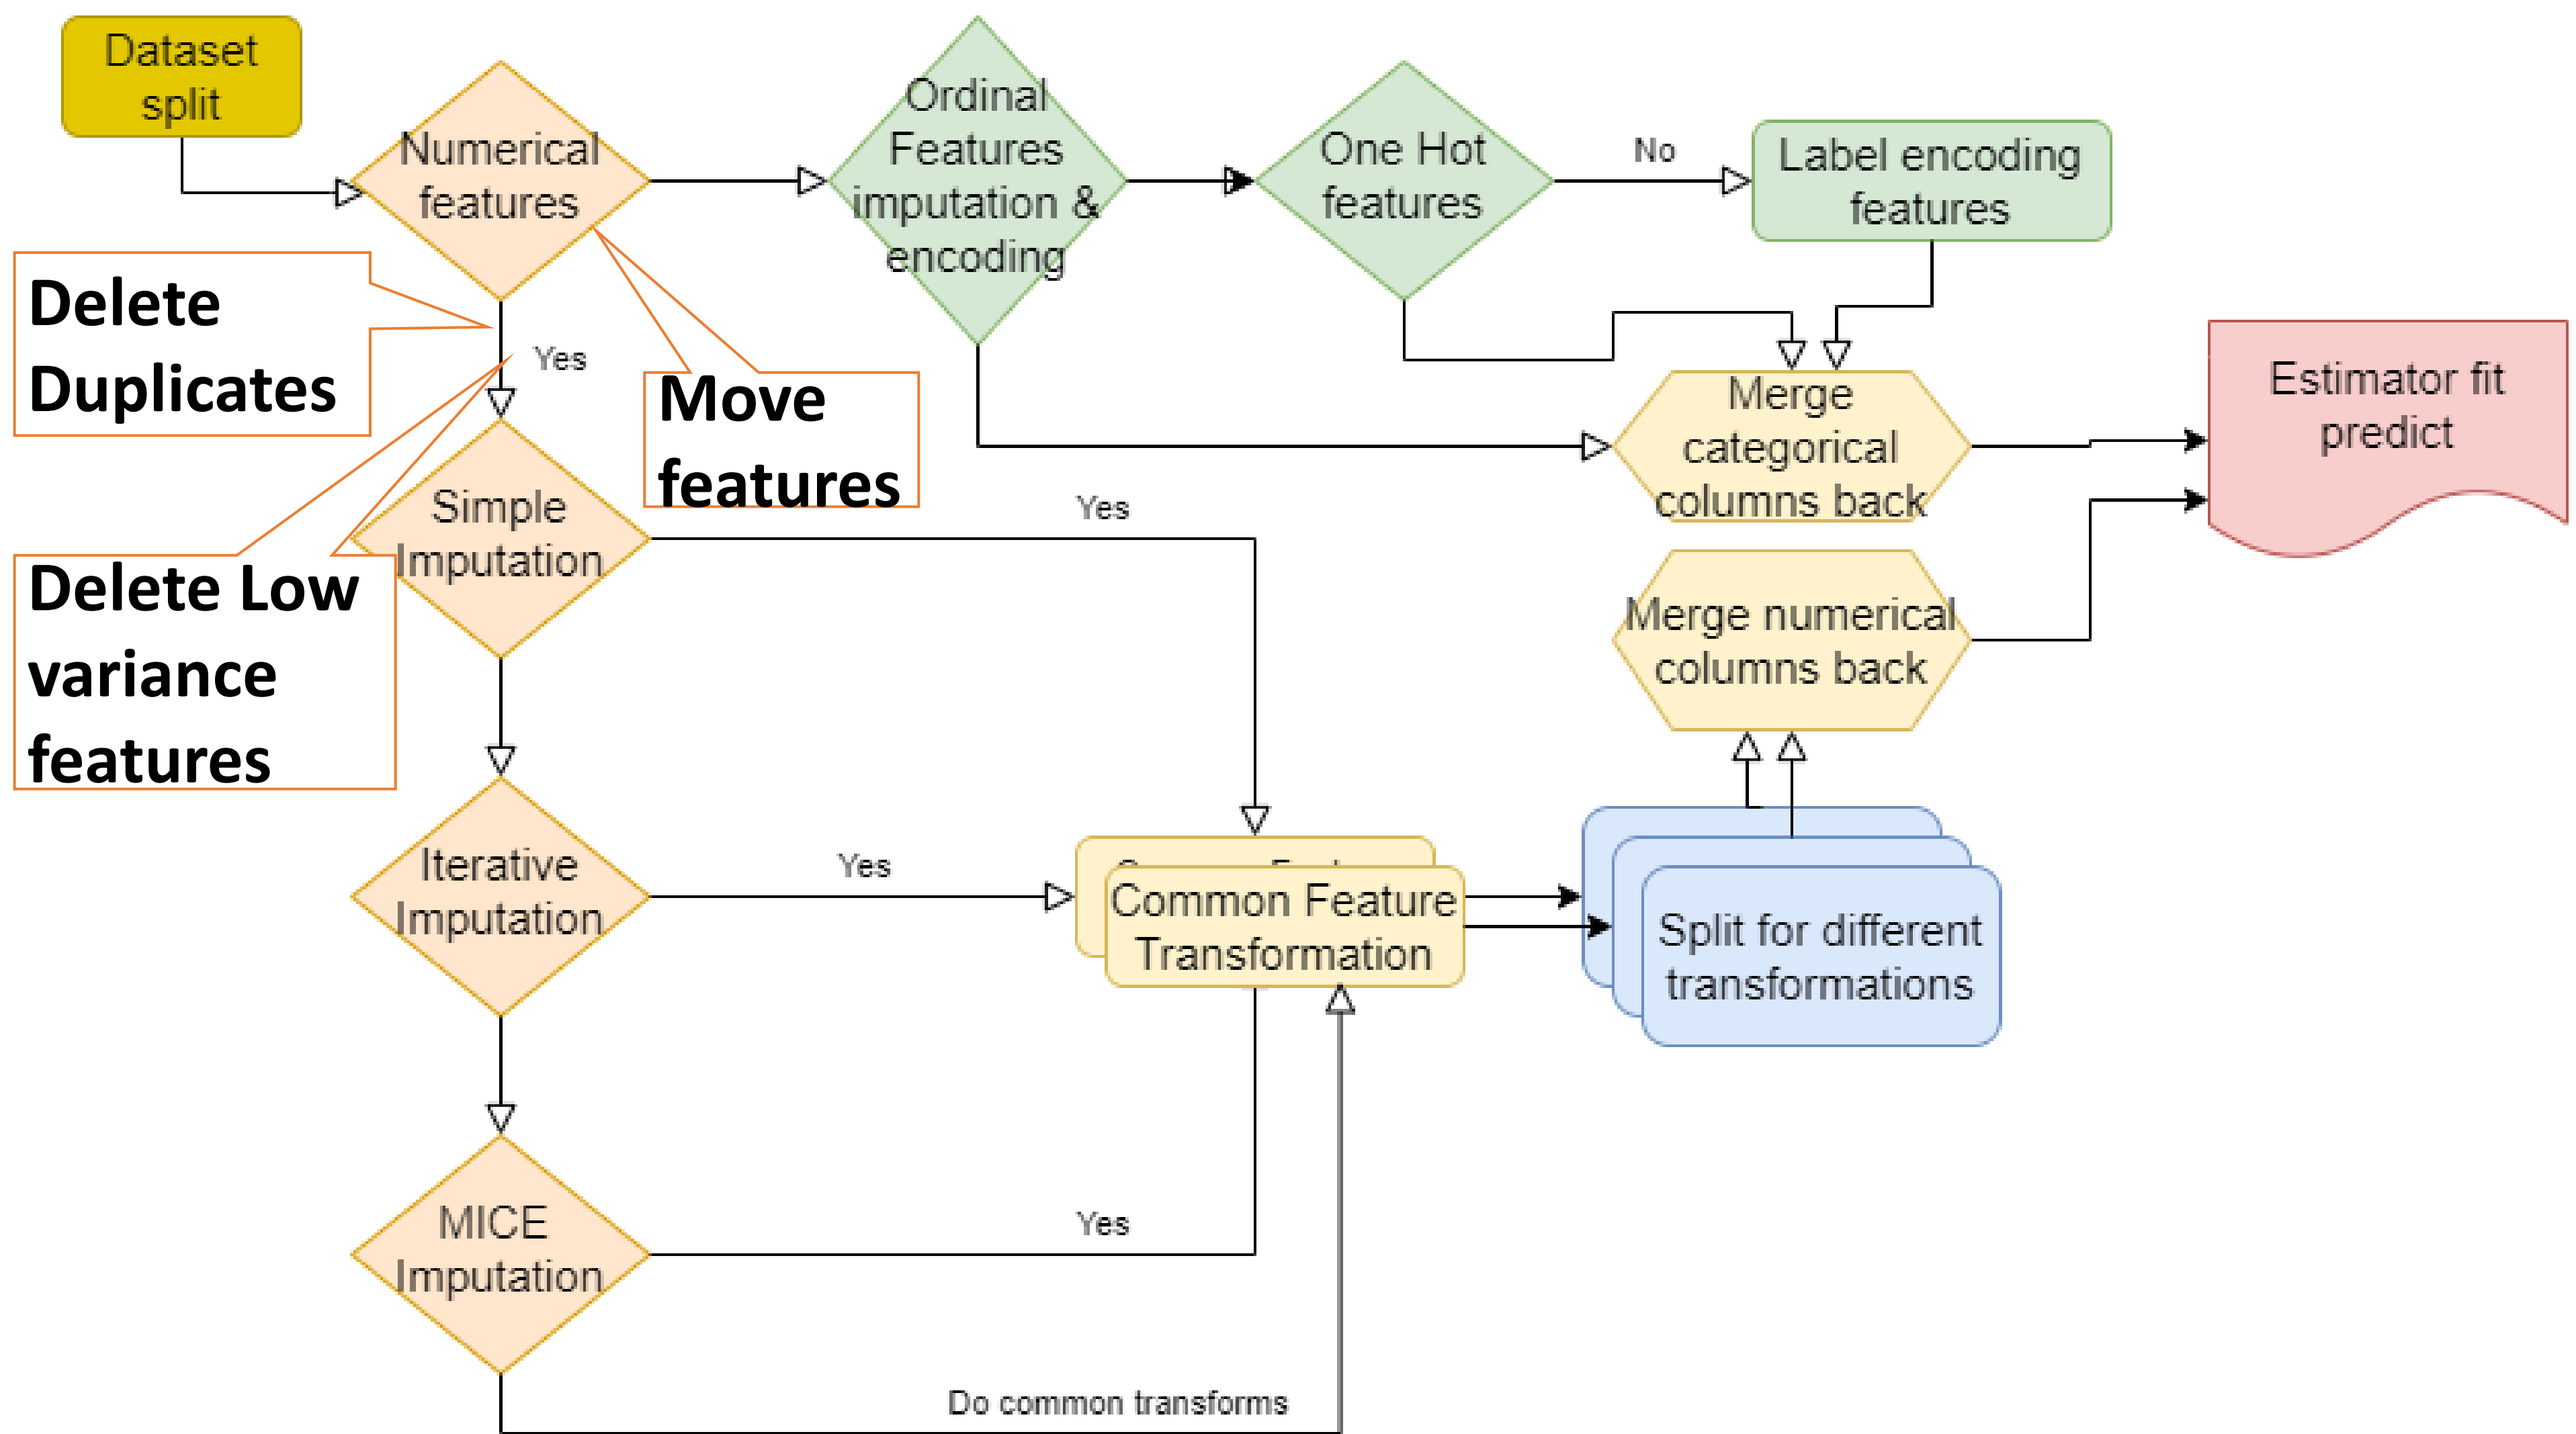
Goal: Prepare data for model



4) Modelling

Goal: Build a model

- 1. Feature Selection**
- 2. Feature Transformation/Engineering**
- 3. Feature Extraction**





Without pipeline


```
#imputer object for missing value
imputer = SimpleImputer(strategy="median")

# scales object the numerical feature
scaler = StandardScaler()

# one-hot object for the categorical features
one_hot=OneHotEncoder(handle_unknown='ignore',sparse=False)

# Applying the fit and transform methods to the Training data
imputer.fit(X_train[numeric_features])
imputed=imputer.transform(X_train[numeric_features])

scaler.fit(imputed)
scaled=scaler.transform( imputed)

one_hot.fit(X_train[categorical_features])
one_hotted_cat=one_hot.transform(X_train[categorical_features])

# Concatenating the scaled and one hot matrixes
Final_train=pd.DataFrame(np.concatenate((scaled,one_hotted_cat), axis=1))
lr.fit(Final_train, y_train)
```

- Preprocessing generally 500-1000 features
 - To be grouped in about 25-50 features
 - E.g. Diff numeric fields need diff imputation
 - Diff groups need diff encoding
 - Encoding as embedding
 - Feature Selection, Elimination, Importance, extraction etc.
 - Distribution transformation
 - Cross validation
- Output of one should be fed into next
- Dynamic code, rearranged often
- Error prone by design

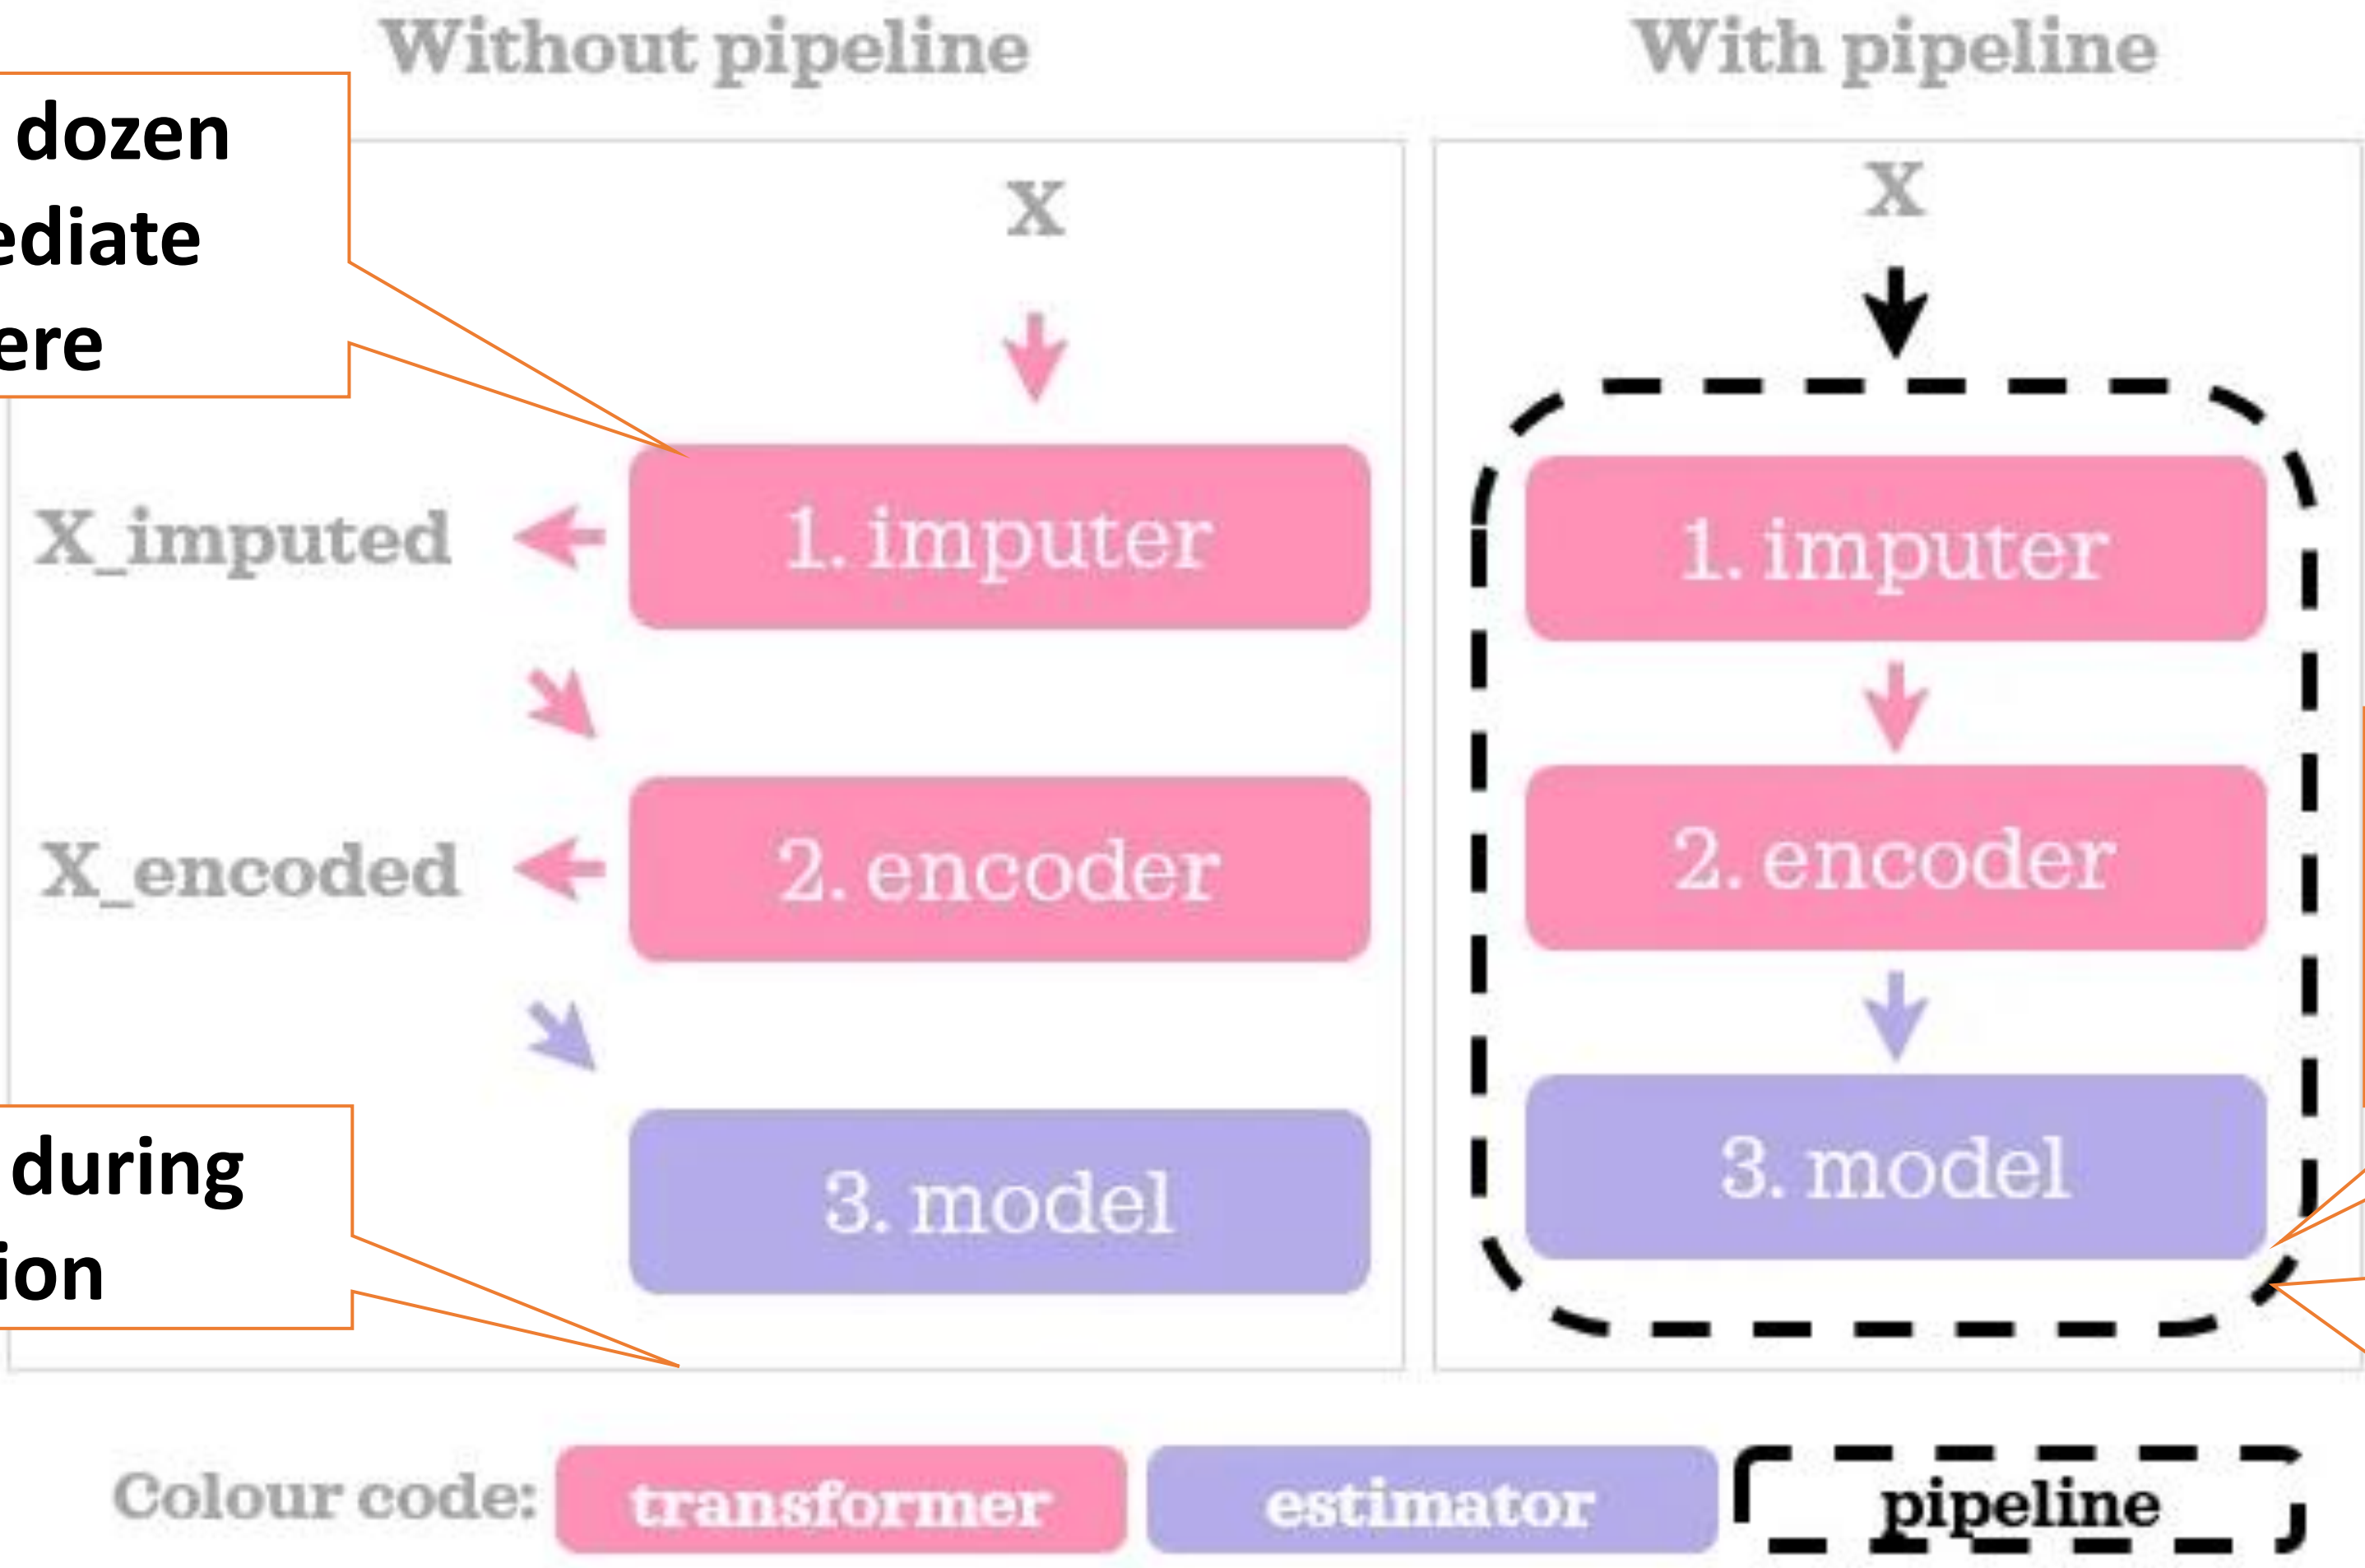

```
# Predict on the test set-using the trained classifier-still need to do the transformations
X_test_filled = imputer.transform(X_test[numeric_features])
X_test_scaled = scaler.transform(X_test_filled)
X_test_one_hot = one_hot.transform(X_test[categorical_features])
X_test=pd.DataFrame(np.concatenate((X_test_scaled, X_test_one_hot), axis=1))
lr.score(X_test,y_test)
```

- Same order in predict
- Predict done by some other team
 - Two dozen models
 - Provide all configured imputers
 - High cost of ownership & maintenance
 - Brittle code - Touch and it breaks
 - Logistics nightmare
- Data leakage (What is it?)



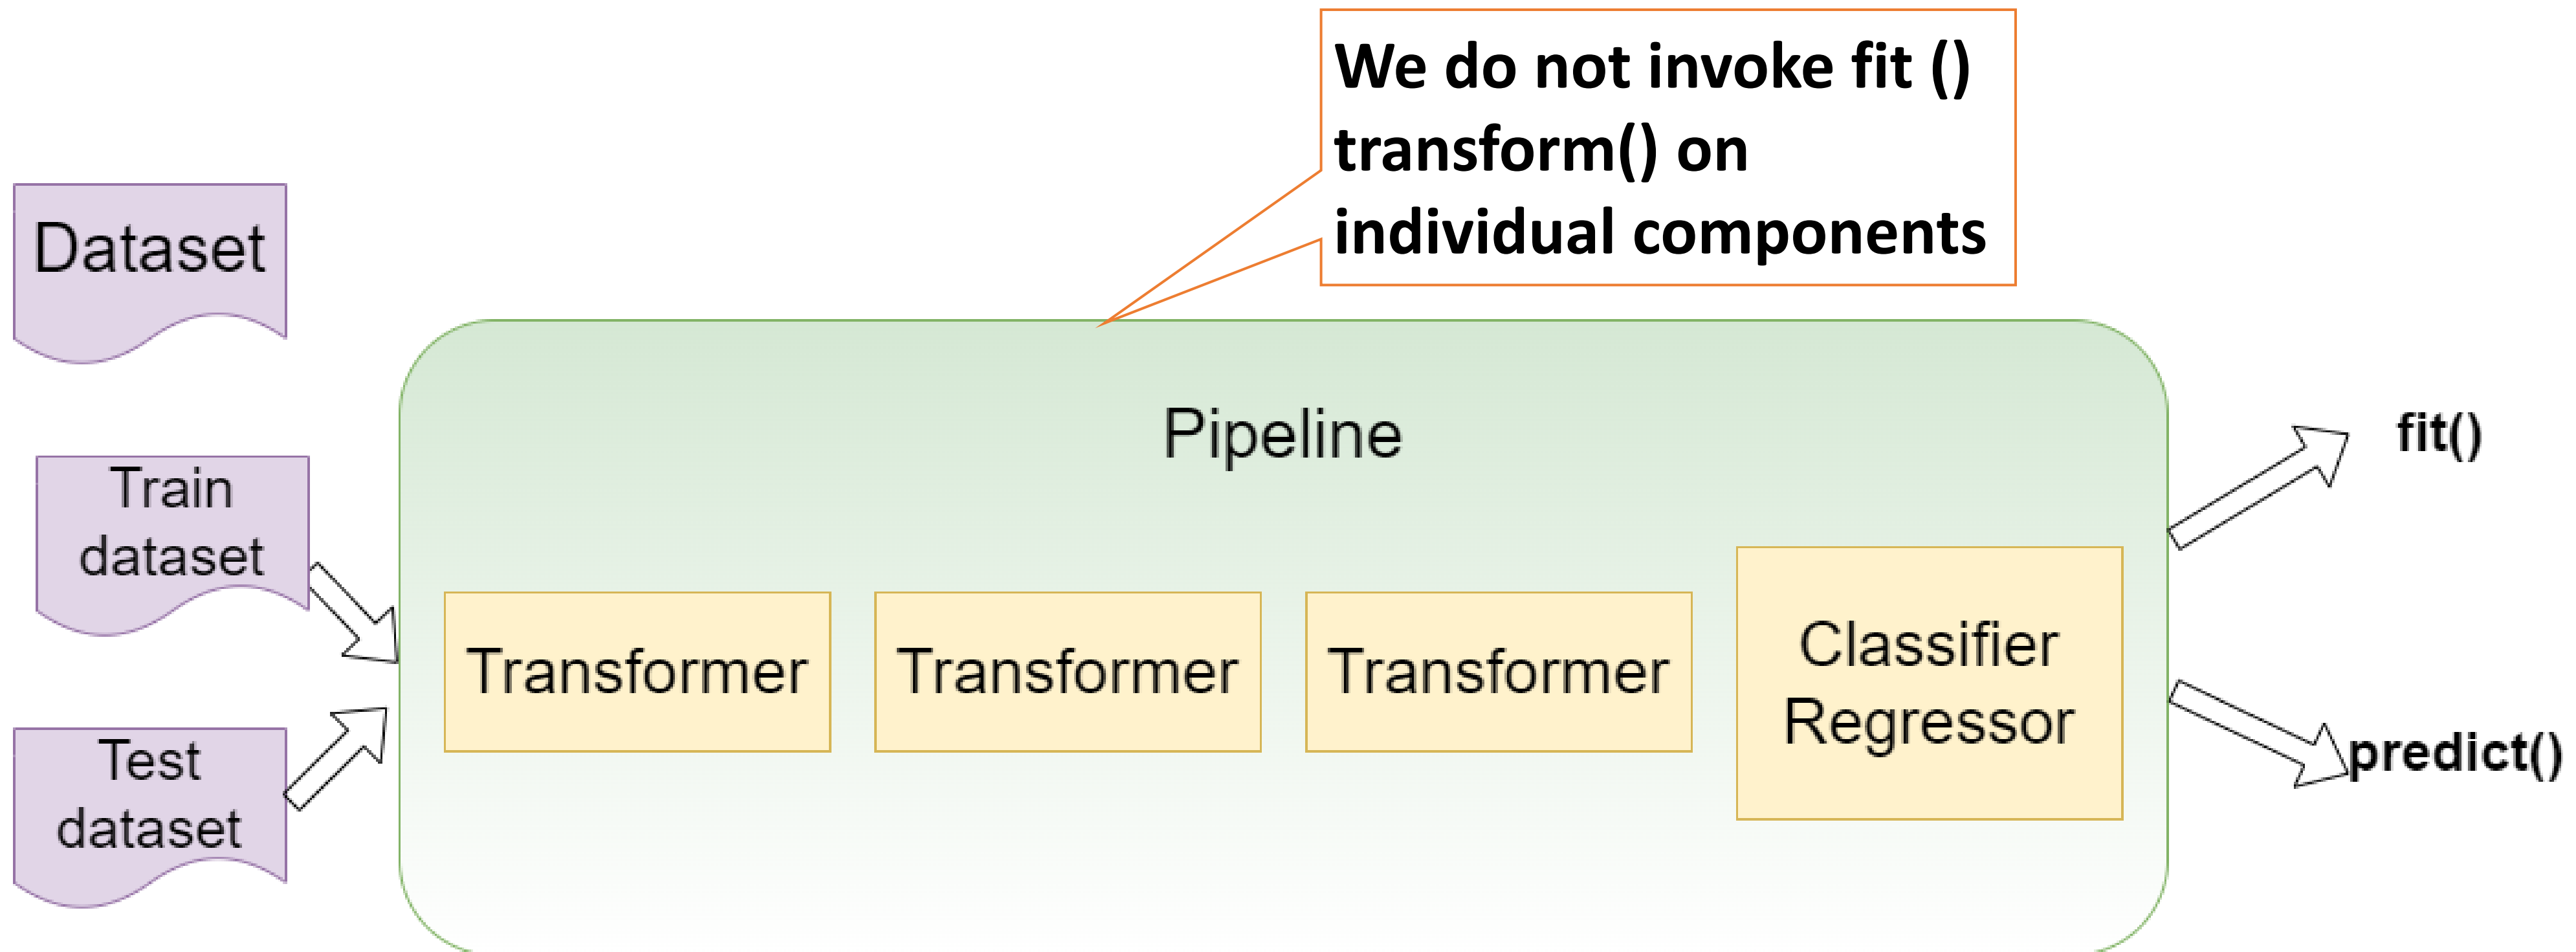
Several dozen intermediate steps here

Repeat during prediction



Last component in pipeline is ML Mixin

Pickle the pipeline & share



```
BaseEstimator:

get_params(self, deep=True)
set_params(self, **params)
```

```
TransformerMixin:

def fit_transform(self, X, y=None, **fit_params):
    if y is None:
        # fit method of arity 1 (unsupervised transformation)
        return self.fit(X, **fit_params).transform(X)
    else:
        # fit method of arity 2 (supervised transformation)
        return self.fit(X, y, **fit_params).transform(X)
```

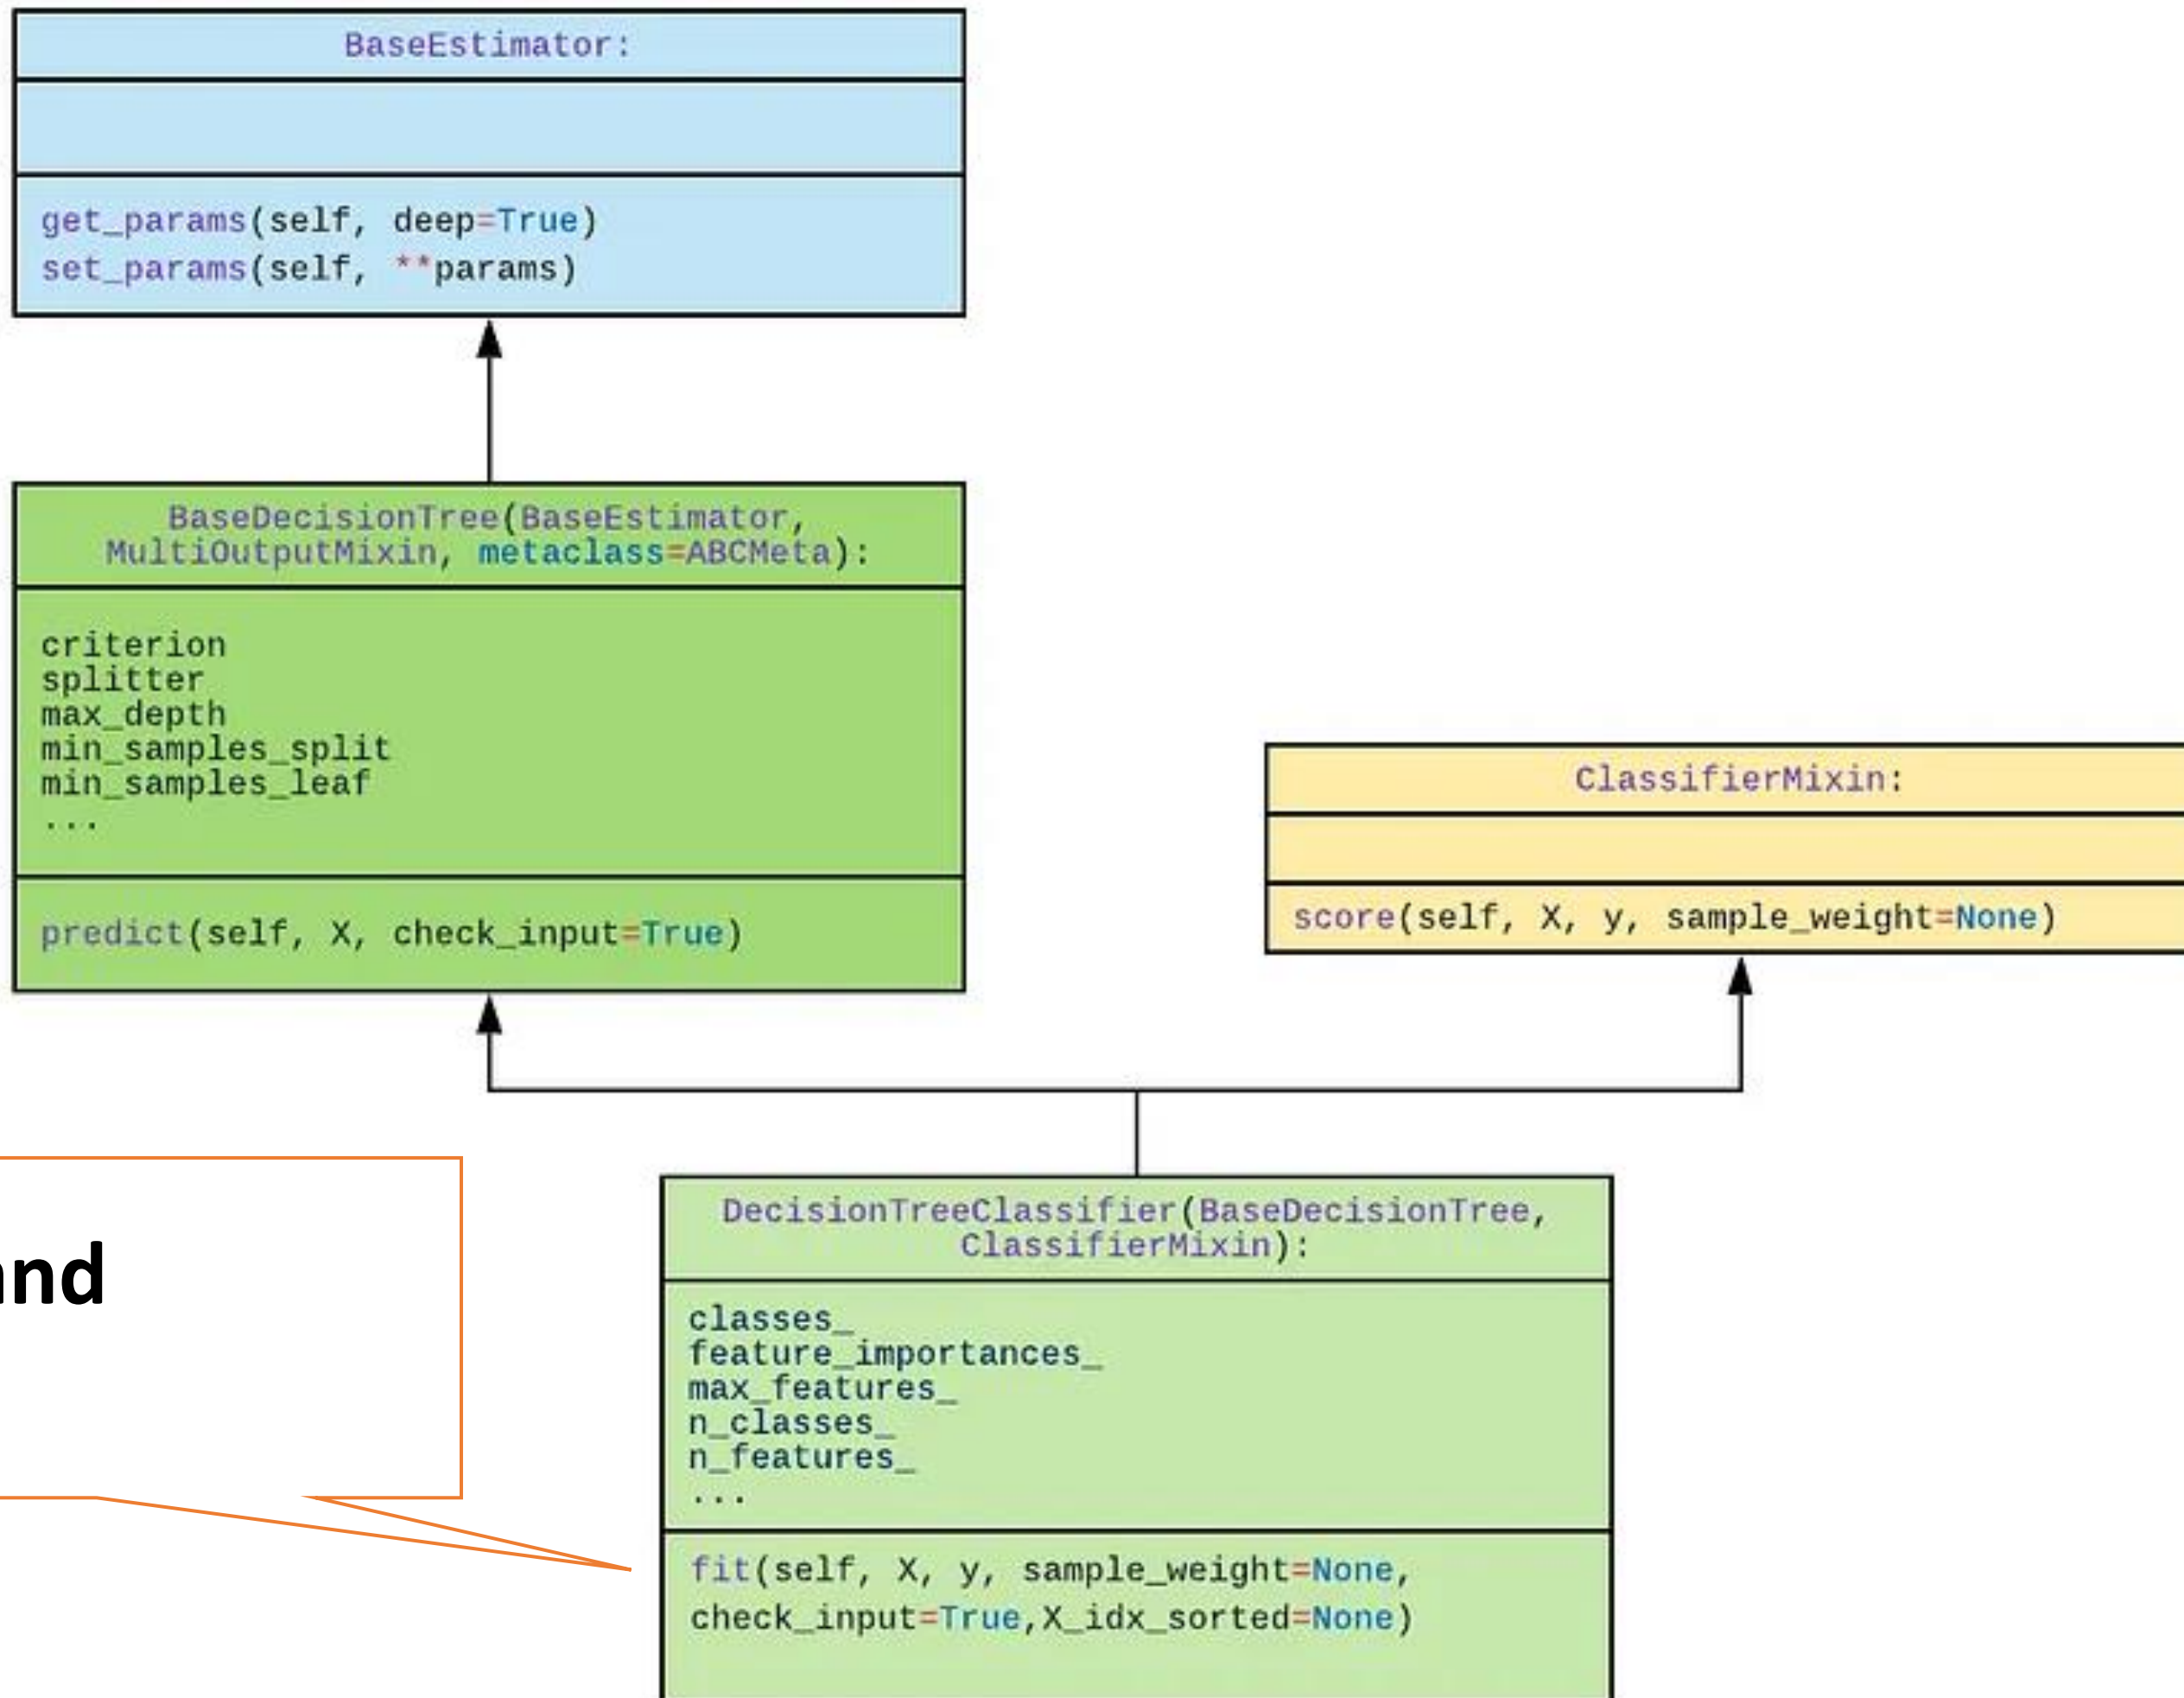
What do fit() and transform() do?

```
StandardScaler(BaseEstimator,
               TransformerMixin):

scale_
mean_
var_
n_samples_seen_

fit(self, X, y=None):
transform(self, X):
```

fit_transform() combines both fit() and transform()



Only fit() and
predict()

fit(), fit_transform() and predict()

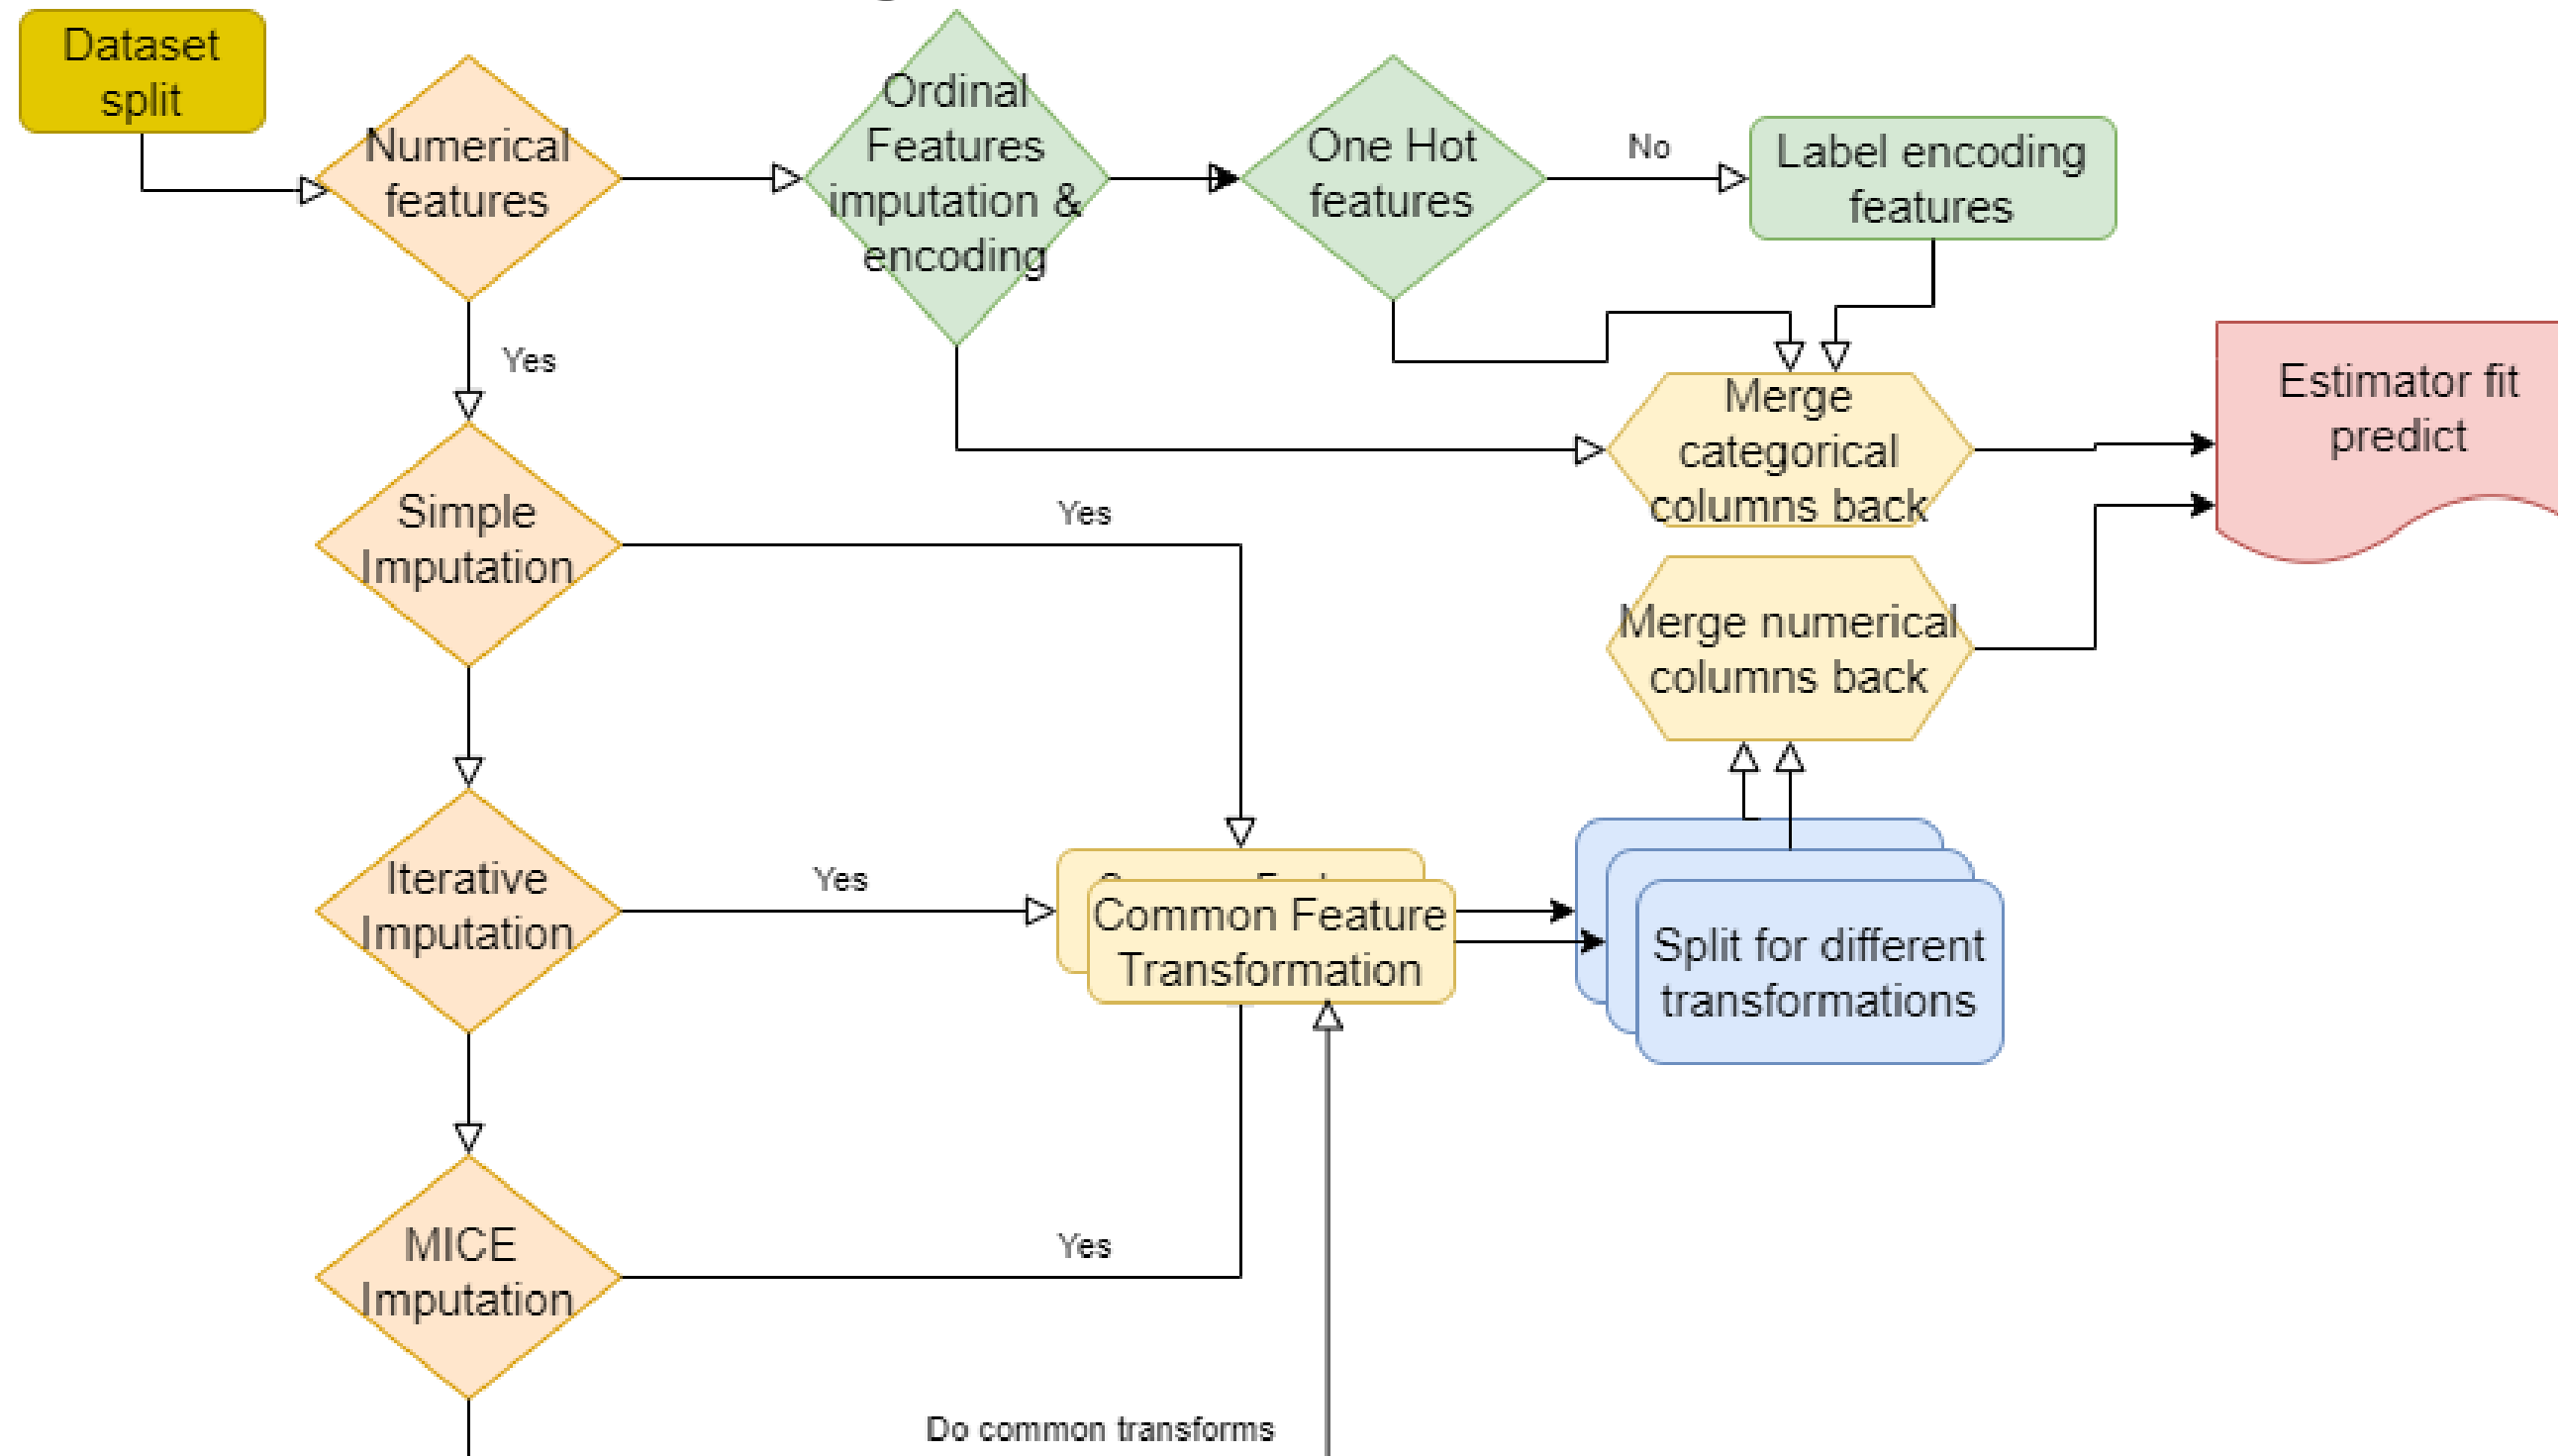
- Always configure pipeline with
 - Last component as Classifier/Regressor/Cluster Mixin
 - Other components are Transformers
- Invoke fit() on pipeline during training
 - Sklearn invokes fit_transform() on all but last component
 - Sklearn invokes fit() on last component
- Invoke predict() on pipeline during test
 - Sklearn invokes transform() on all but last component
 - Sklearn invokes predict on last component
- Sklearn relies on python duck typing than class hierarchy

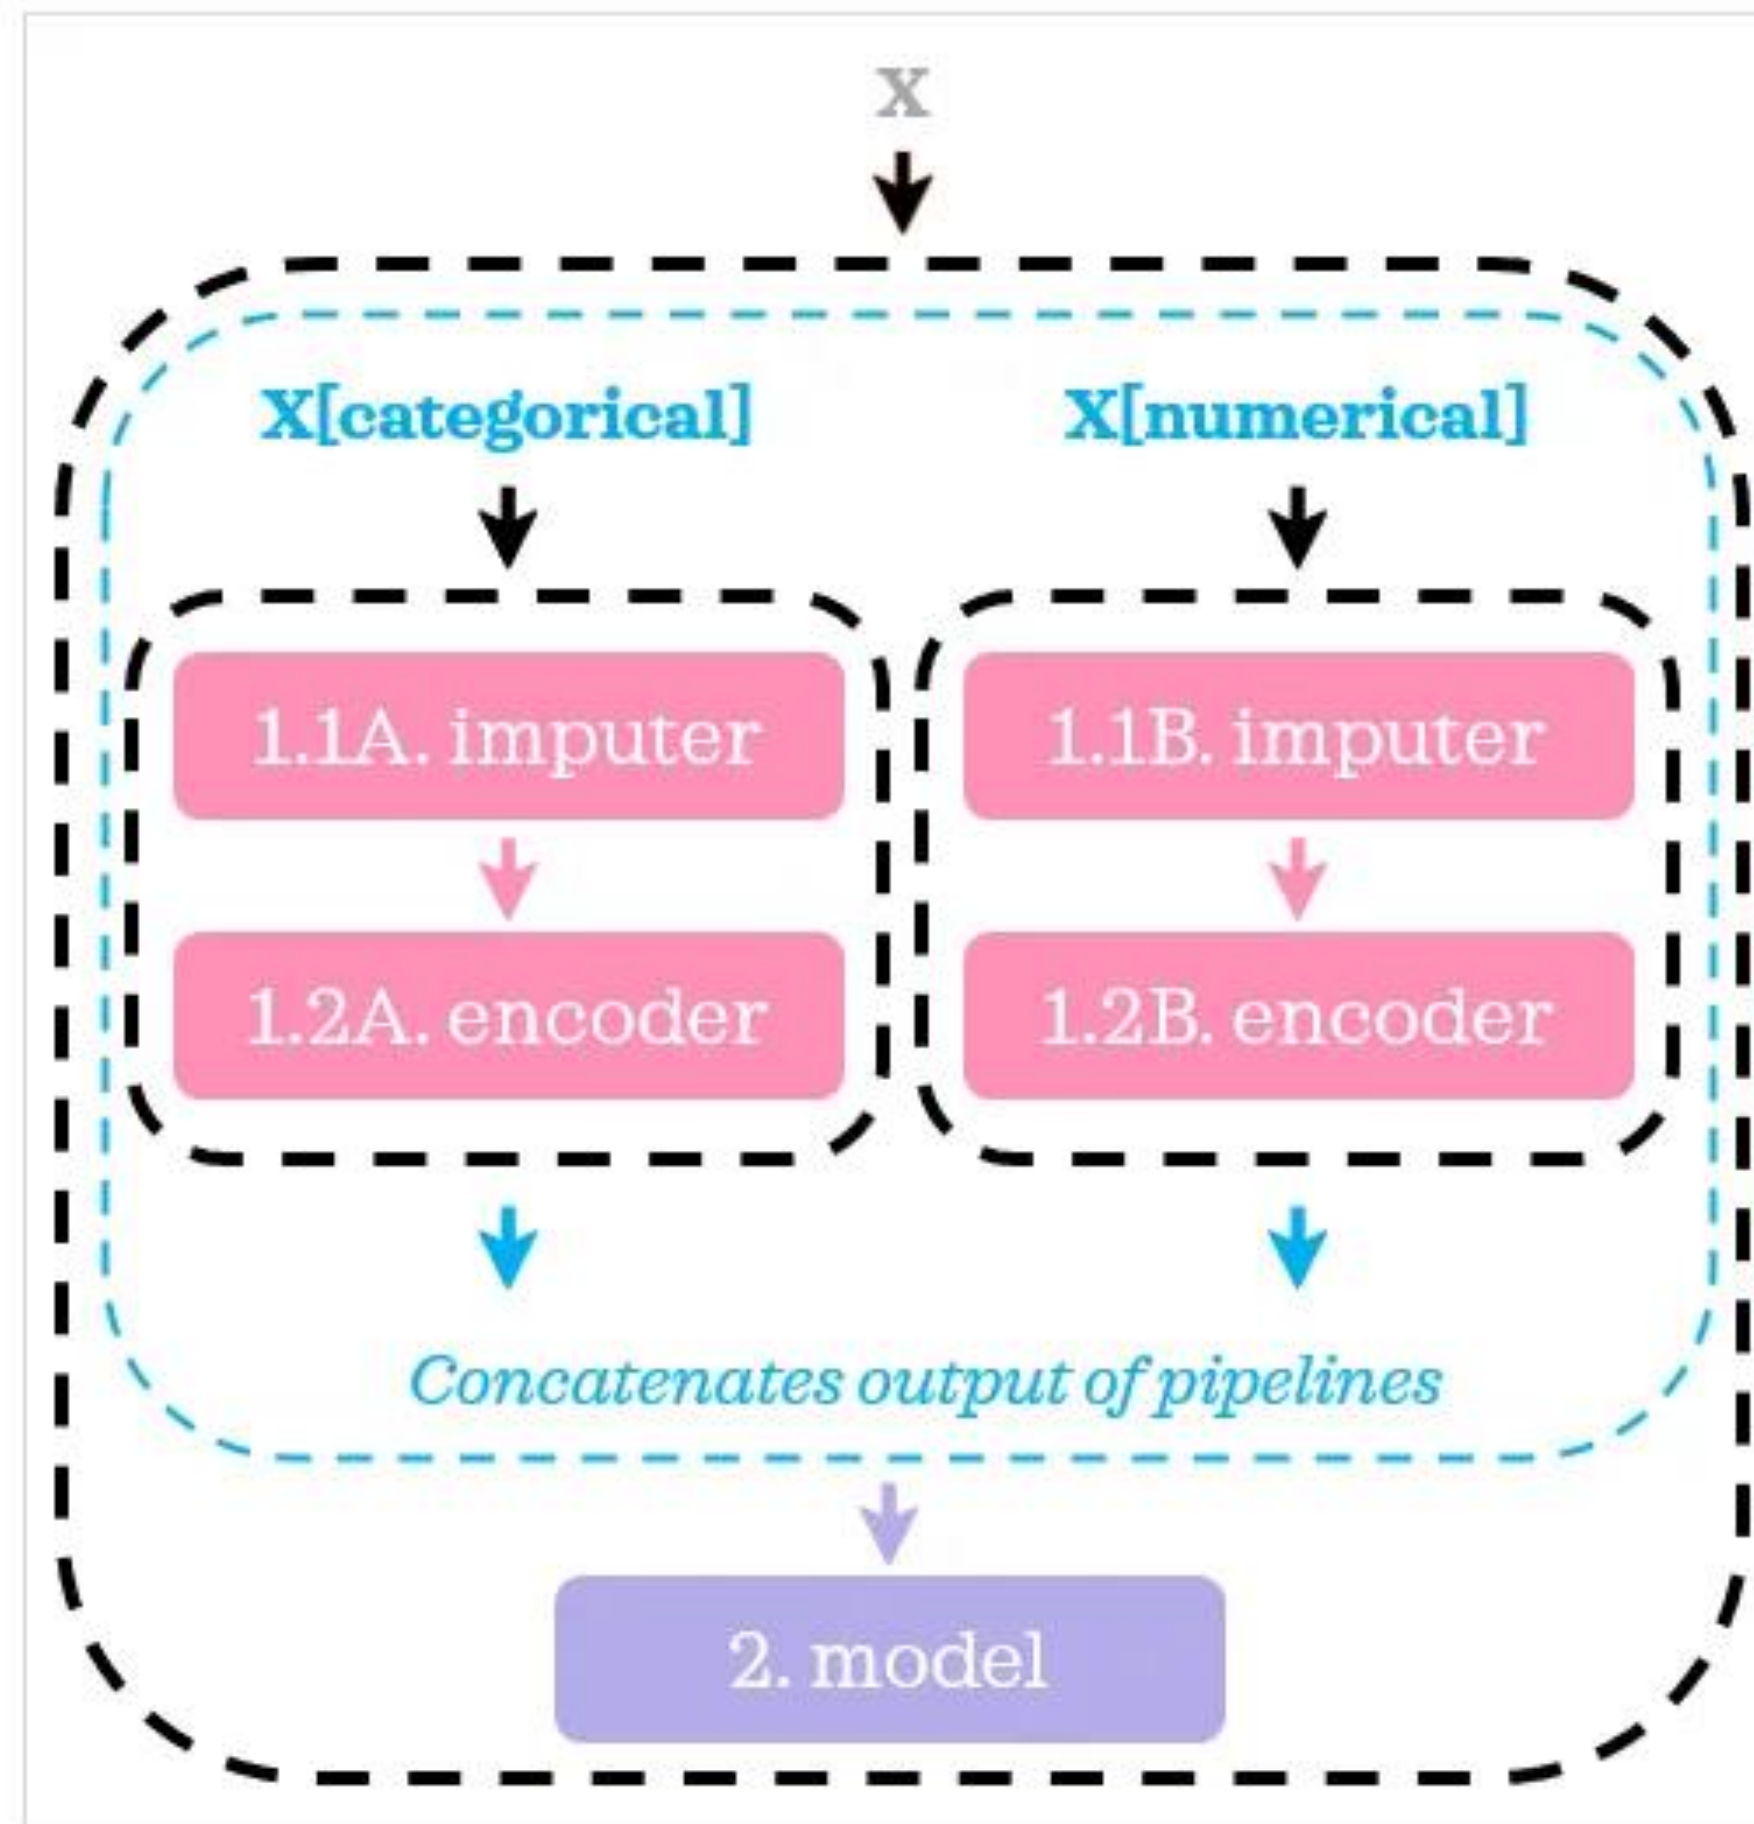


Parallel pipelines

Mixture Model clustering

- Apply different transformations on numerical & categorical features. Then merge back





Colour
code:

transformer

estimator

pipeline

**column
transformer**

```
from sklearn.pipeline import Pipeline
from sklearn.compose import ColumnTransformer

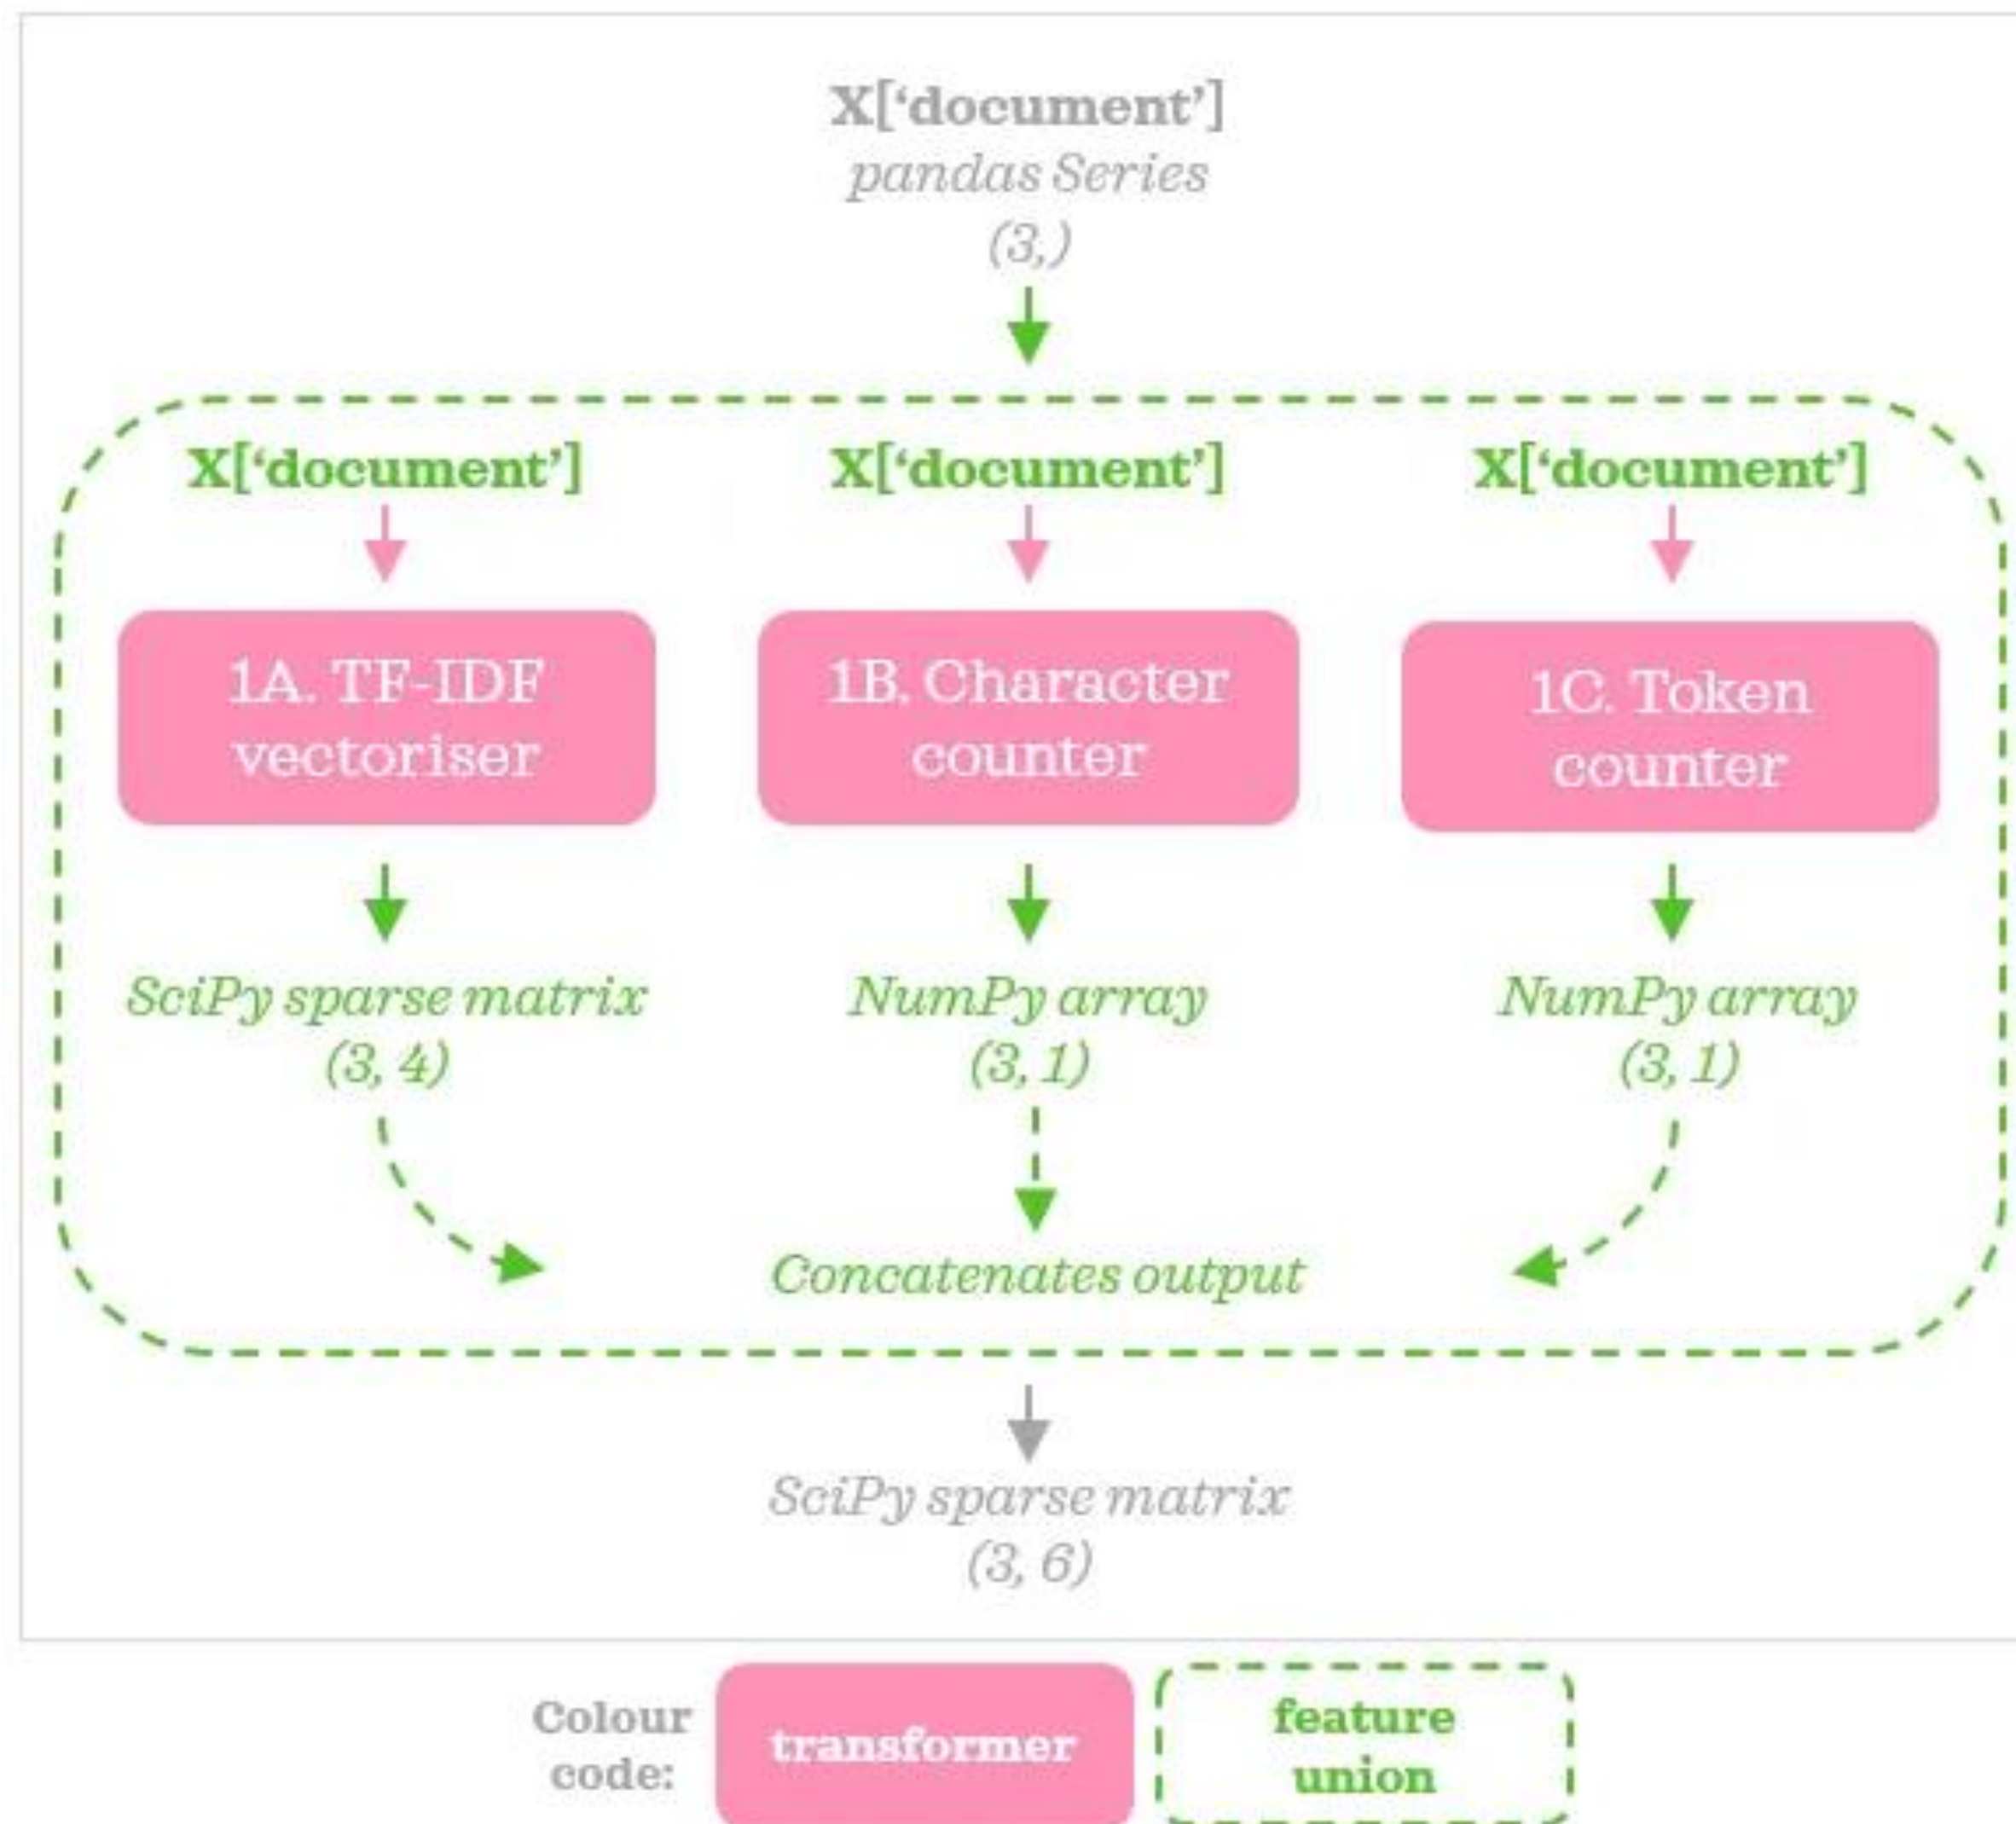
numeric_transformer = Pipeline(steps=[
    ('meanimputer', SimpleImputer(strategy='mean')),
    ('stdscaler', StandardScaler())
])

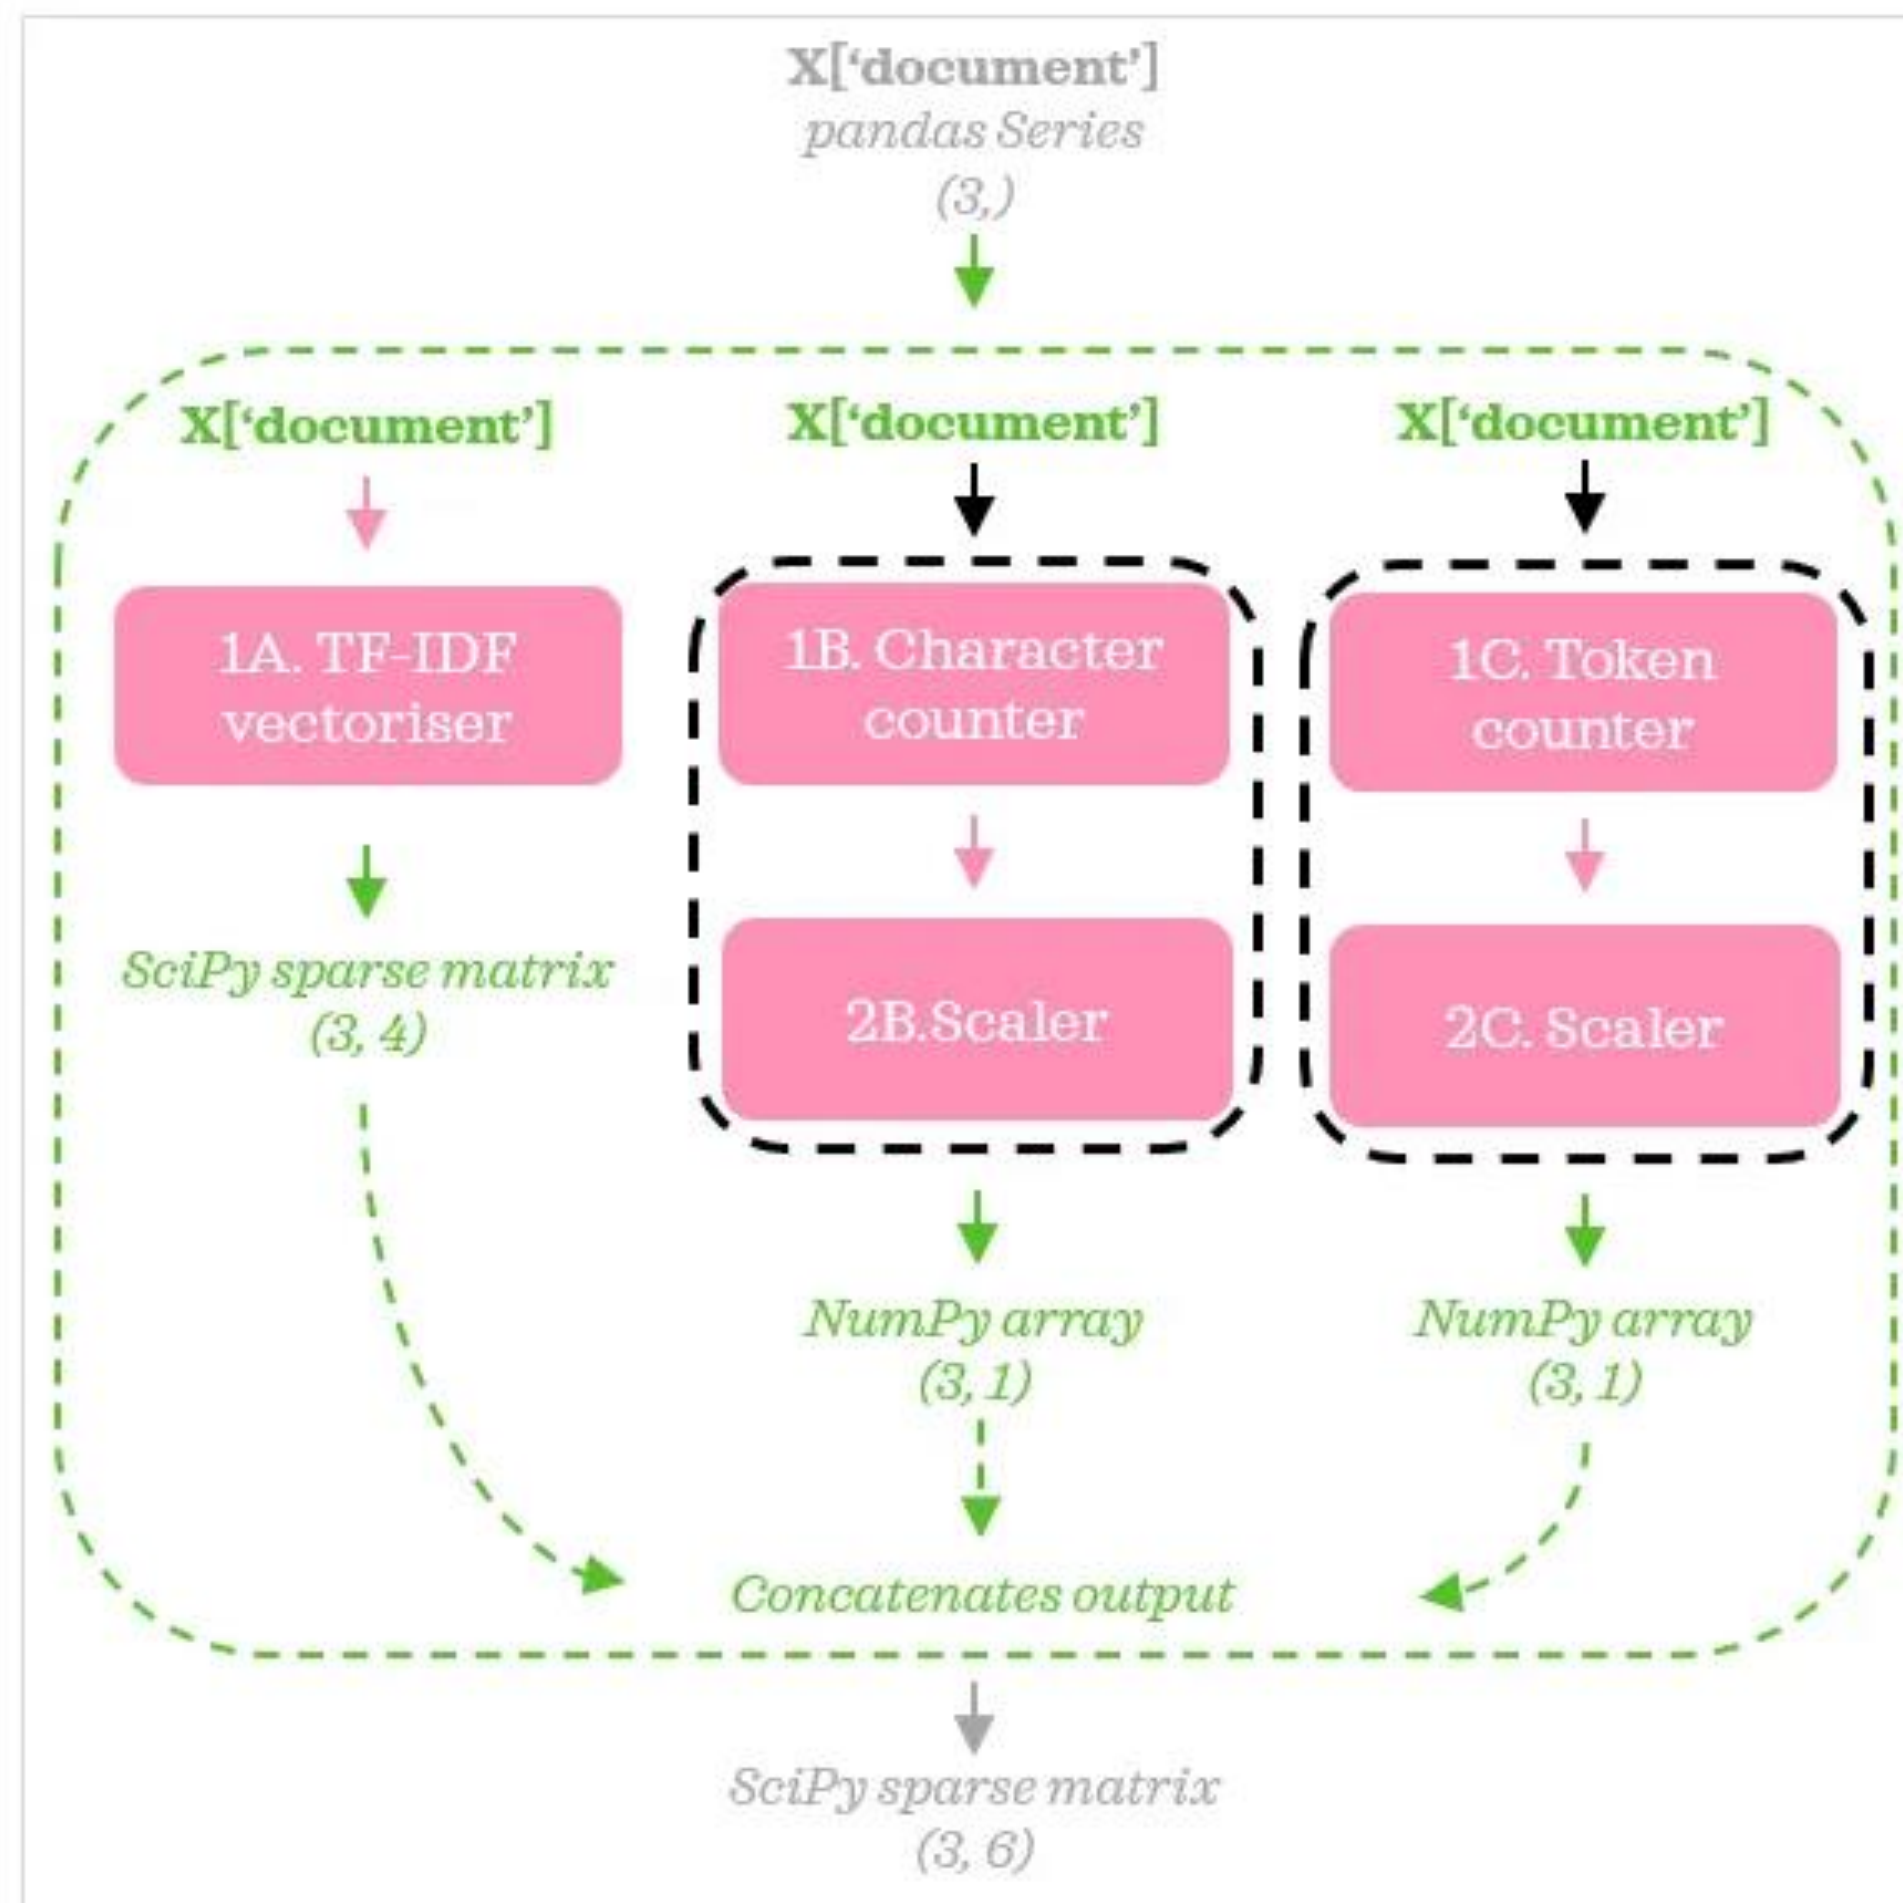
categorical_transformer = Pipeline(steps=[
    ('onehotenc', OneHotEncoder(handle_unknown='ignore'))
])

col_transformer = ColumnTransformer(transformers=[('numeric_processing', numeric_transformer,
                                                  numeric_features),
                                                  ('categorical_processing', categorical_transformer,
                                                  categorical_features)])

pipeline = Pipeline([
    ('transform_column', col_transformer),
    ('logistics', LogisticRegression())
])

pipeline.fit(X_train, y_train)
pipeline.score(X_test, y_test)
```



Colour
code:

transformer

pipeline

feature
union

GridSearch with Pipeline

```
from sklearn.model_selection import GridSearchCV

X_train, X_test, y_train, y_test = train_test_split(df, y, random_state=42)

#create new a knn model
knn = KNeighborsClassifier()
#create a dictionary of all values we want to test for n_neighbors
param_grid = {"n_neighbors": np.arange(1, 25),
              "weights": ["uniform", "distance"]}

loocv = LeaveOneOut()

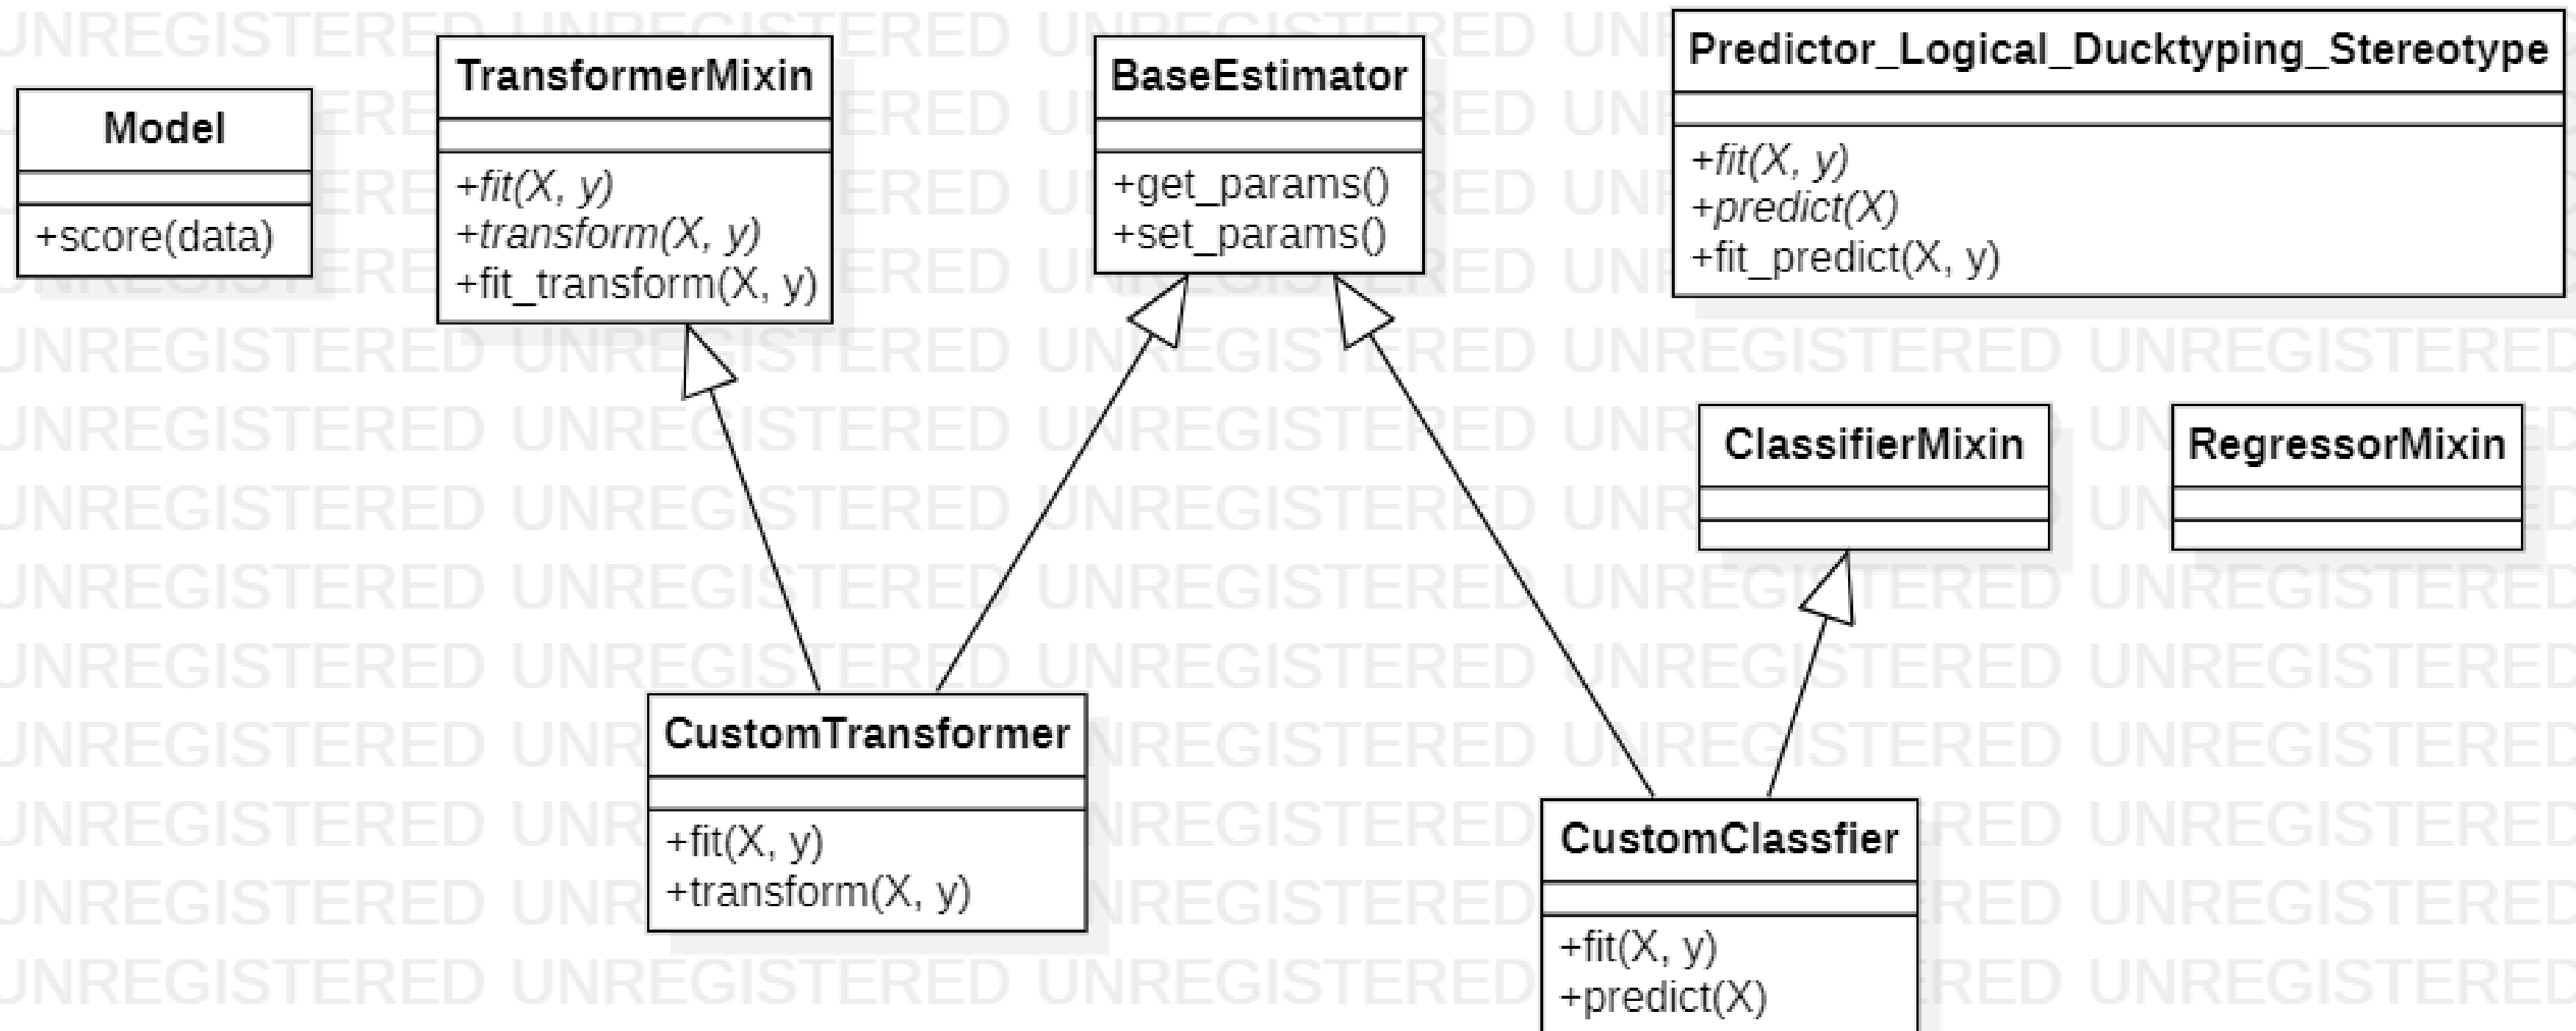
#use gridsearch to test all values for n_neighbors
knn_gscv = GridSearchCV(knn, param_grid, cv=loocv)
#fit model to data
knn_gscv.fit(X_train, y_train)
```

**Replace with
pipeline**

Odds and ends

- PowerTransformer
- FunctionTransformer
- BYOT – Bring Your Own Transformer
 - Custom Transformer, Custom Predictor

Custom Transformer & Classifier/Regressor





QUESTIONS