



# Lecture 21 & 22: Linear Regression

## Part 1

# Recap

- Evaluation metrics
- Entropy
- Join Entropy
- Conditional Entropy
- Mutual Information (Information Gain)
- Use in Decision Trees

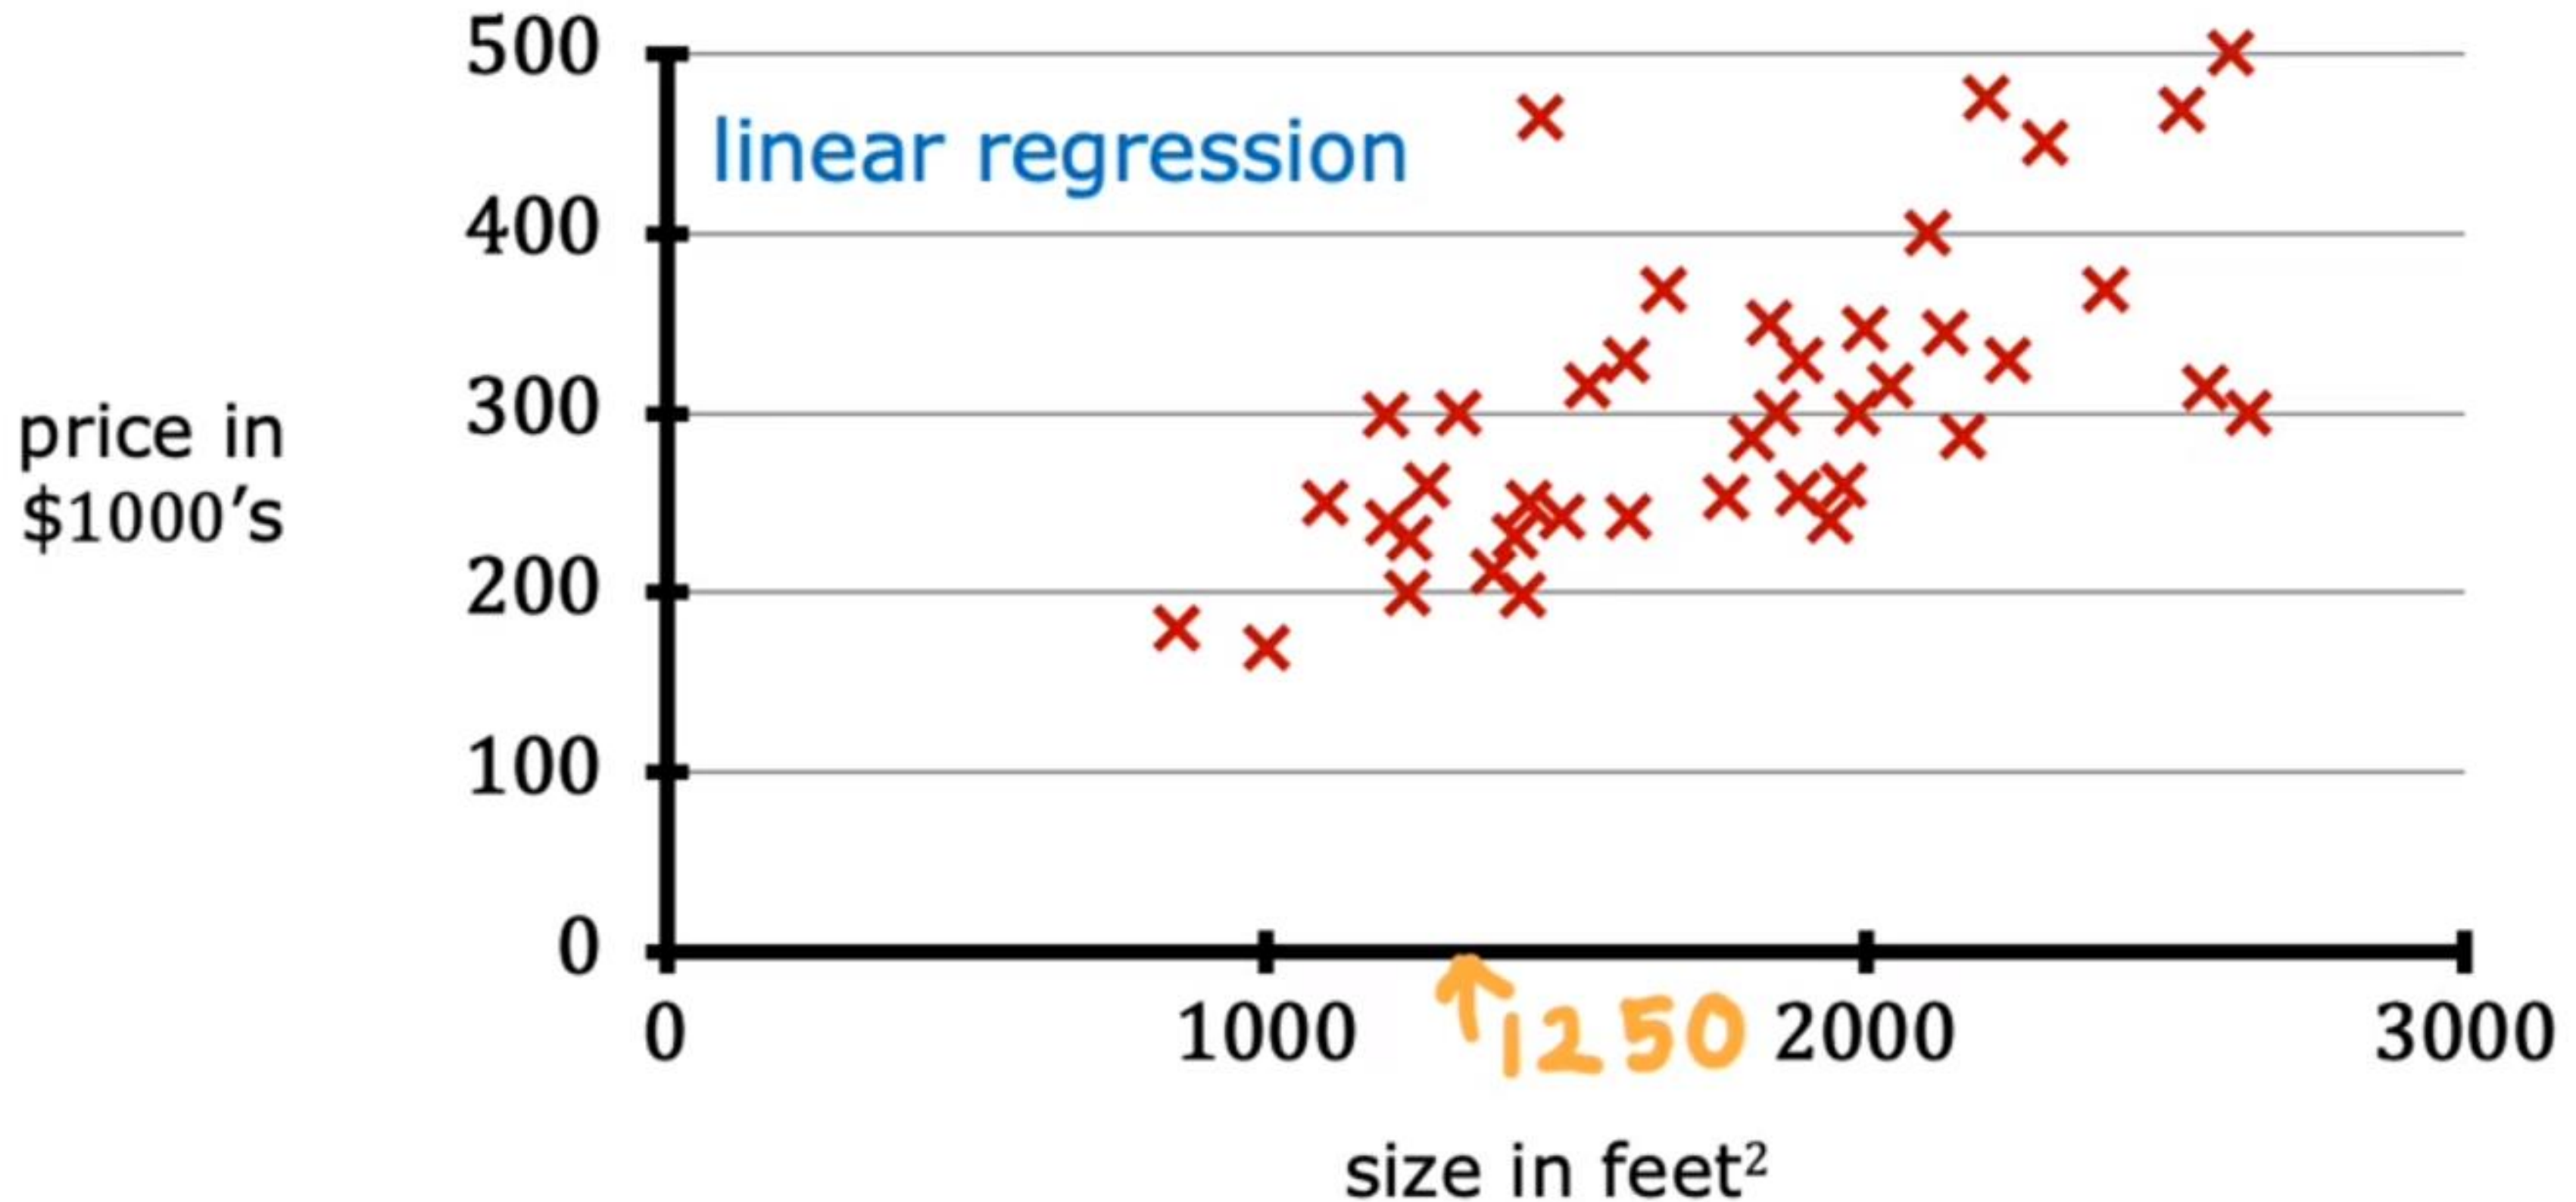




# Simple Linear Regression

# Univariate Linear Regression

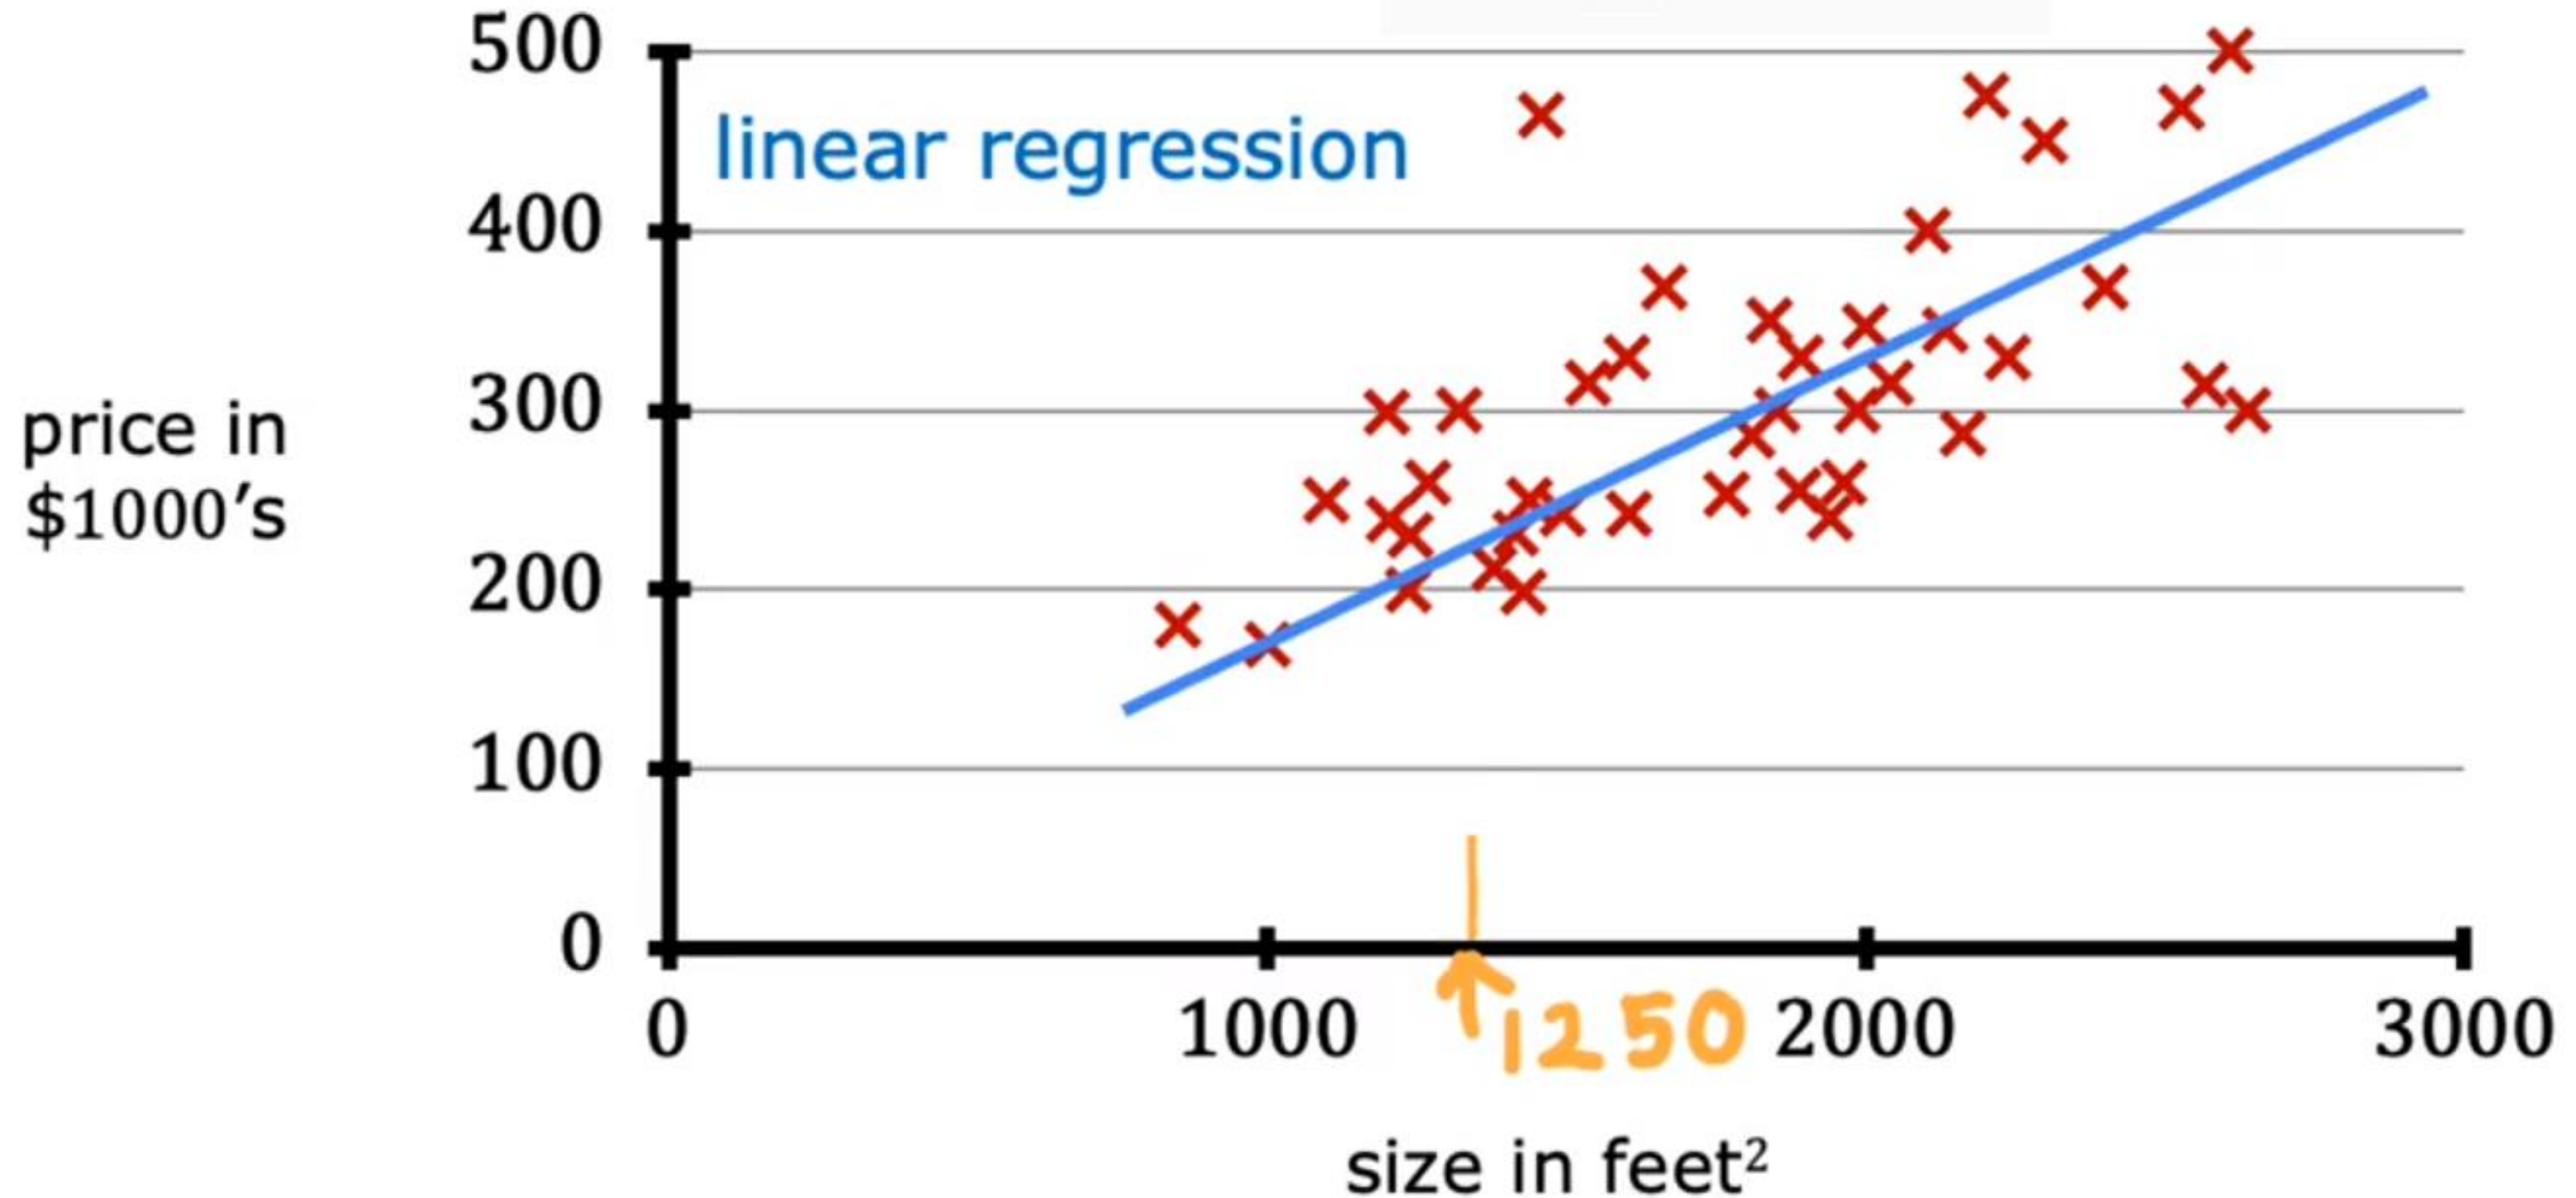
House sizes and prices





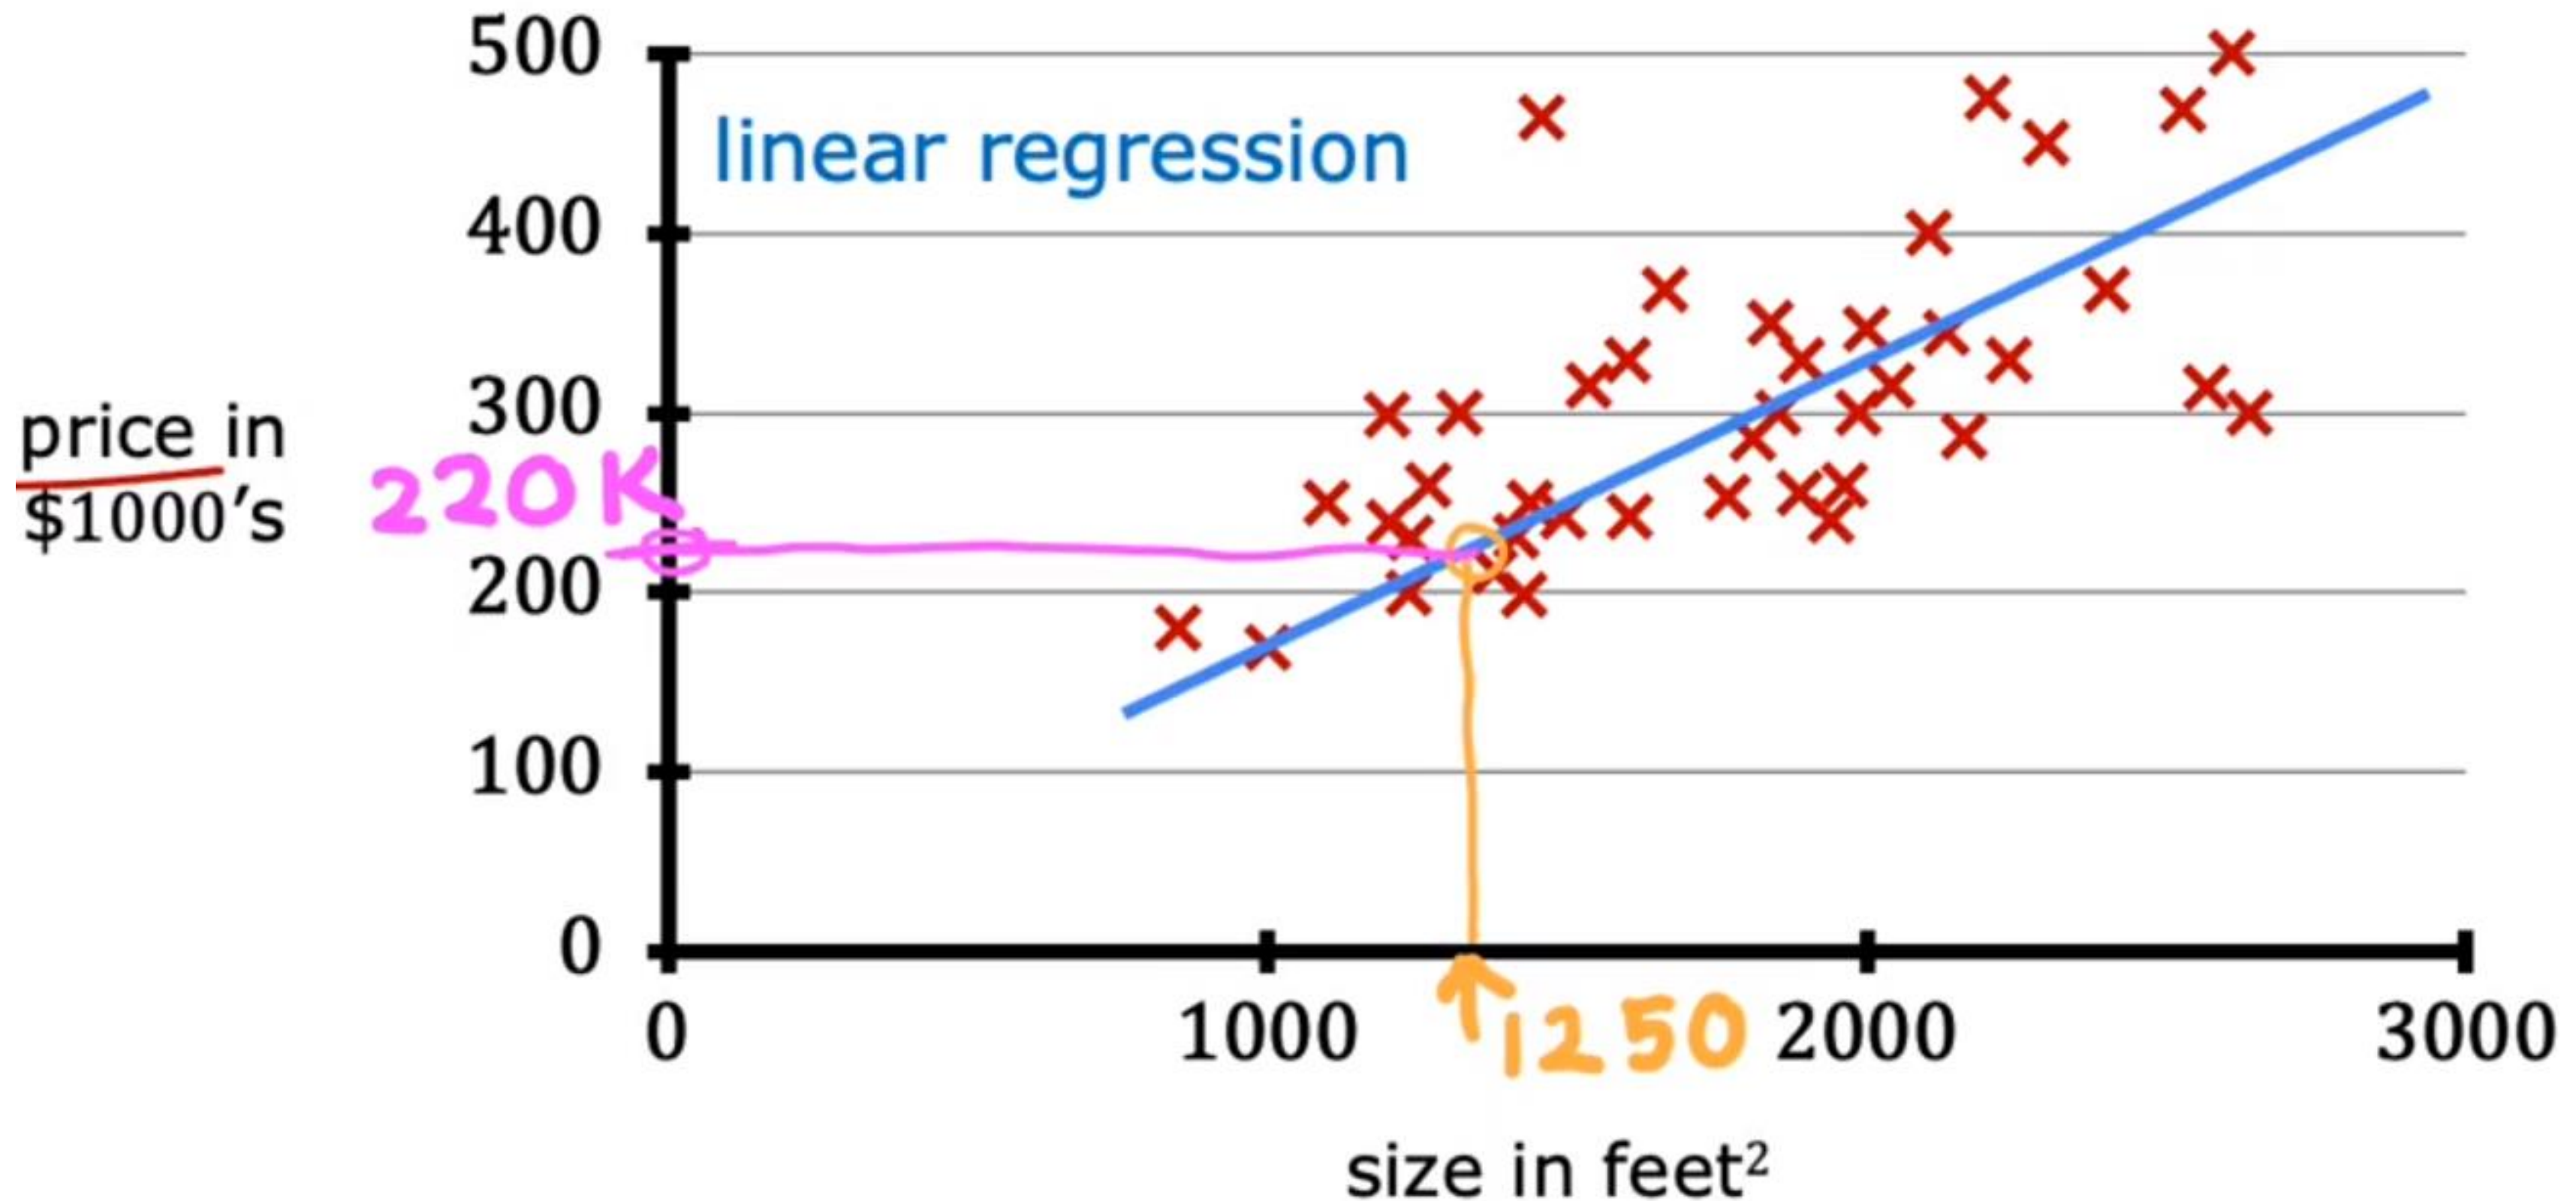
# Univariate Linear Regression

## House sizes and prices



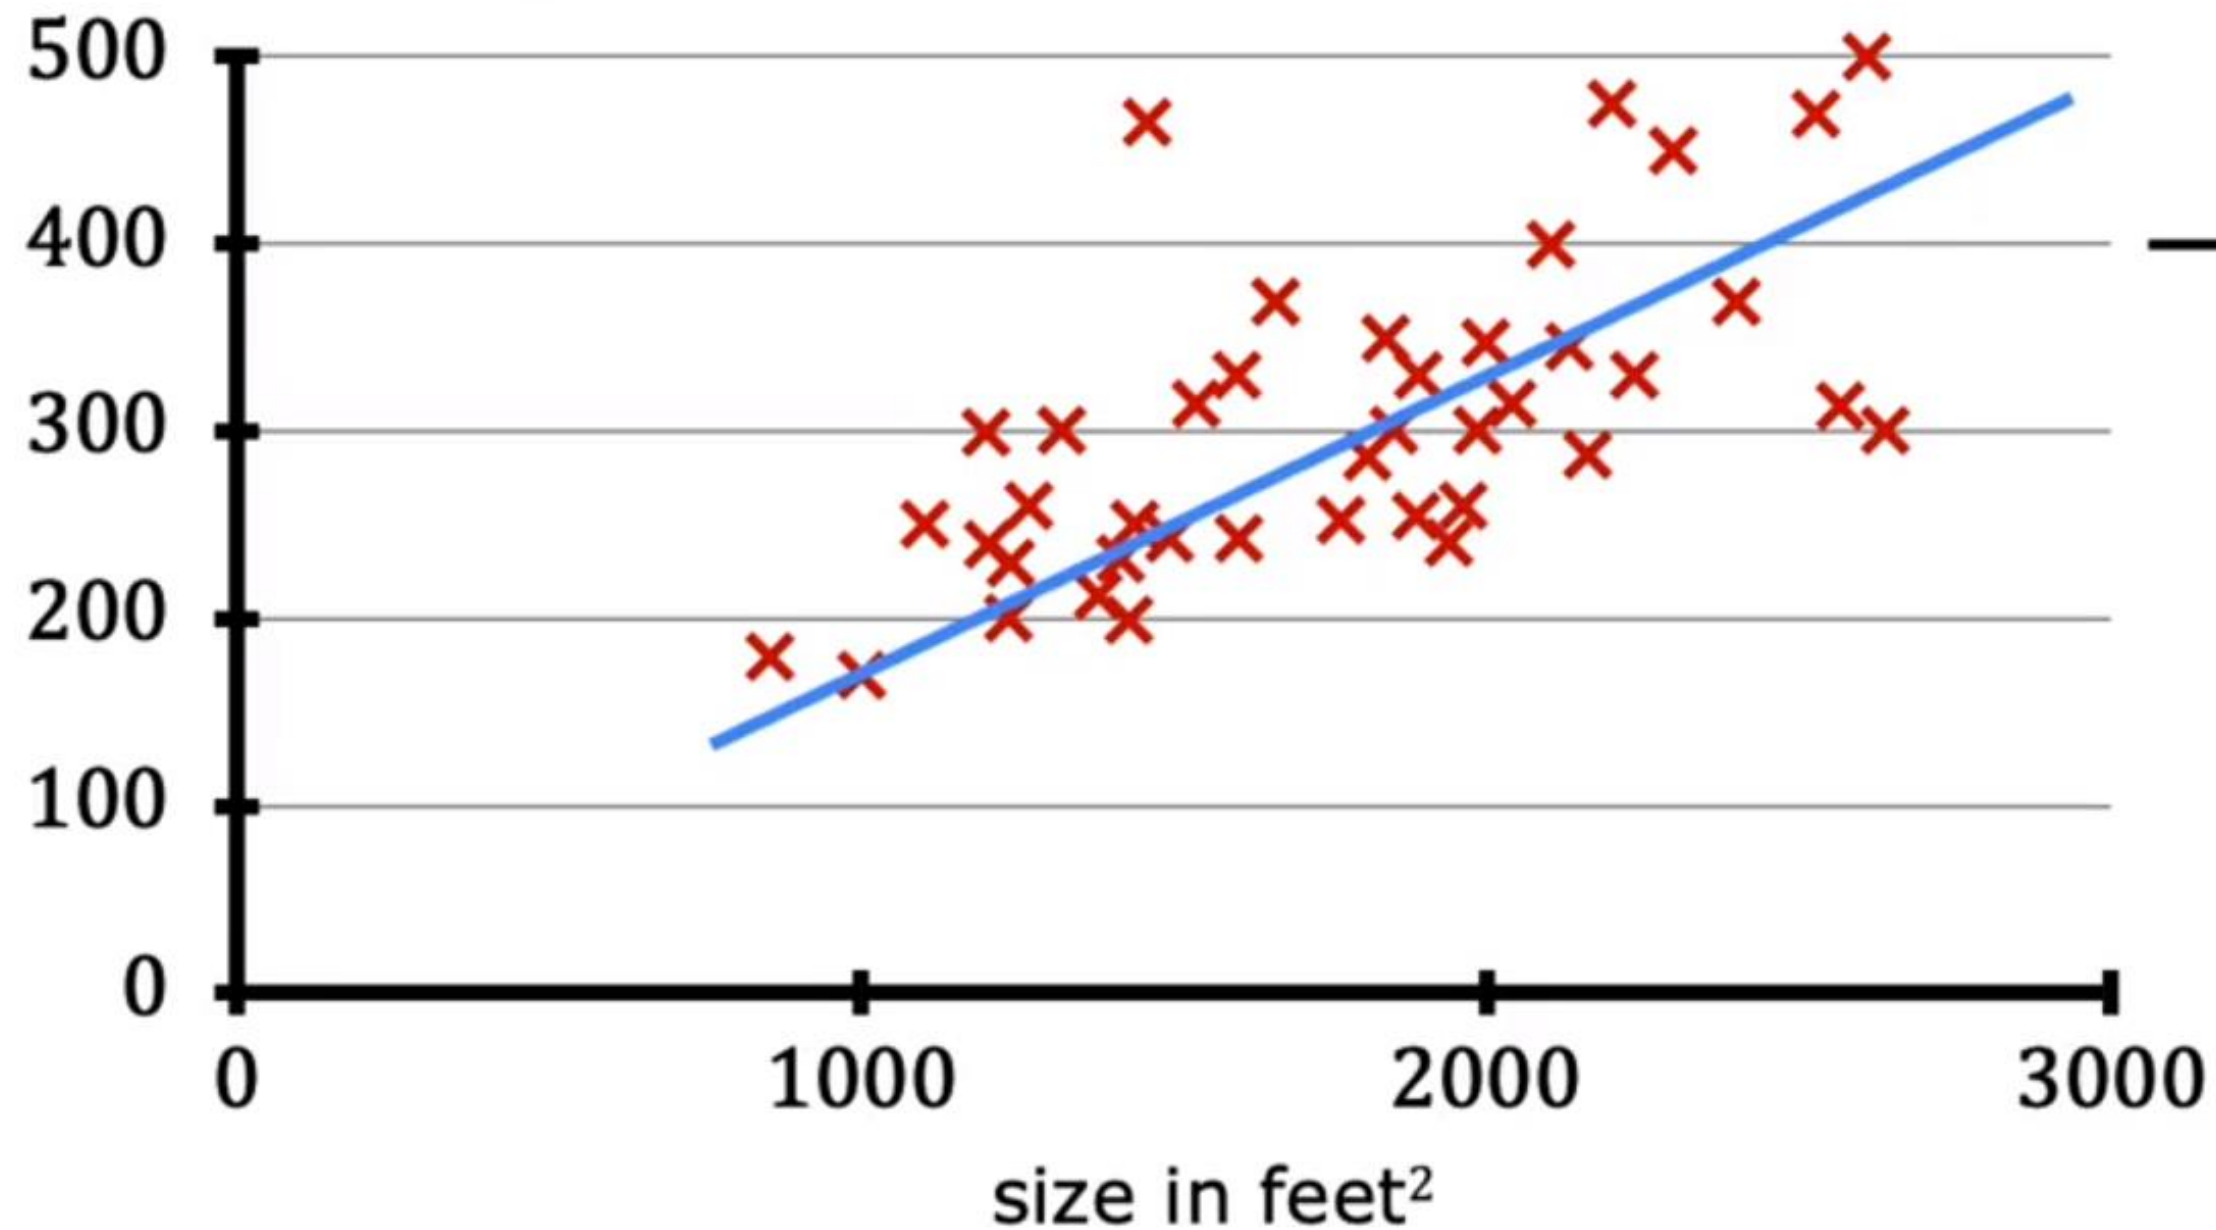
# Univariate Linear Regression

House sizes and prices



# Univariate Linear Regression

House sizes and prices



Data table

size in feet <sup>2</sup>	price in \$1000's
2104	400
1416	232
1534	315
852	178
...	...
3210	870

# Population versus Sample

- Population Regression Function

- Deterministic Component  $y = f(x)$

- Stochastic Component  $y = f(x) + \epsilon$

- Normally distributed error component

- Univariate function  $y = wx + b + \epsilon$

- Multivariate function

$$y = w_n x_n + \dots w_1 x_1 + w_0 + \epsilon$$

$$= \mathbf{w}^T \mathbf{x} + \epsilon$$

- Population Regression Line/Plane/Hyperplane

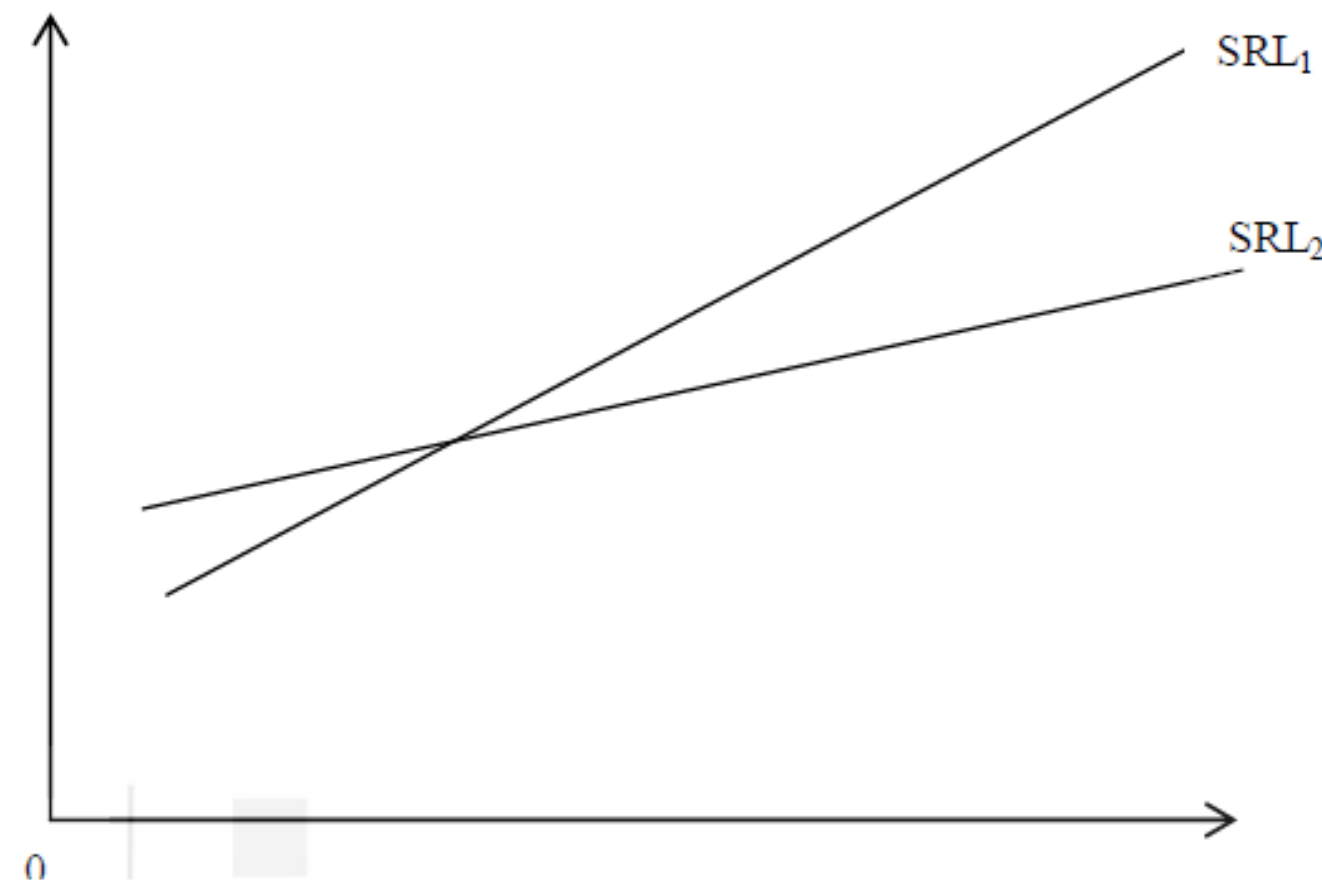
$$\mathbf{w} = \begin{bmatrix} w_0 \\ w_1 \\ \dots \\ w_n \end{bmatrix}$$

$$\mathbf{x} = \begin{bmatrix} 1 \\ x_1 \\ \dots \\ x_n \end{bmatrix}$$



# Population versus Sample

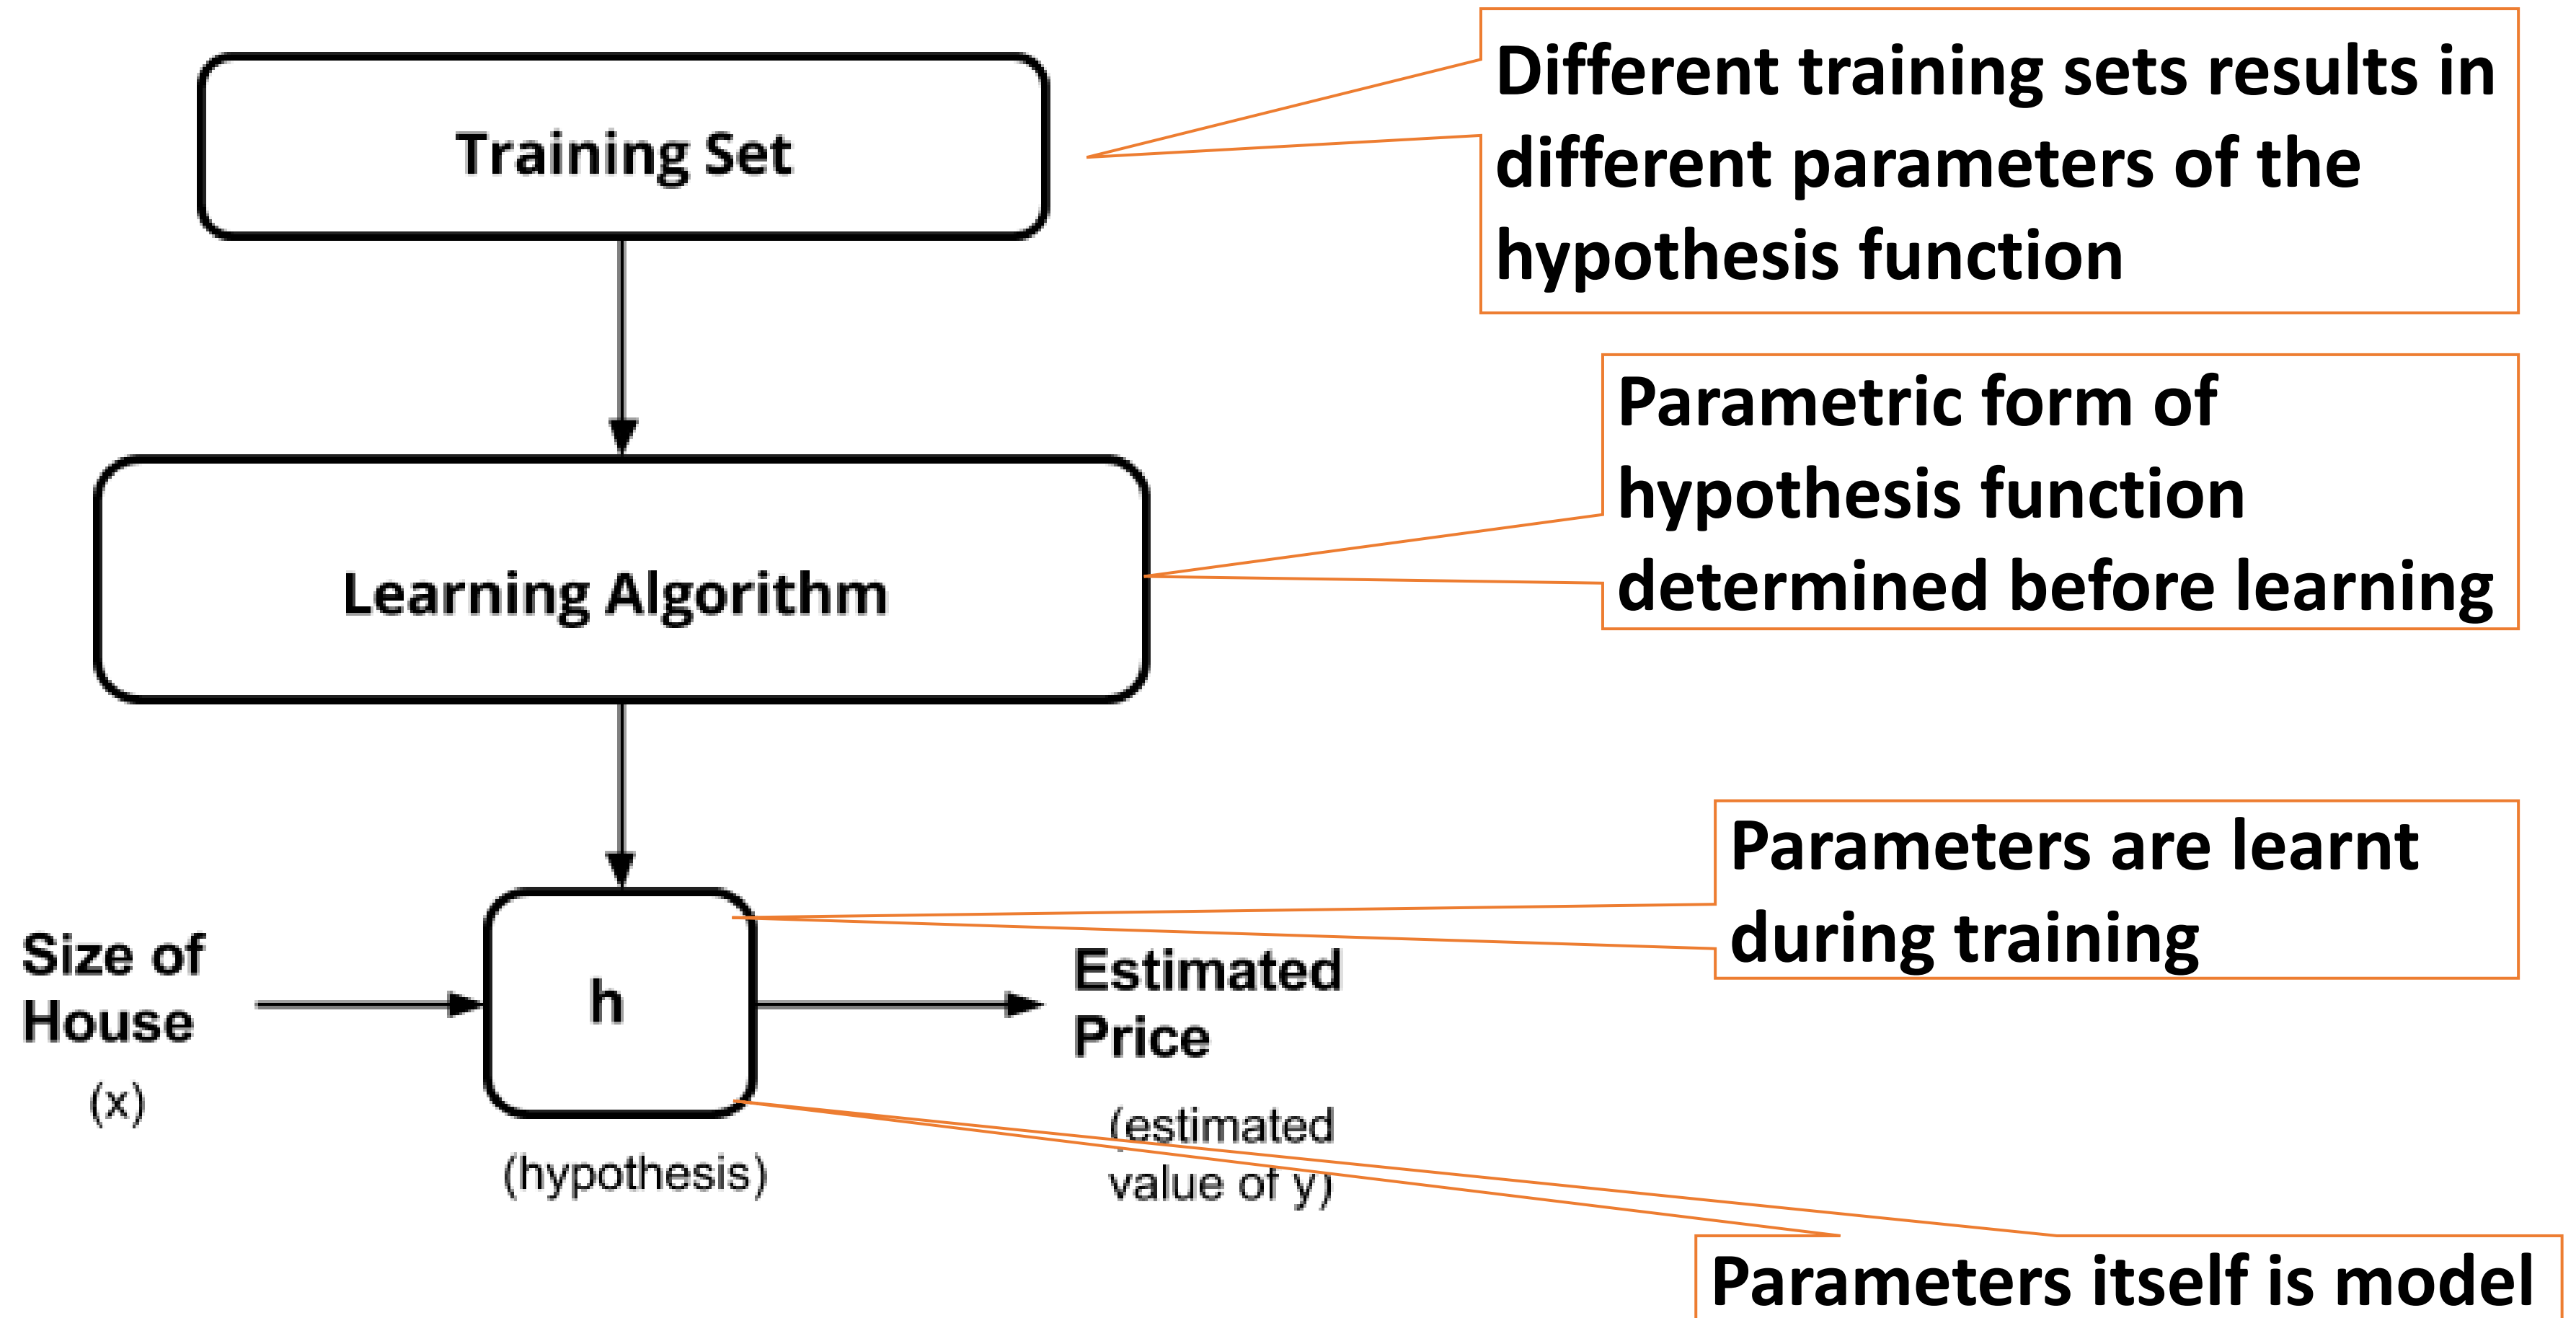
- Sample Regression Functions
  - Different Regression Line/Plane/Hyperplane



Hypothesis  
function(s)

- Univariate function  $\hat{y} = h(x) = wx + b$
- Multivariate function  $\hat{y} = h(x) = \mathbf{w}^T \mathbf{x}$

# Hypothesis function(s)





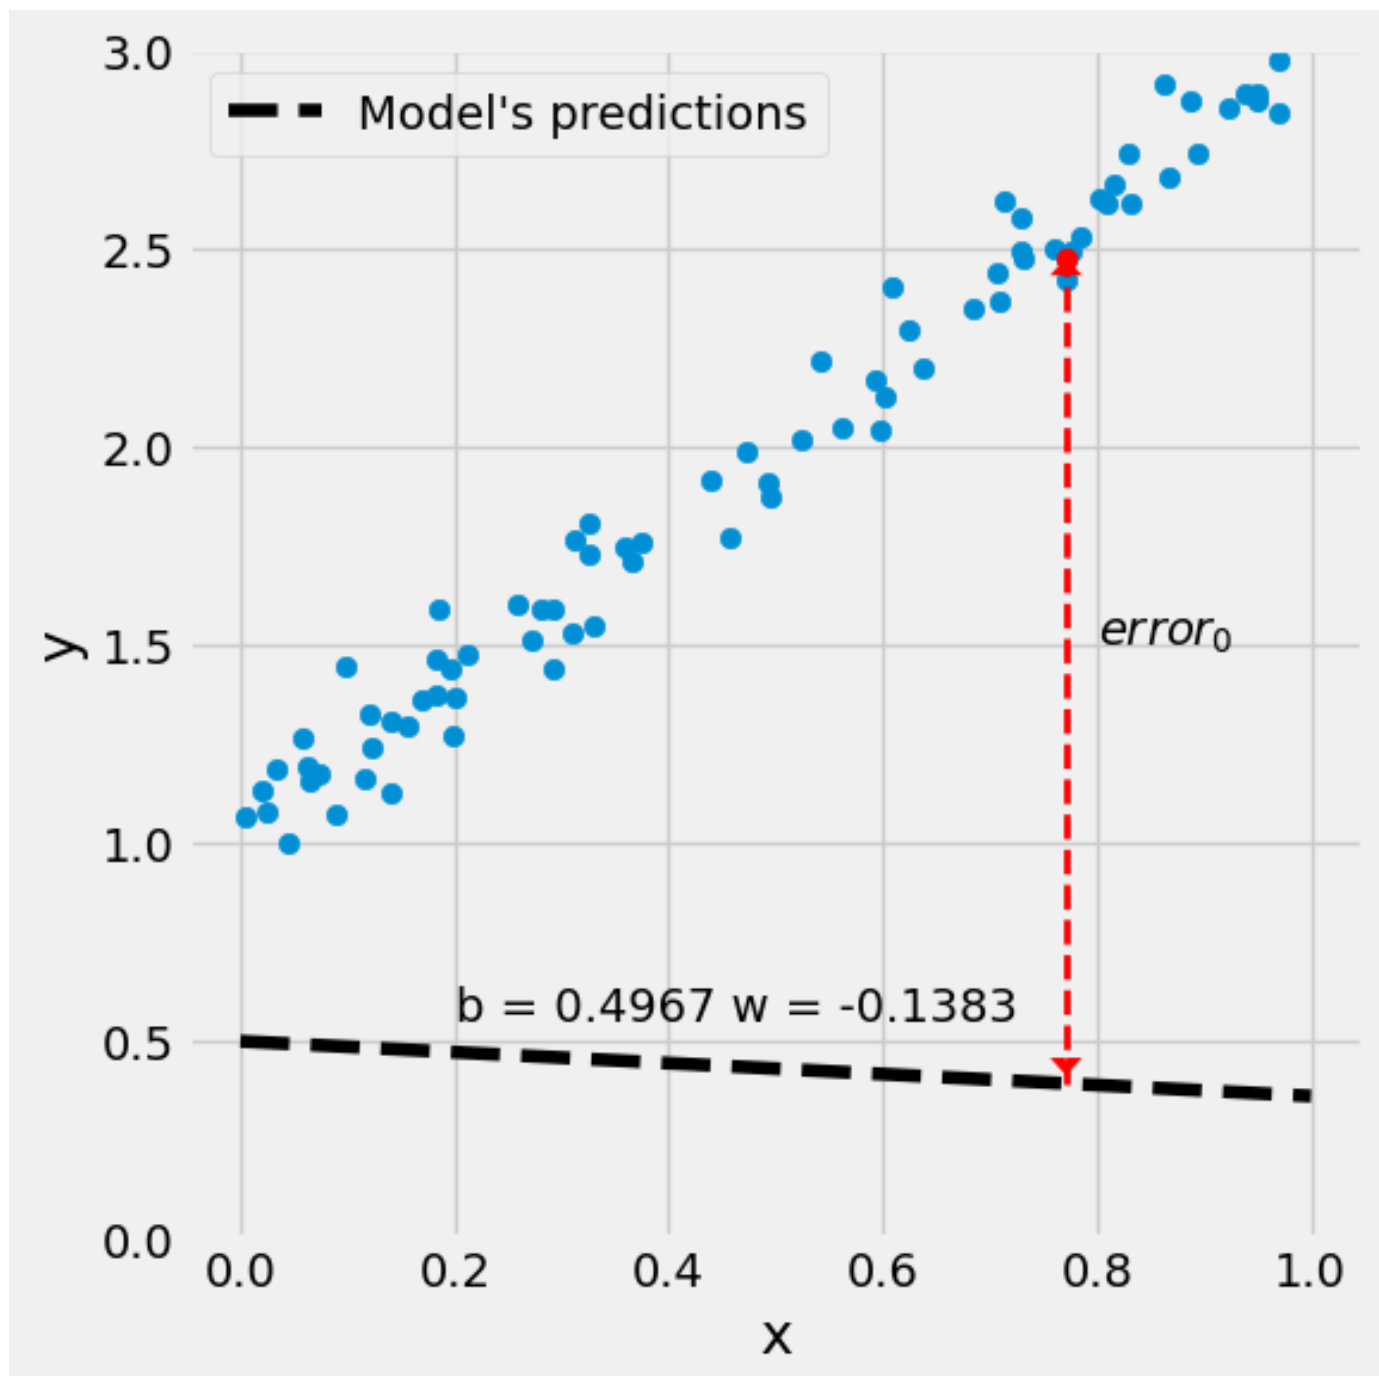
# Linear Regression algorithm



# Step 1. Initialization

- Assume parametric form  $\hat{y} = wx + b$
- Assign random values for  $w$  and  $b$

Parametric  
form of the  
Hypothesis  
function



- Calculate error using the initial  $w$  &  $b$

$$error^{(i)} = \hat{y}^{(i)} - y^{(i)}$$

# Formulate Objective function

Also called  
cost / loss  
function

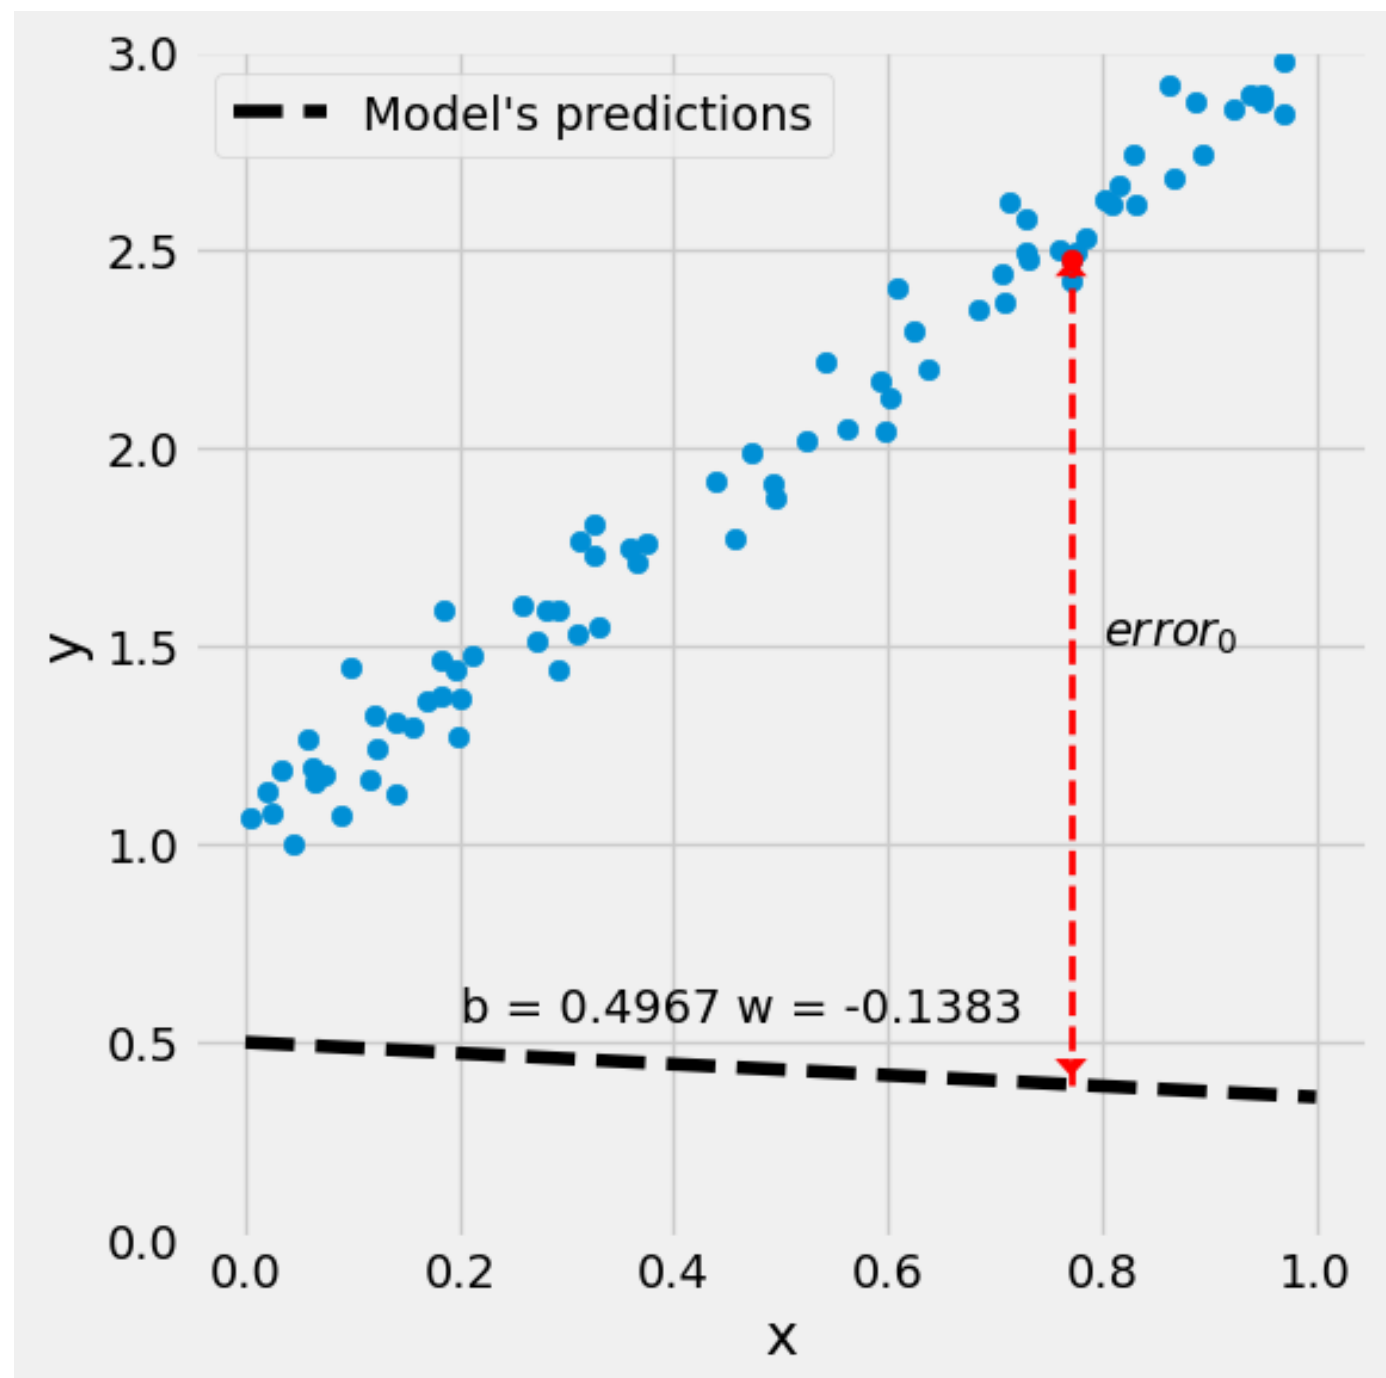
$$error^{(i)} = \hat{y}^{(i)} - y^{(i)}$$

$$MSE = \frac{1}{n} \sum_{i=1}^n error^{(i)2}$$

How do I  
quantify my  
unhappiness

$$= \frac{1}{n} \sum_{i=1}^n (\hat{y}^{(i)} - y^{(i)})^2$$


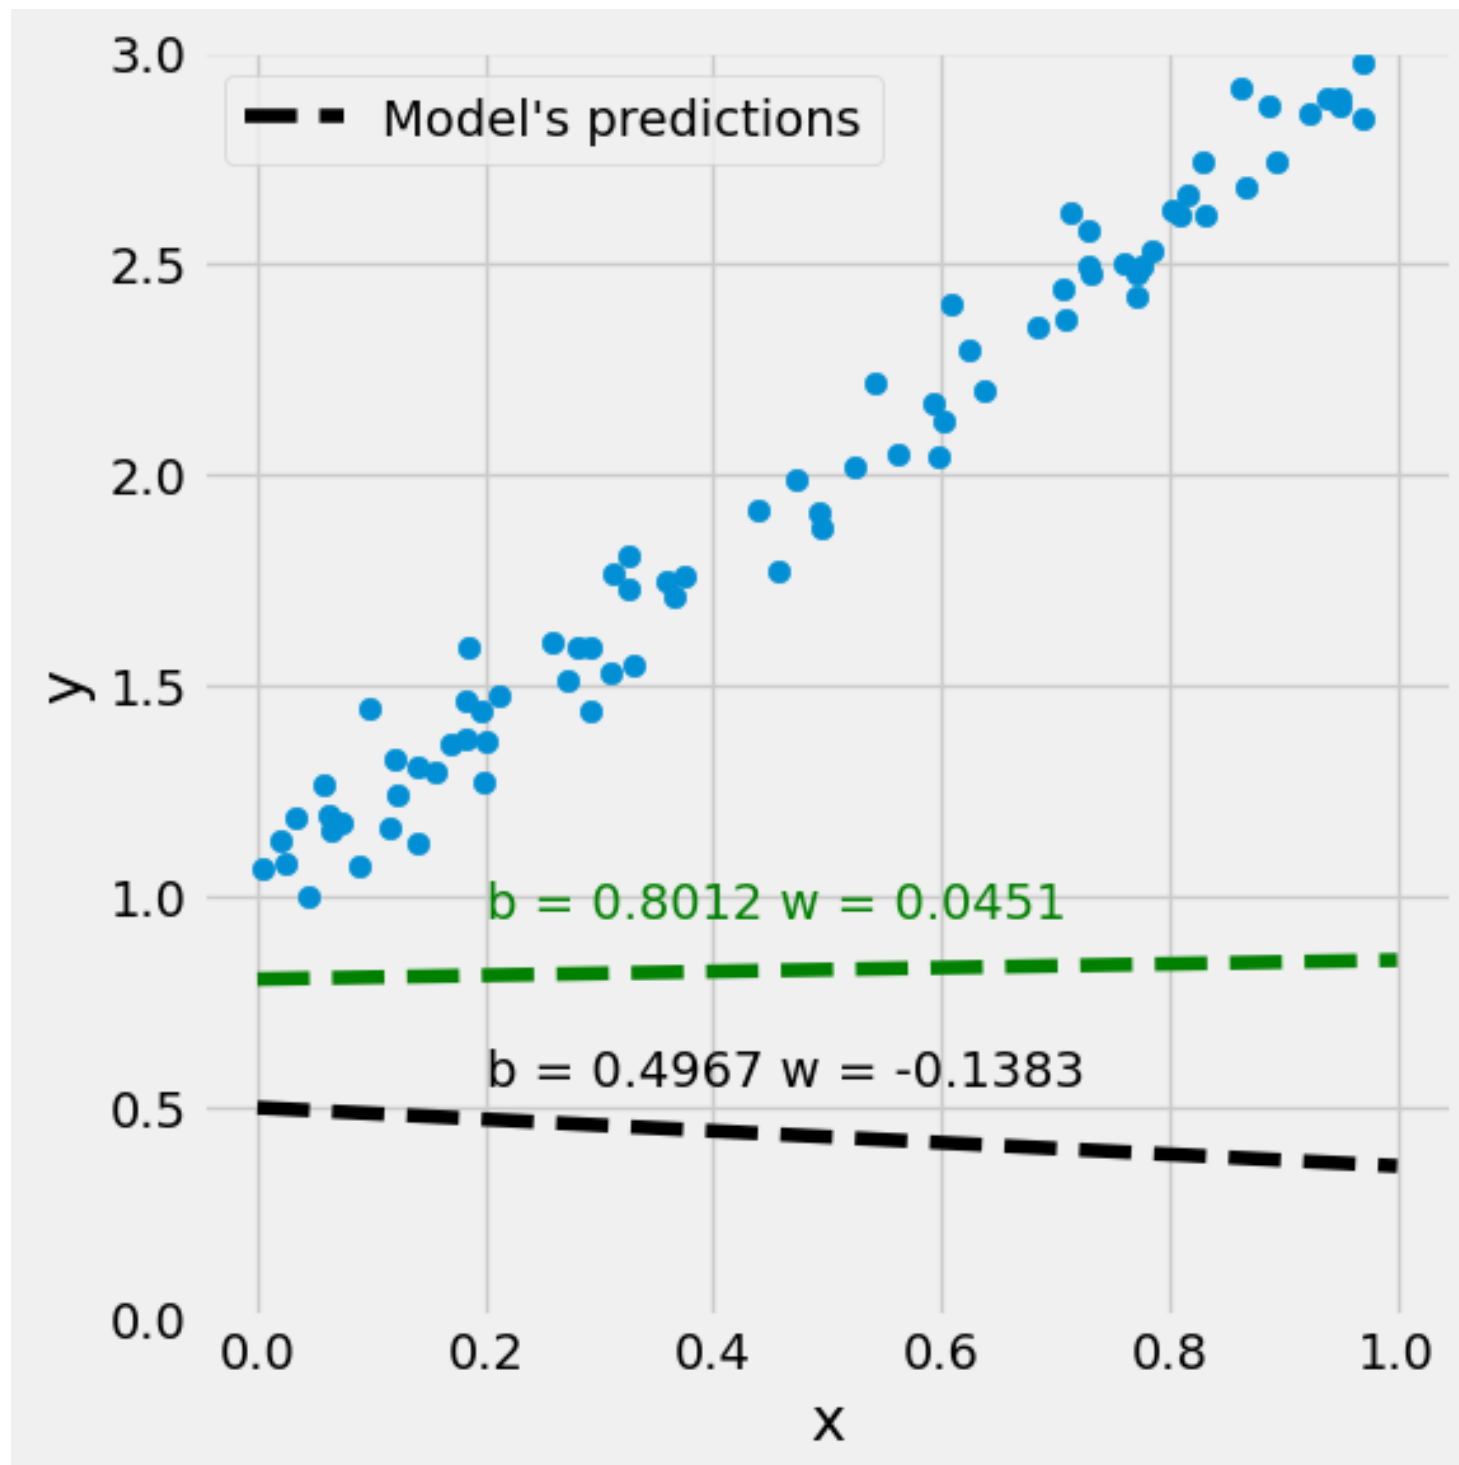
$$\mathcal{J}(w, b) = \frac{1}{n} \sum_{i=1}^n (b + wx^{(i)} - y^{(i)})^2$$



# Feature space versus parameter space

$$\hat{y} = wx + b$$

$$\mathcal{J}(w, b) = \frac{1}{n} \sum_{i=1}^n (b + wx^{(i)} - y^{(i)})^2$$



Find the  $w$  &  $b$   
when the cost is  
minimum

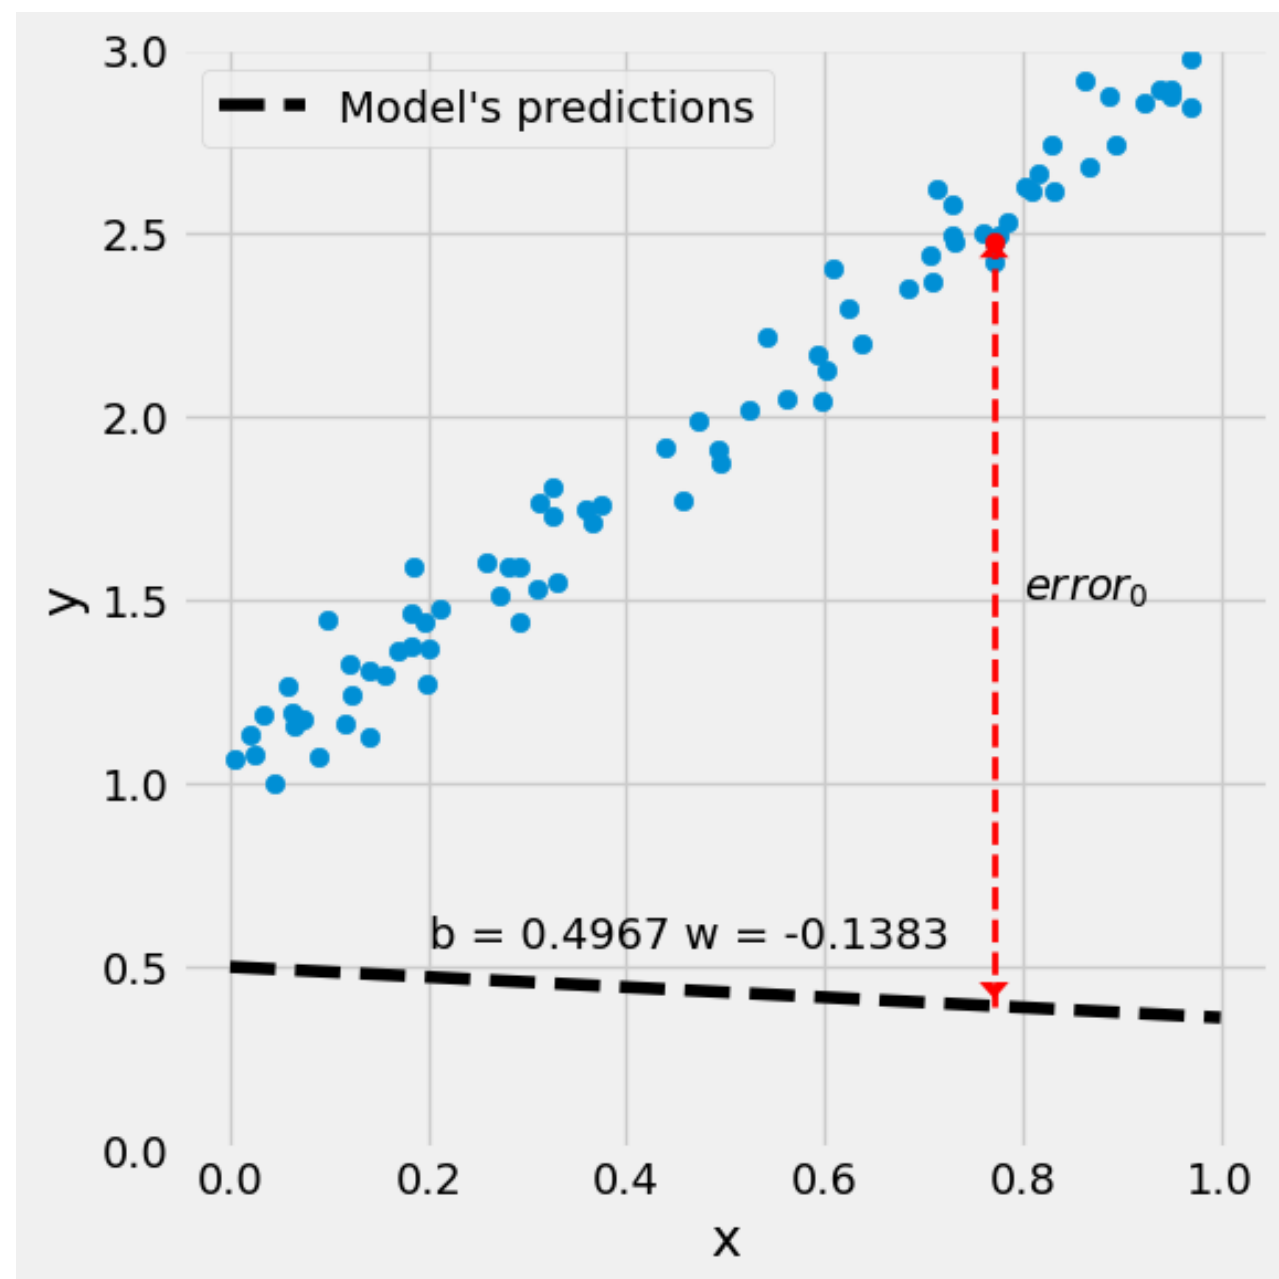


## Step 2. Evaluate $\hat{y}$ & evaluate objective function

- Evaluate  $\hat{y}$   $\hat{y} = wx + b$
- Evaluate Objective function
- Also known as forward pass

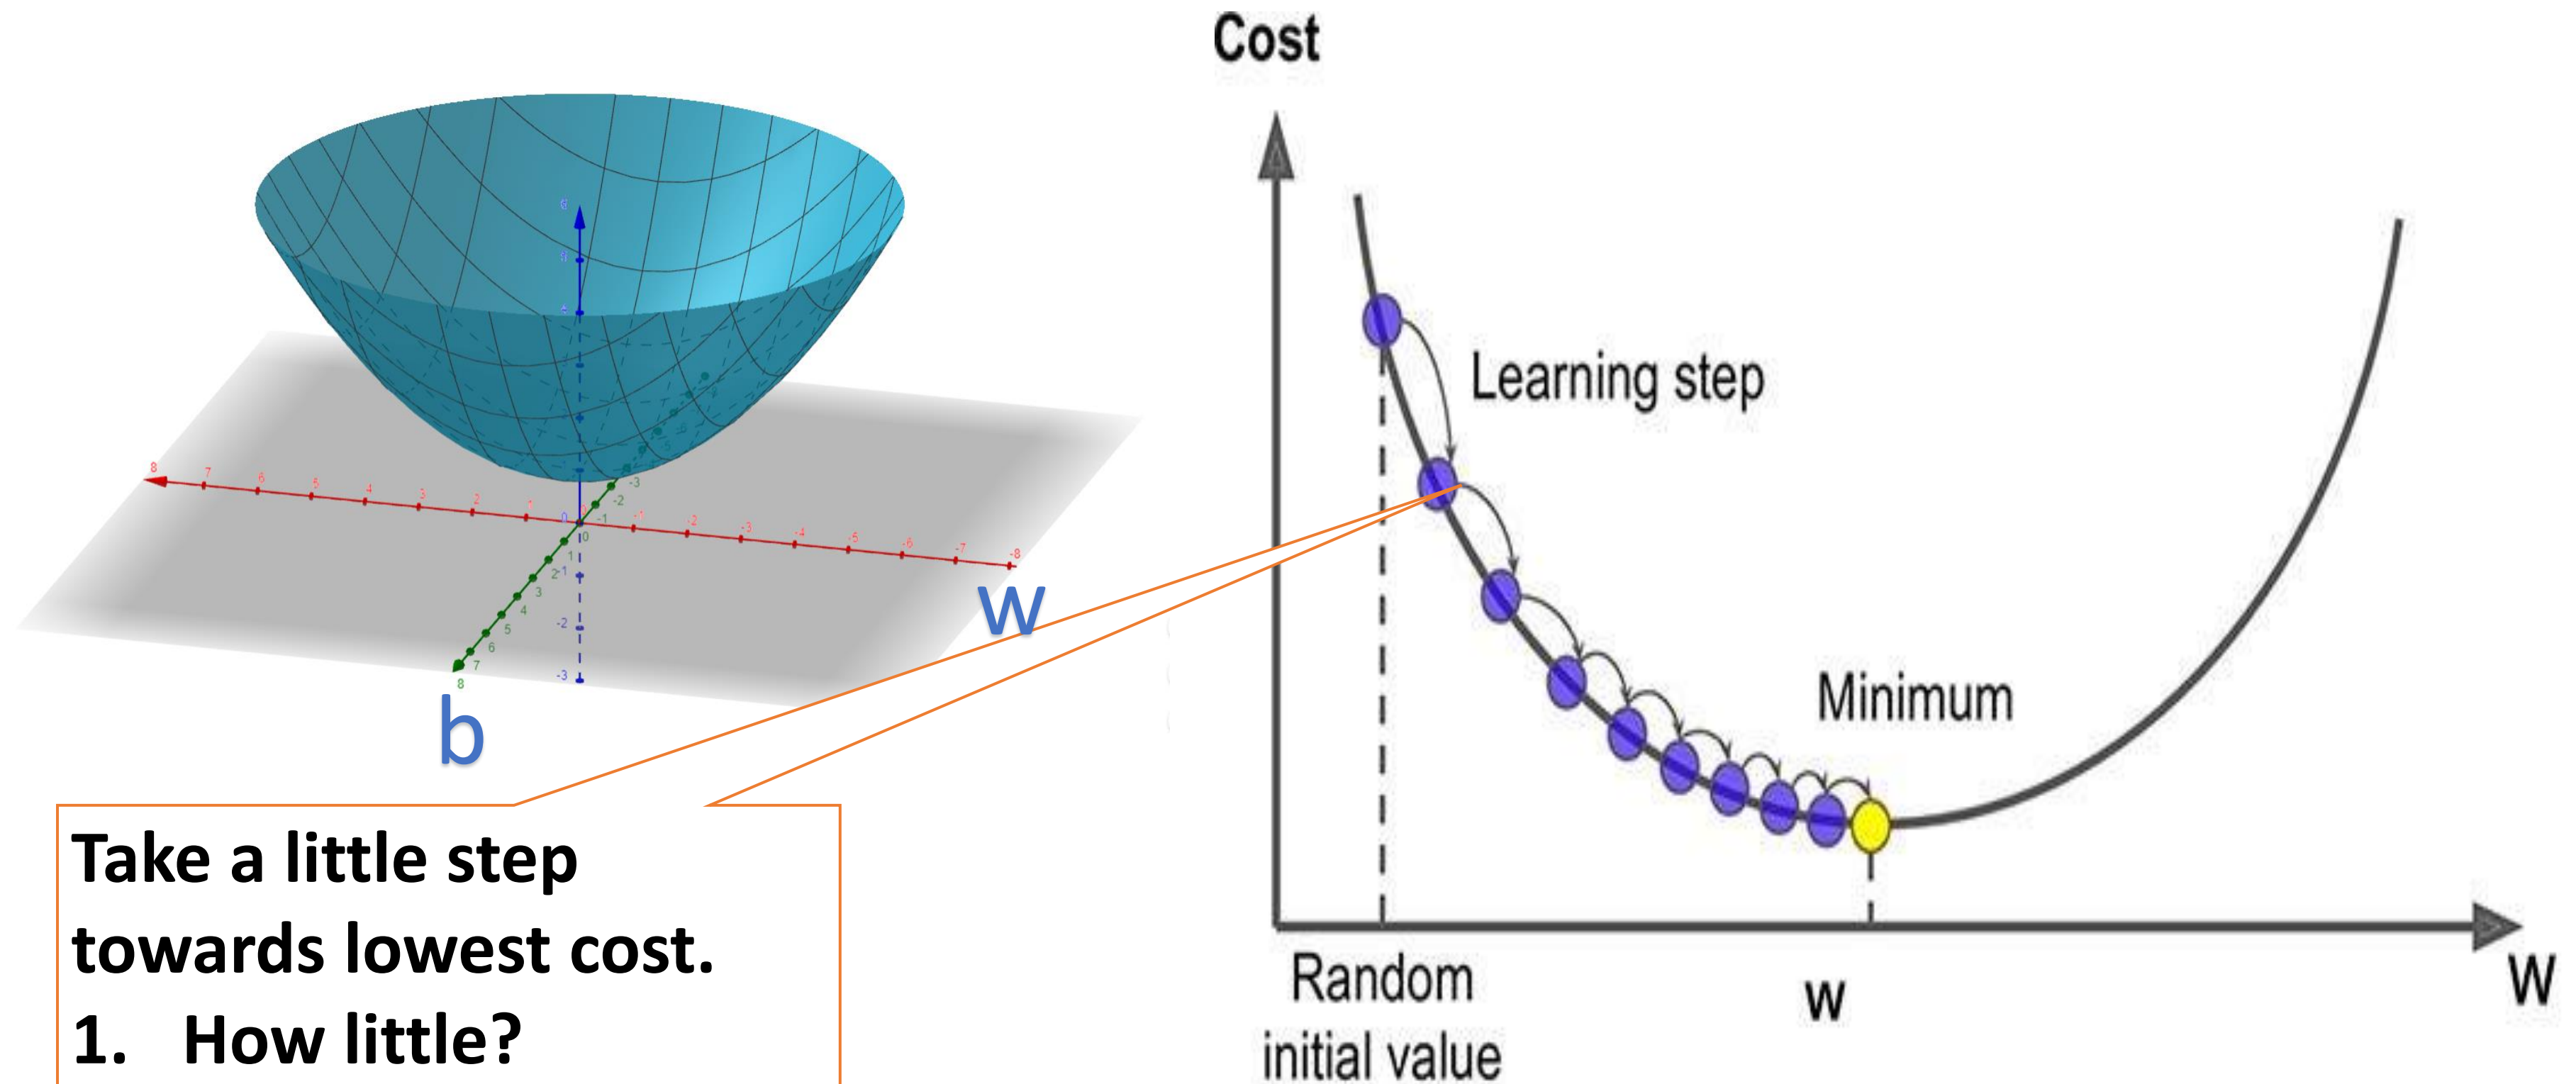
$$\mathcal{J}(w, b) = \frac{1}{n} \sum_{i=1}^n (\hat{y}^{(i)} - y^{(i)})^2$$

$$\mathcal{J}(w, b) = \frac{1}{n} \sum_{i=1}^n (b + wx^{(i)} - y^{(i)})^2$$



# Objective function plot

- <https://www.geogebra.org/calculator/ua52fqtr>



Take a little step  
towards lowest cost.

1. How little?
2. Which direction?

## Step 3. Calculate analytical gradients

$$\mathcal{J}(w, b) = \frac{1}{n} \sum_{i=1}^n (b + wx^{(i)} - y^{(i)})^2$$

$$\mathcal{J}(w, b) = \frac{1}{n} \sum_{i=1}^n (\hat{y}^{(i)} - y^{(i)})^2$$


### • Calculate gradient

$$\frac{\partial \mathcal{J}}{\partial b} = \frac{\partial \mathcal{J}}{\partial \hat{y}^{(i)}} \frac{\partial \hat{y}^{(i)}}{\partial b}$$

$$\frac{\partial \mathcal{J}}{\partial w} = \frac{\partial \mathcal{J}}{\partial \hat{y}^{(i)}} \frac{\partial \hat{y}^{(i)}}{\partial w}$$

$$\frac{\partial \mathcal{J}}{\partial b} = \frac{1}{m} \sum_{i=1}^m 2(b + wx^{(i)} - y^{(i)})$$

$$\frac{\partial \mathcal{J}}{\partial w} = \frac{1}{m} \sum_{i=1}^m 2x^{(i)}(b + wx^{(i)} - y^{(i)})$$





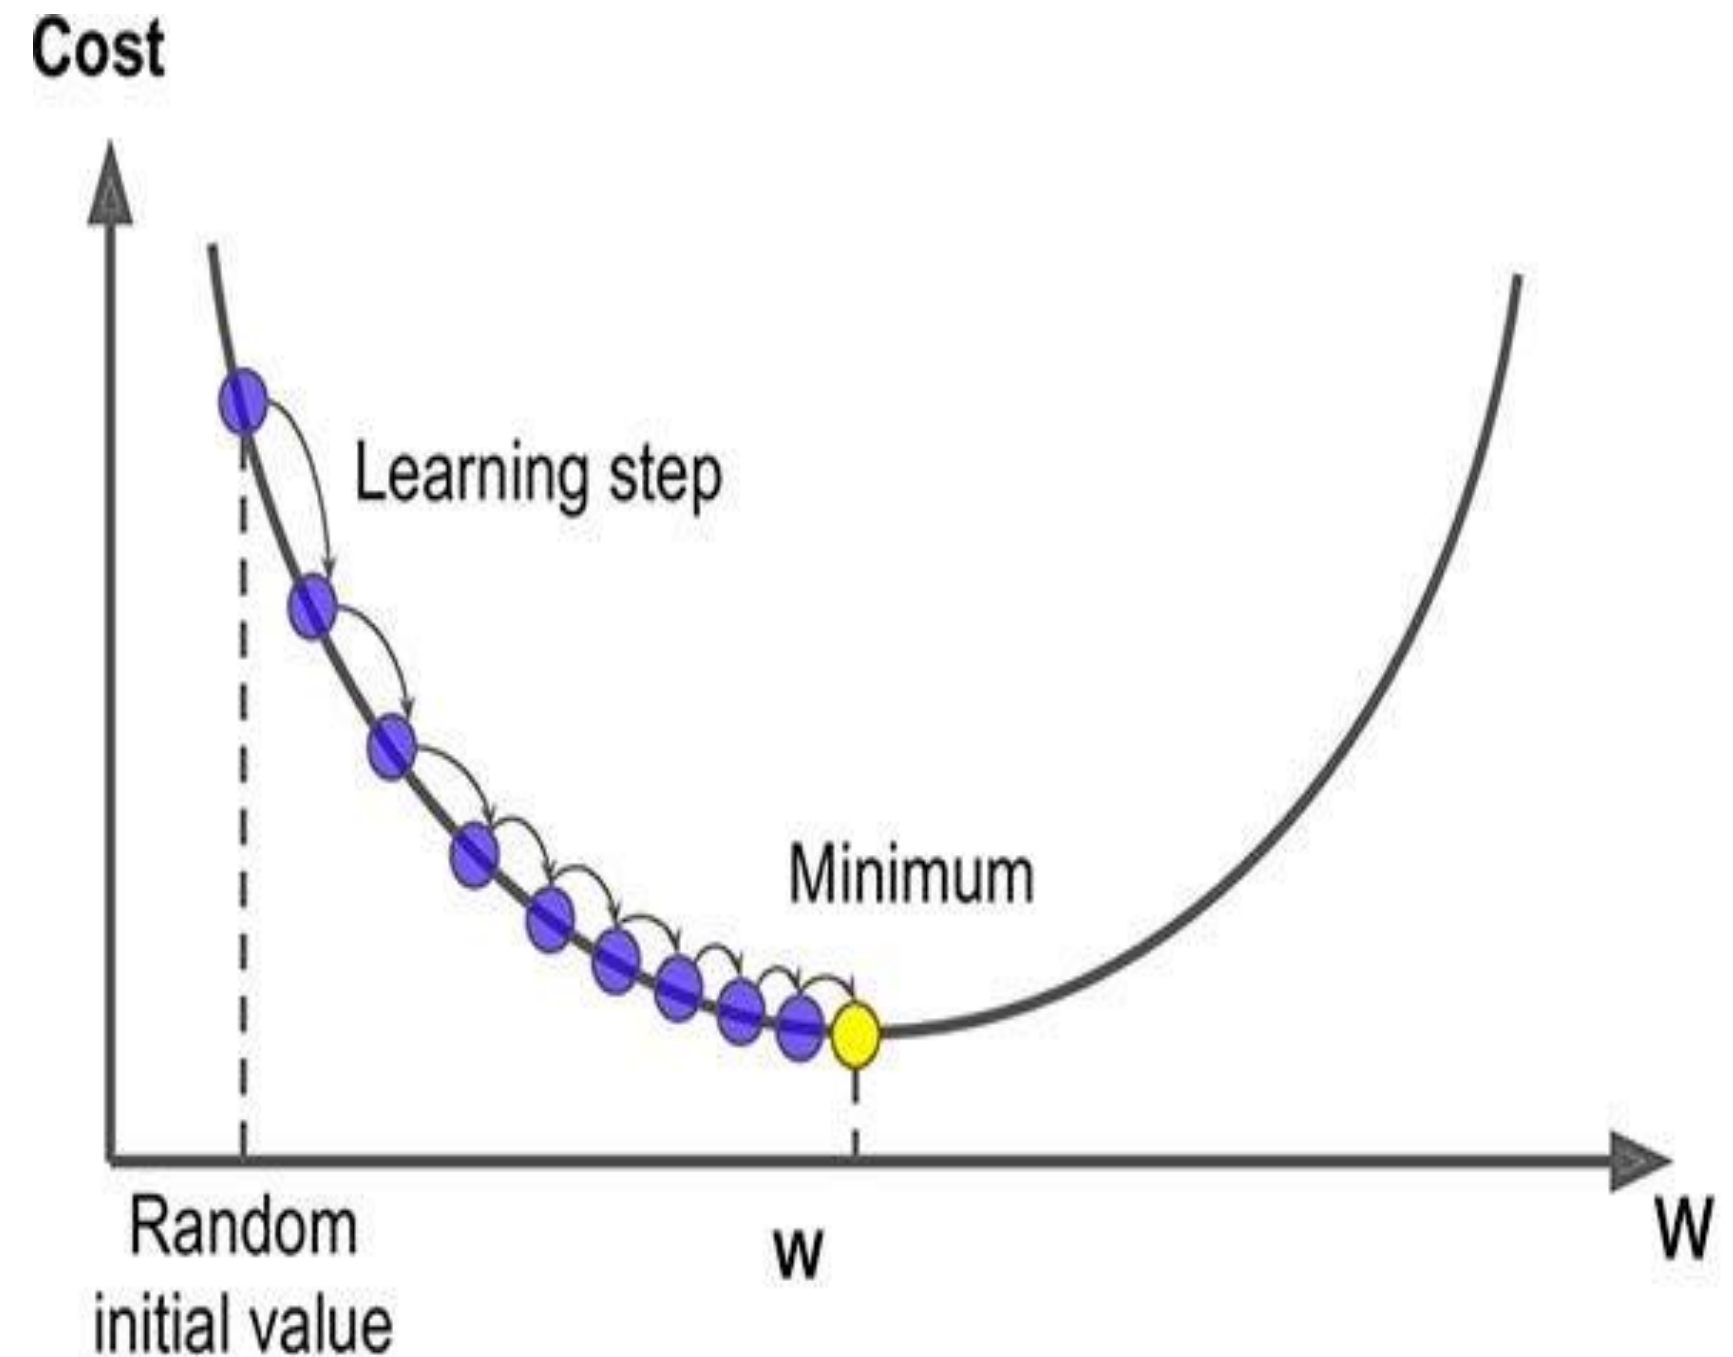
## Step 4. Perform numerical gradient descent

- Also known as backward pass

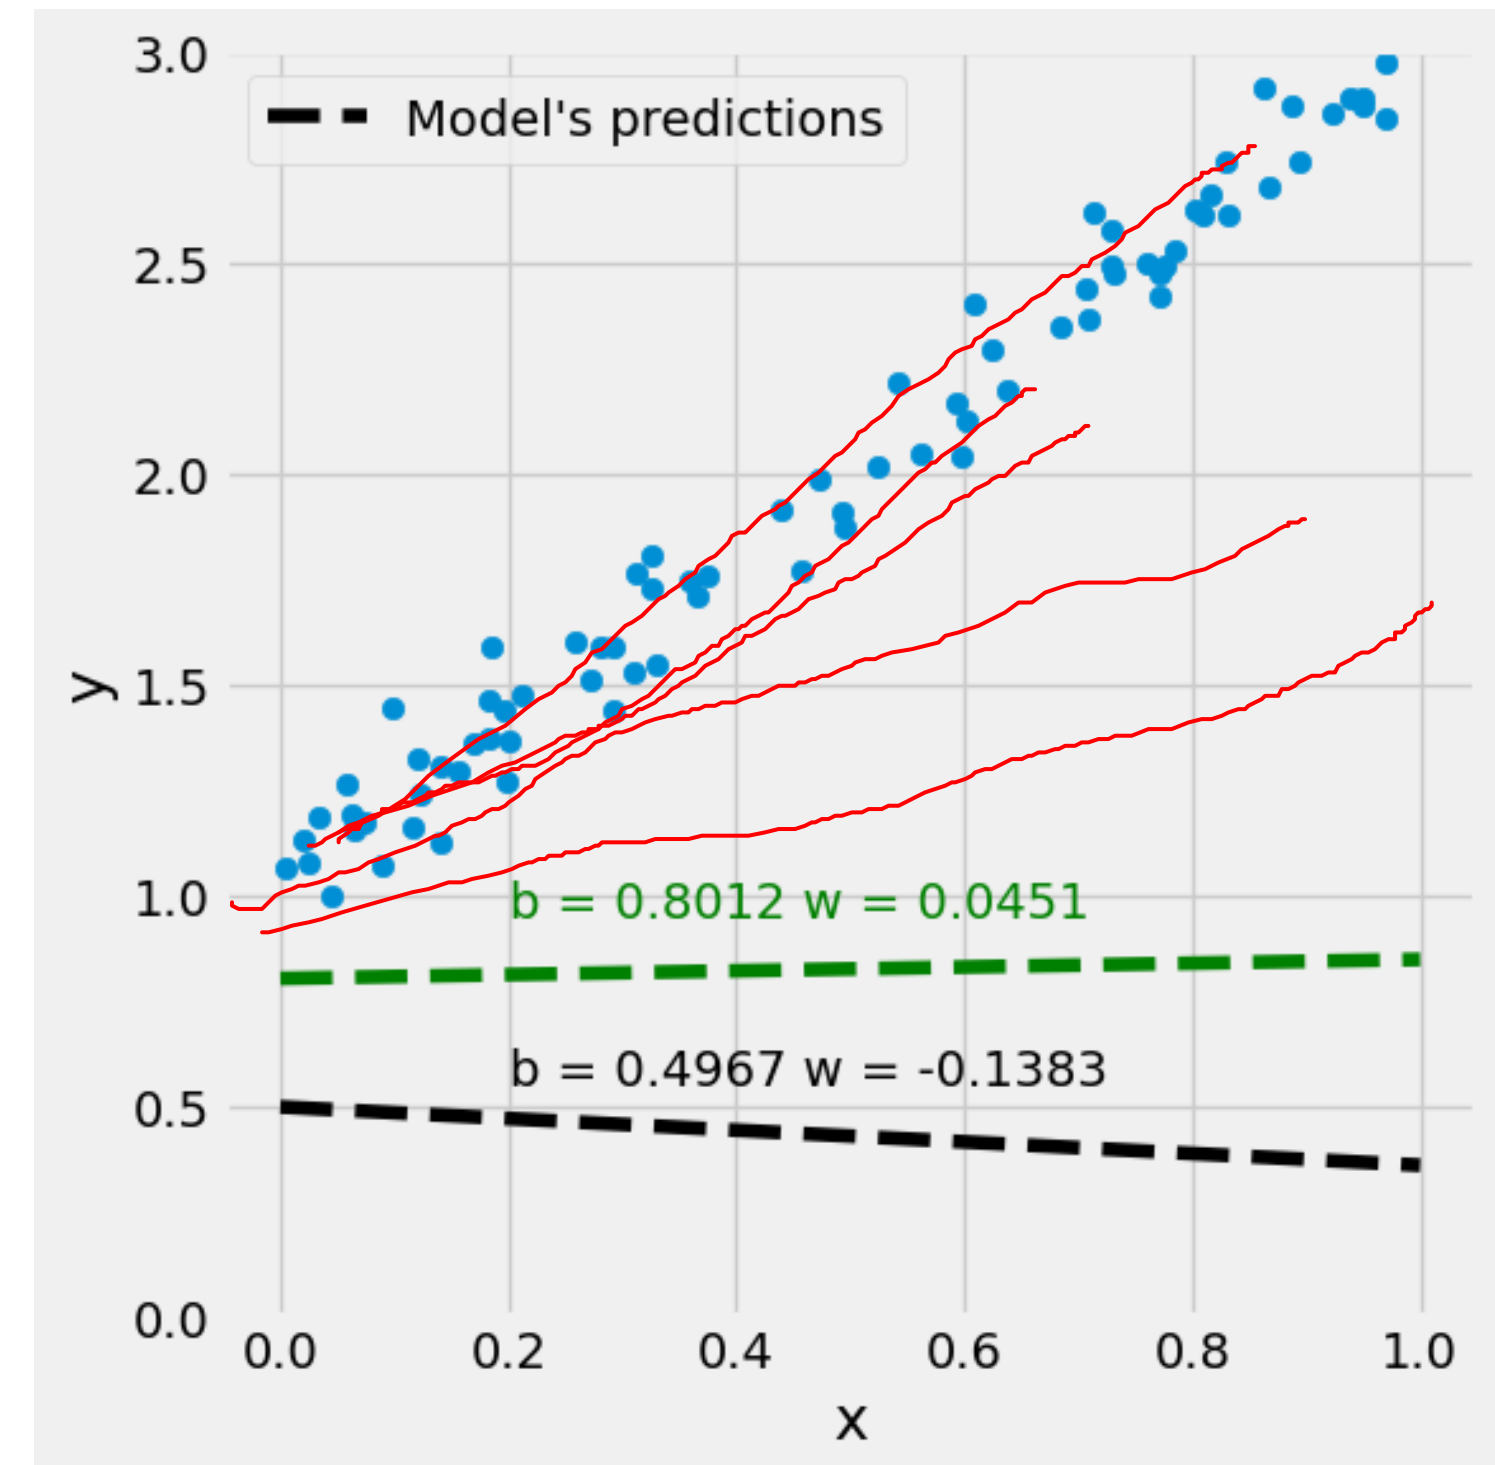
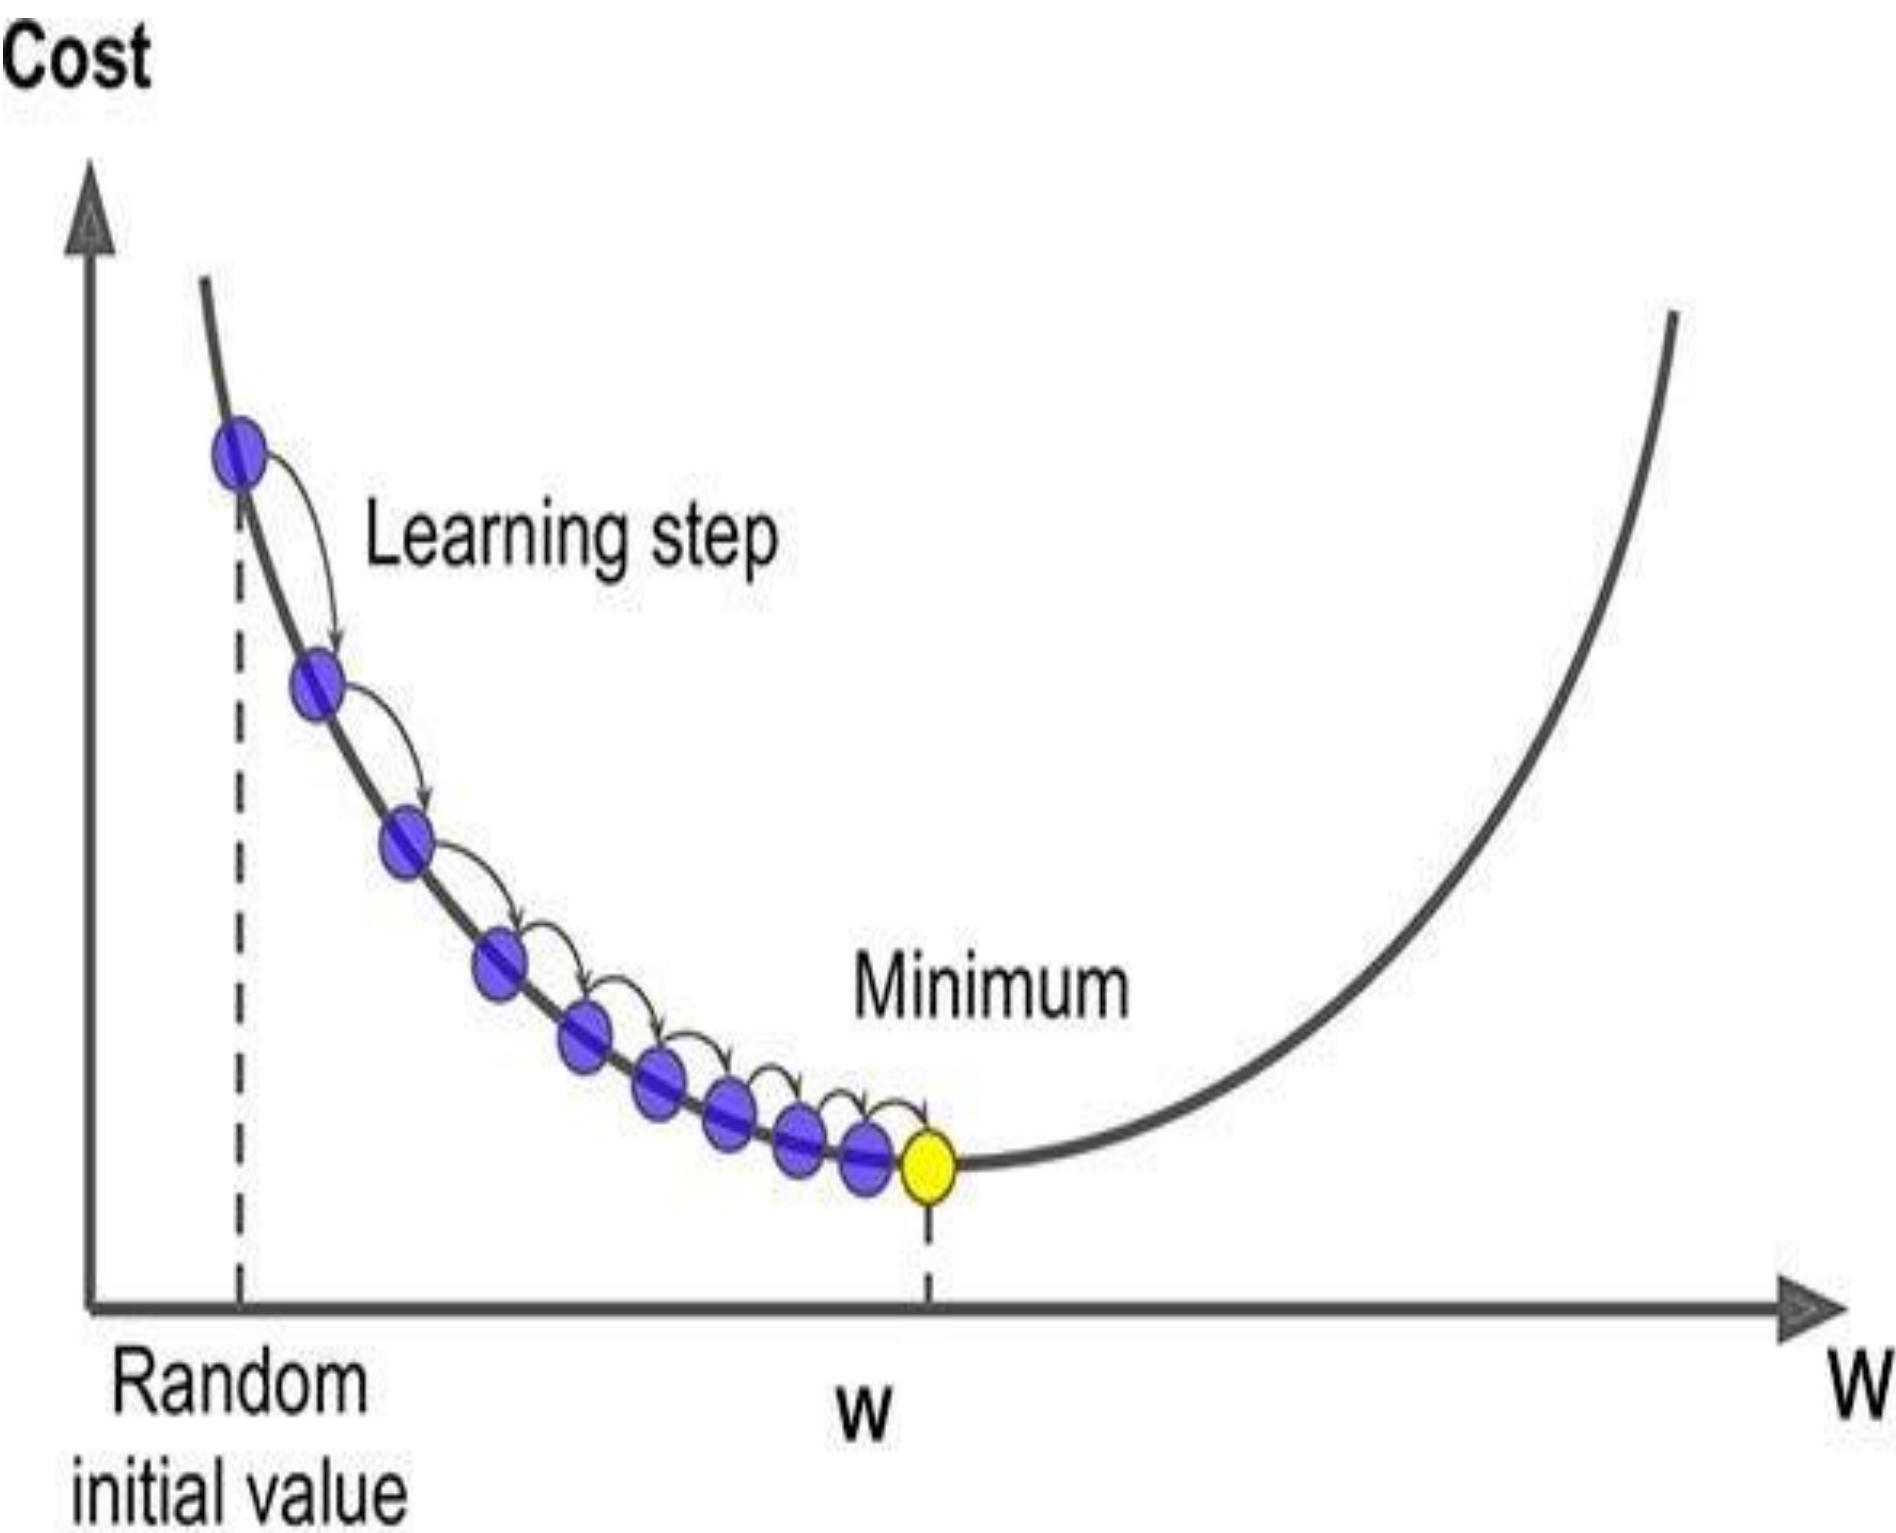
$$b = b - \eta \frac{\partial \mathcal{J}}{\partial b}$$

$$w = w - \eta \frac{\partial \mathcal{J}}{\partial w}$$

- Repeat Step 2, 3, 4



# Change in $w$ and $b$ with gradient descent

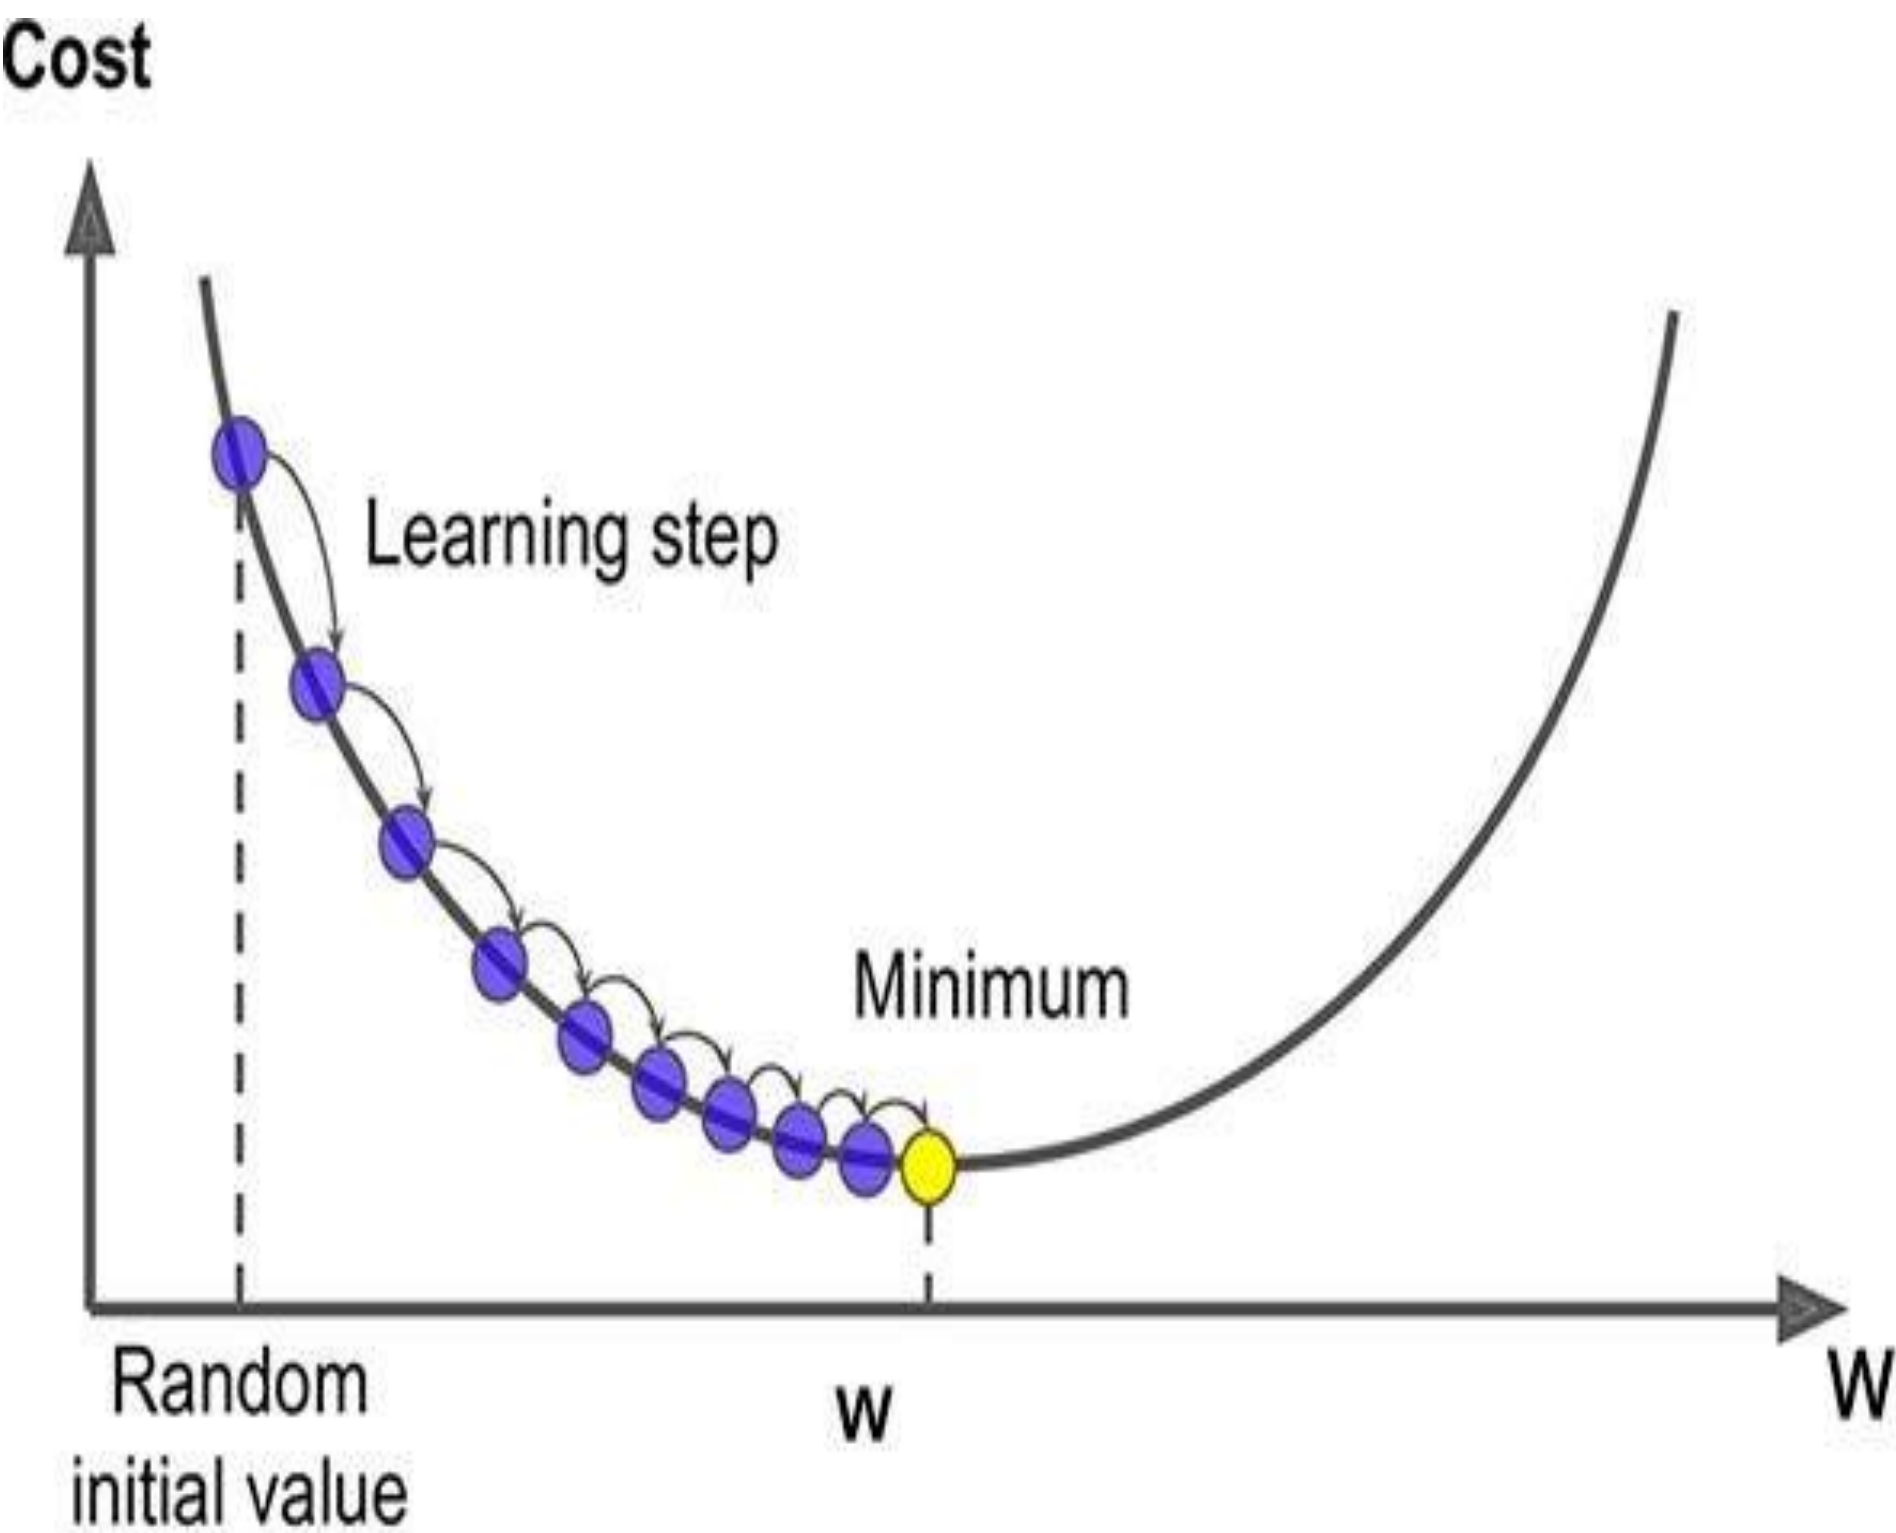


# Review Linear Regression

- Initialization: Select Random  $w$  &  $b$ , choose learning rate  $\eta$
- Loop
  - Calculate new-cost for given  $w$  &  $b$   $\mathcal{J}(w, b) = \frac{1}{n} \sum_{i=1}^n (b + wx^{(i)} - y^{(i)})^2$
  - Break If iter == max or new-cost – old cost < threshold
  - Calculate gradients wrt  $w$  and  $b$   $\frac{\partial \mathcal{J}}{\partial b} = \frac{1}{m} \sum_{i=1}^m 2(b + wx^{(i)} - y^{(i)})$
  - Do gradient descent  $\frac{\partial \mathcal{J}}{\partial w} = \frac{1}{m} \sum_{i=1}^m 2x^{(i)}(b + wx^{(i)} - y^{(i)})$   
$$b = b - \eta \frac{\partial \mathcal{J}}{\partial b}$$
$$w = w - \eta \frac{\partial \mathcal{J}}{\partial w}$$
  - Old Cost = new cost



# Gradient descent summary



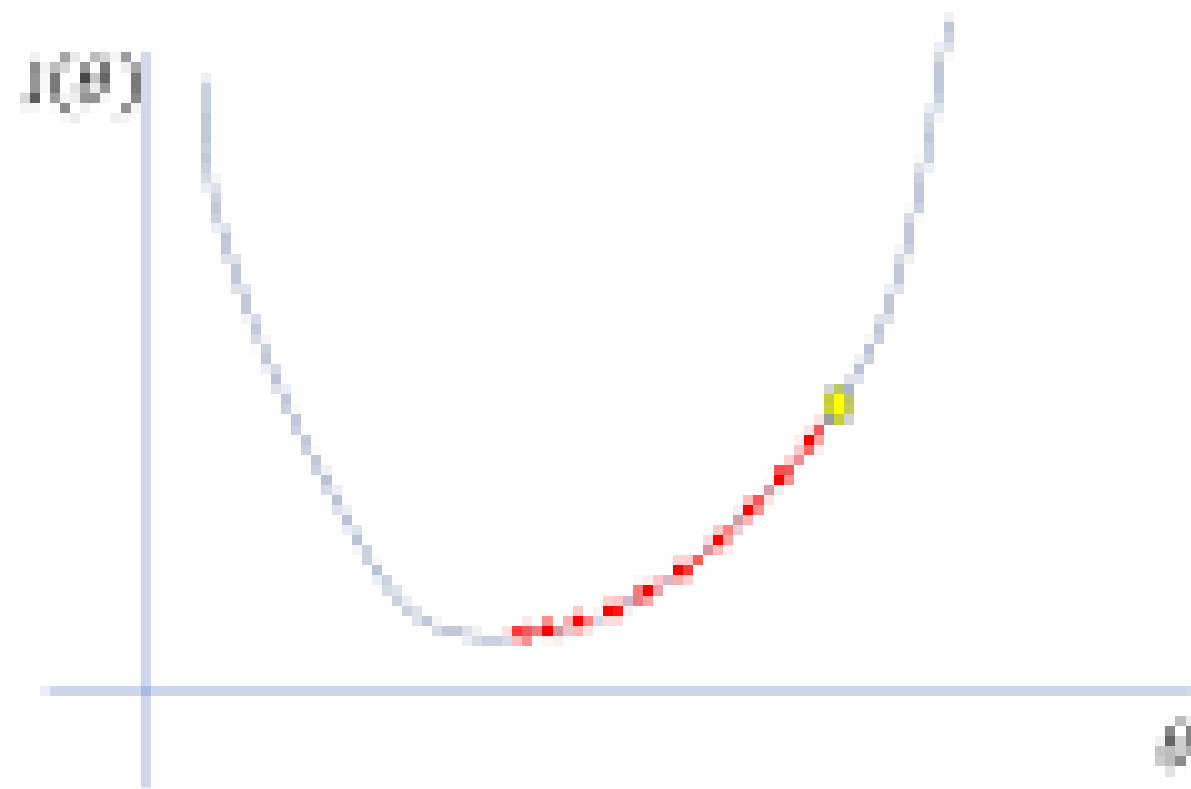
Iteration	w	b	Cost	dJ/dw	dJ/db
0	0.0110	0.0195	2.0443	-1.1077	-1.9538
1000	1.4309	1.2985	0.0178	-0.0407	0.02078
2000	1.7162	1.1527	0.0071	-0.0191	0.00976
3000	1.8502	1.0842	0.0047	-0.0090	0.00459
4000	1.9132	1.0520	0.0042	-0.0042	0.00215
5000	1.9430	1.0369	0.0041	-0.0020	0.00101
6000	1.9567	1.0298	0.0040	-0.0009	0.00047
7000	1.9632	1.0265	0.0040	-0.0004	0.00022
8000	1.9663	1.0249	0.0040	-0.0002	0.00010
9000	1.9677	1.0242	0.0040	-9.637e-05	4.925e-05



# Demo in Jupyter Notebook


# Impact of Learning rate on gradient descent

Too low



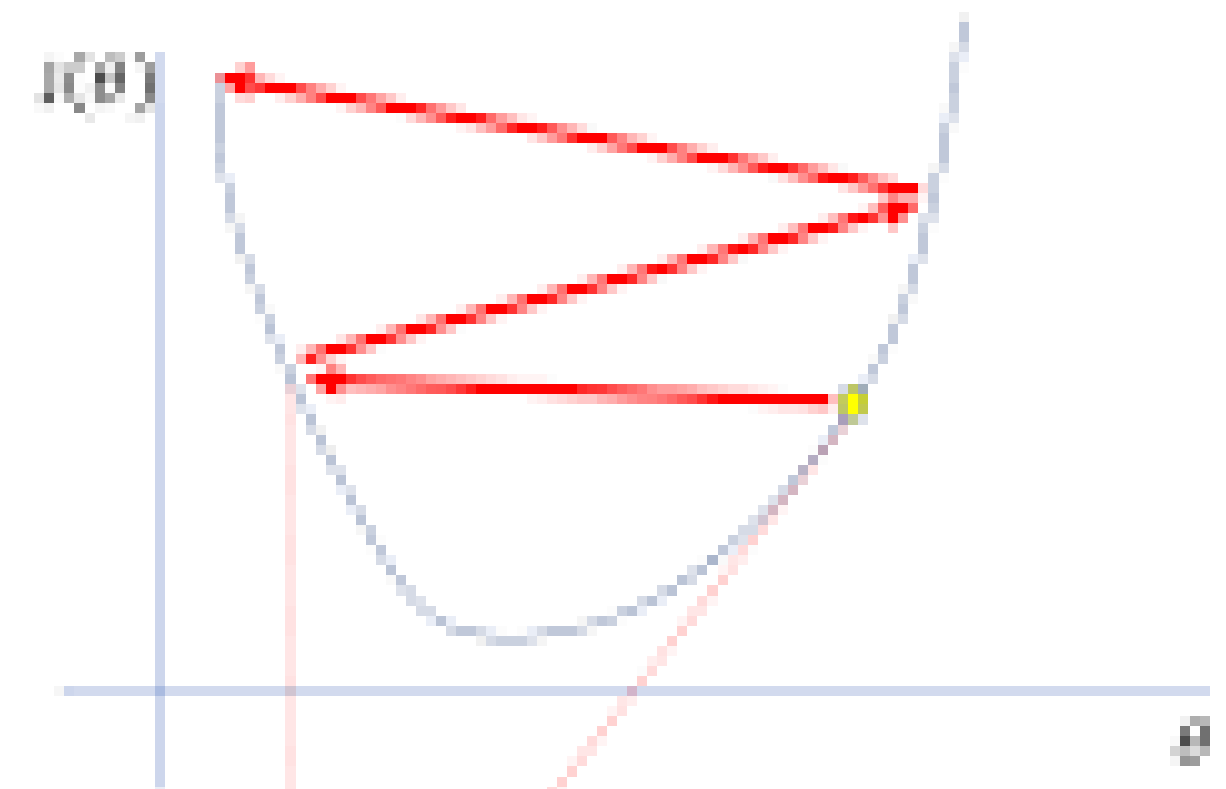
A small learning rate requires many updates before reaching the minimum point

Just right



The optimal learning rate swiftly reaches the minimum point

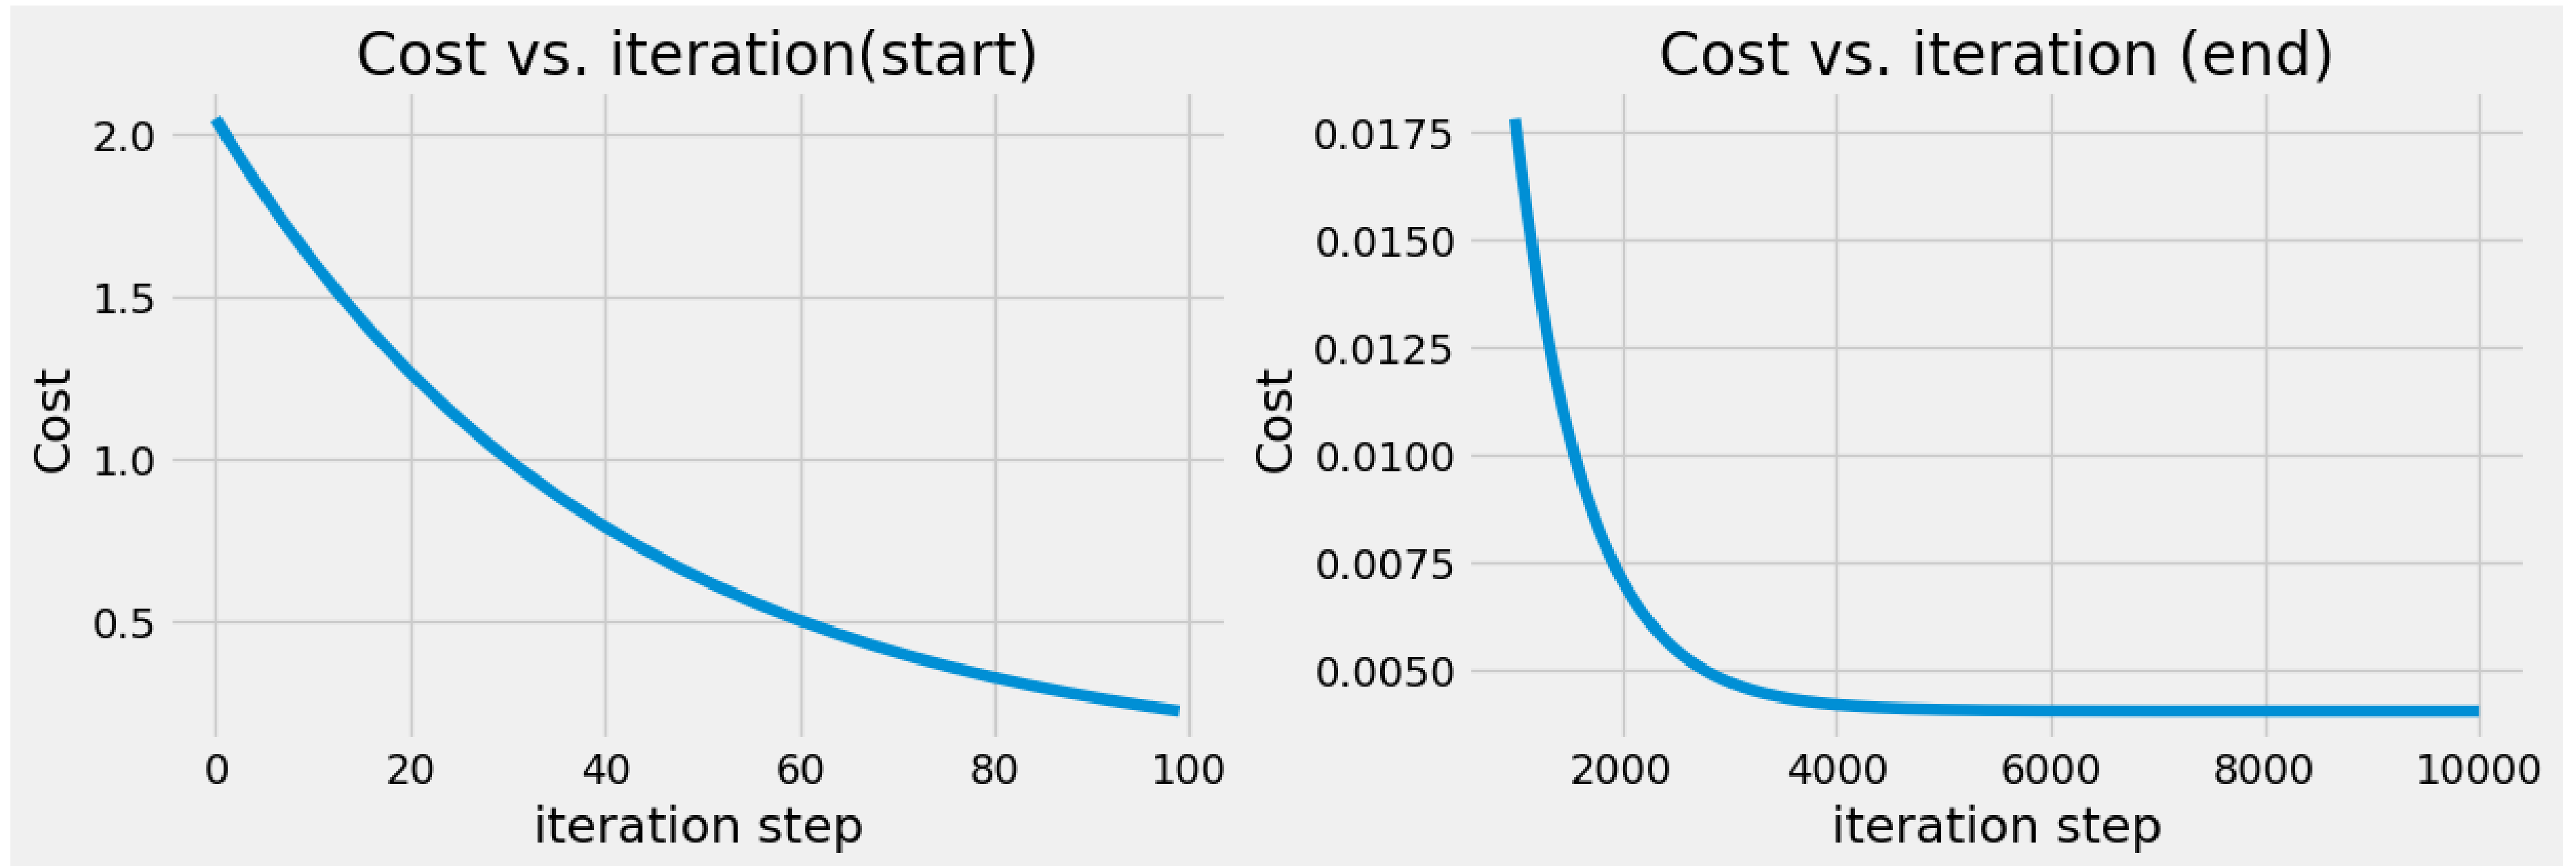
Too high



Too large of a learning rate causes drastic updates which lead to divergent behaviors

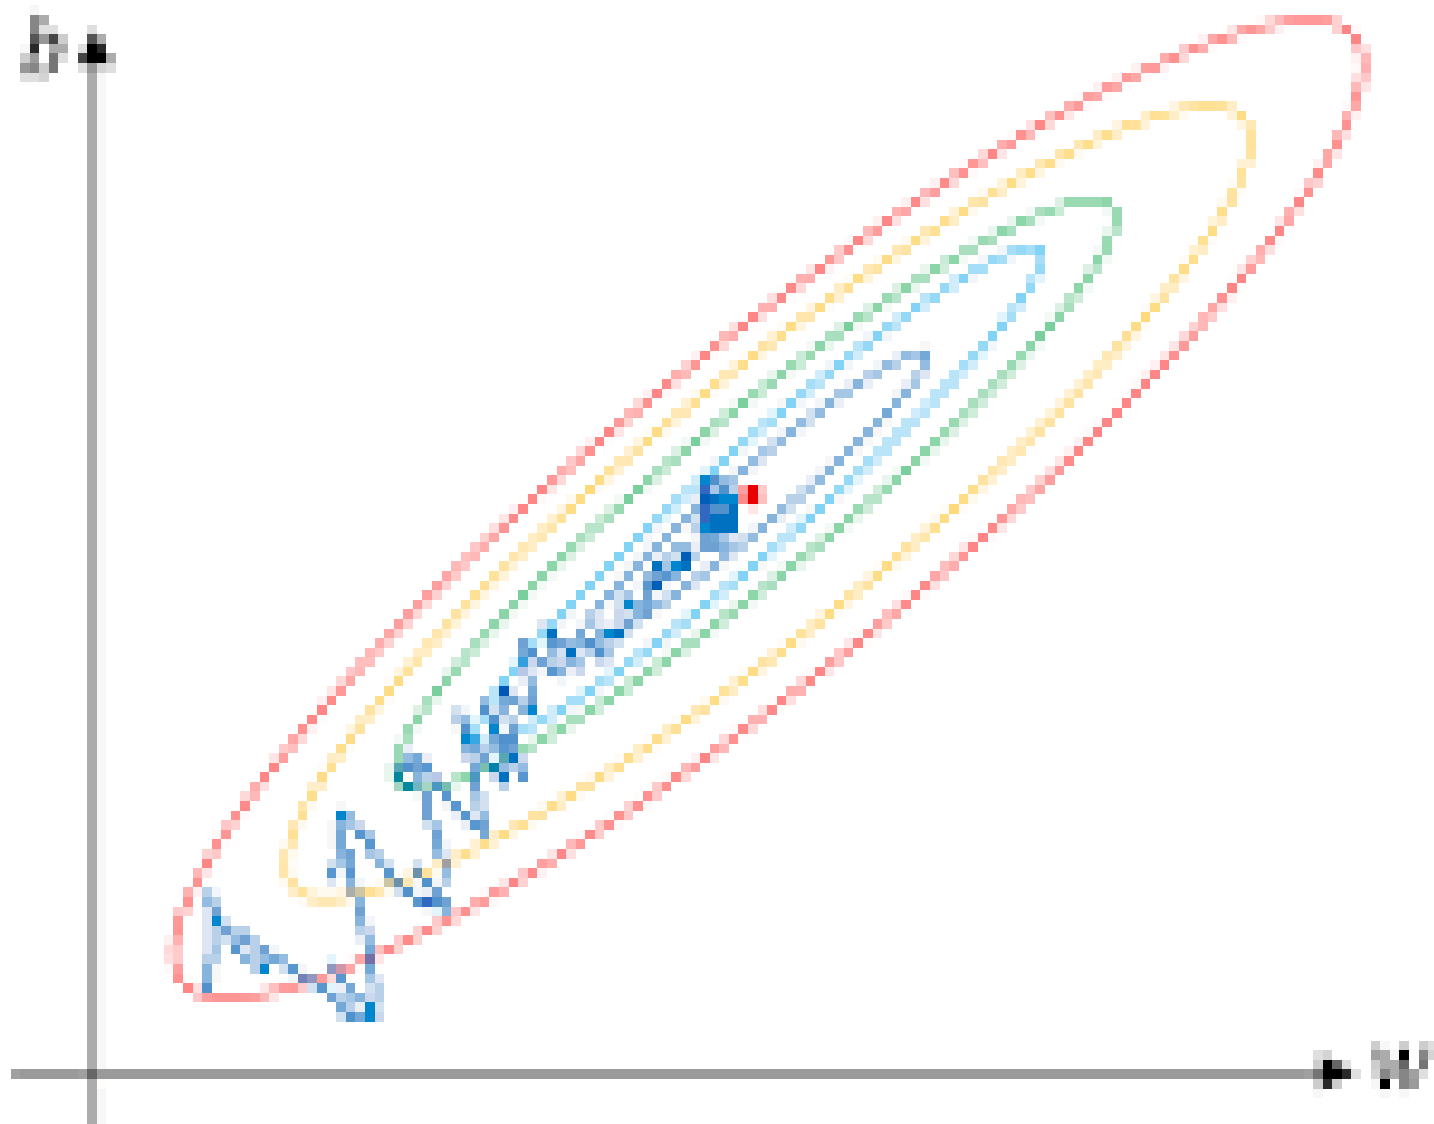


# Impact of Learning rate on gradient descent

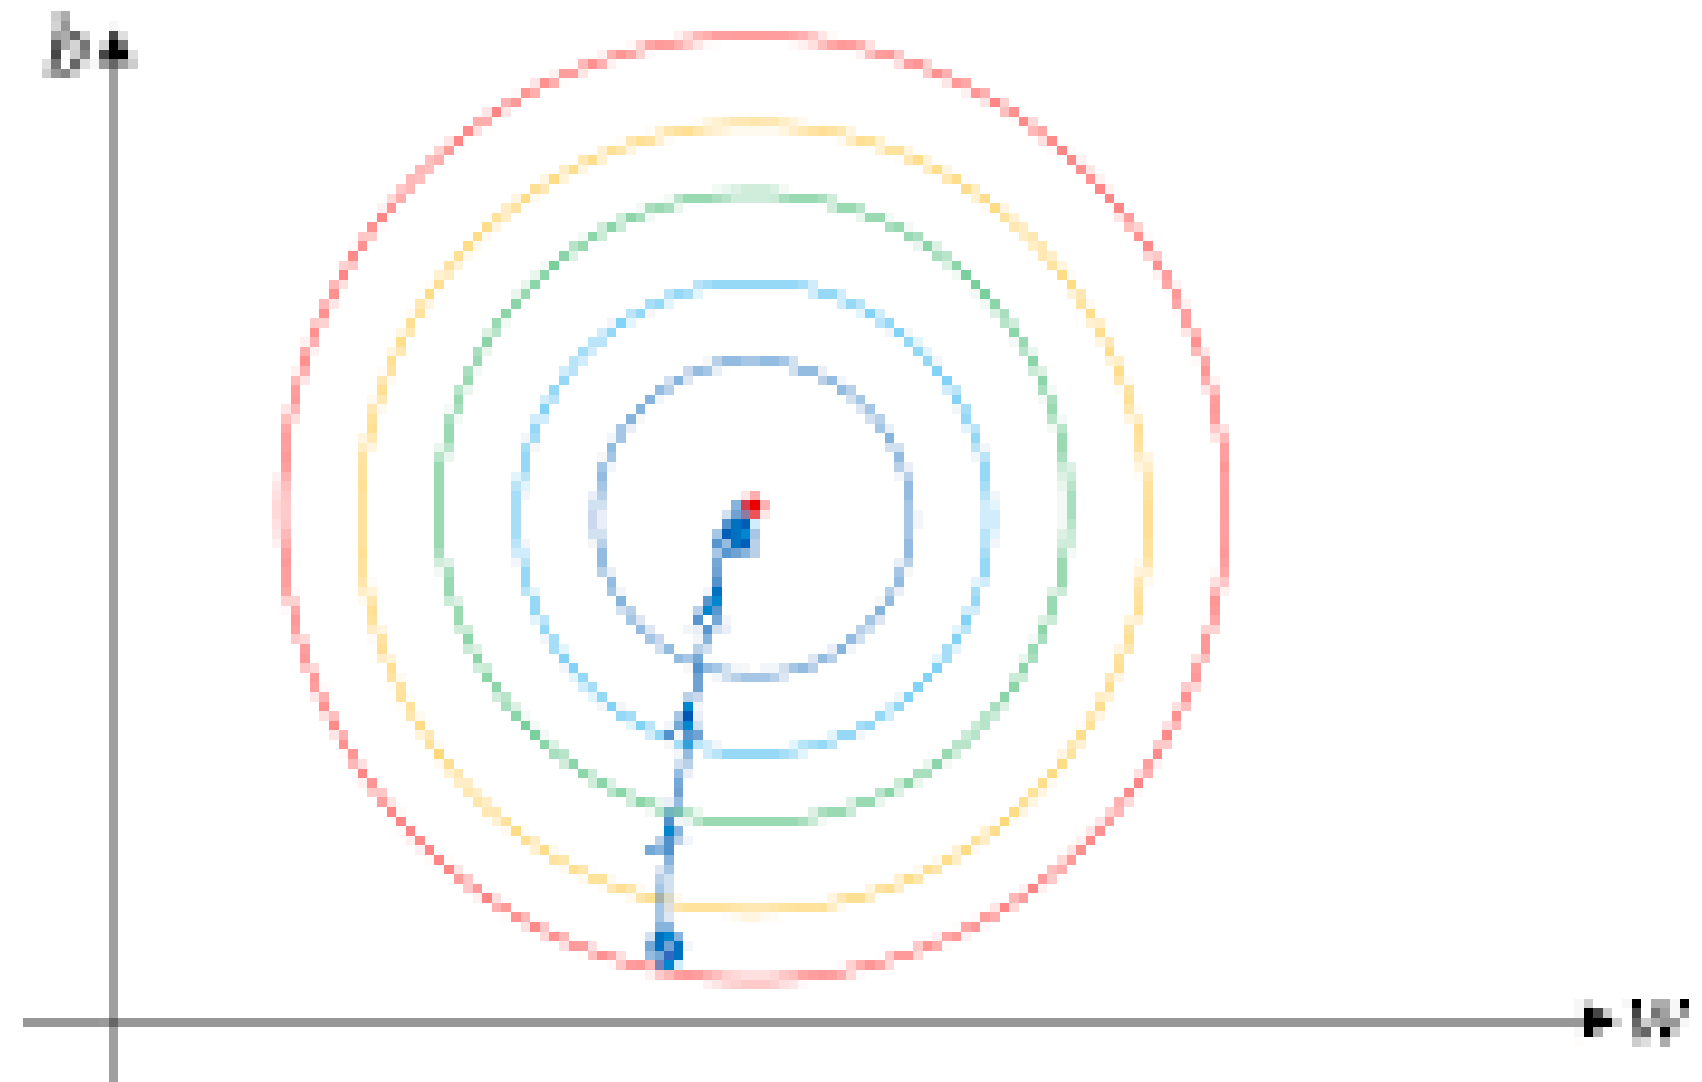


# Importance of Scaling/Normalization

Unnormalized



Normalized



- Faster convergence (Less steps to minima)
- Robust convergence (will not wander away)





# Closed form analytical solution

$$y = w_1 x + w_0 + \epsilon \qquad \hat{y} = w_1 x + w_0$$

$$\mathcal{J}(w_1, w_0) = \frac{1}{n} \sum_{i=1}^n (w_0 + w_1 x^{(i)} - y^{(i)})^2$$

$$\frac{\partial \mathcal{J}}{\partial w_0} = \frac{\partial \mathcal{J}}{\partial \hat{y}^{(i)}} \frac{\partial \hat{y}^{(i)}}{\partial w_0} = \frac{1}{n} \sum_{i=1}^n 2(w_0 + w_1 x^{(i)} - y^{(i)}) = 0$$

$$\frac{\partial \mathcal{J}}{\partial w_1} = \frac{\partial \mathcal{J}}{\partial \hat{y}^{(i)}} \frac{\partial \hat{y}^{(i)}}{\partial w_1} = \frac{1}{n} \sum_{i=1}^n 2(w_0 + w_1 x^{(i)} - y^{(i)}) x^{(i)} = 0$$

# Closed form analytical solution

$$w_1 = \frac{n \sum_{i=1}^m x^{(i)} y^{(i)} - \sum_{i=1}^m x^{(i)} \sum_{i=1}^m y^{(i)}}{n \sum_{i=1}^m x^{(i)2} - \left(\sum_{i=1}^m x^{(i)}\right)^2}$$

$$w_0 = \frac{\sum_{i=1}^m y^{(i)} - w_1 \sum_{i=1}^m x^{(i)}}{n}$$

- Formula starts getting complicated with inter-dependencies
- Needs to load all data at once
  - What happens when there are million+ records?



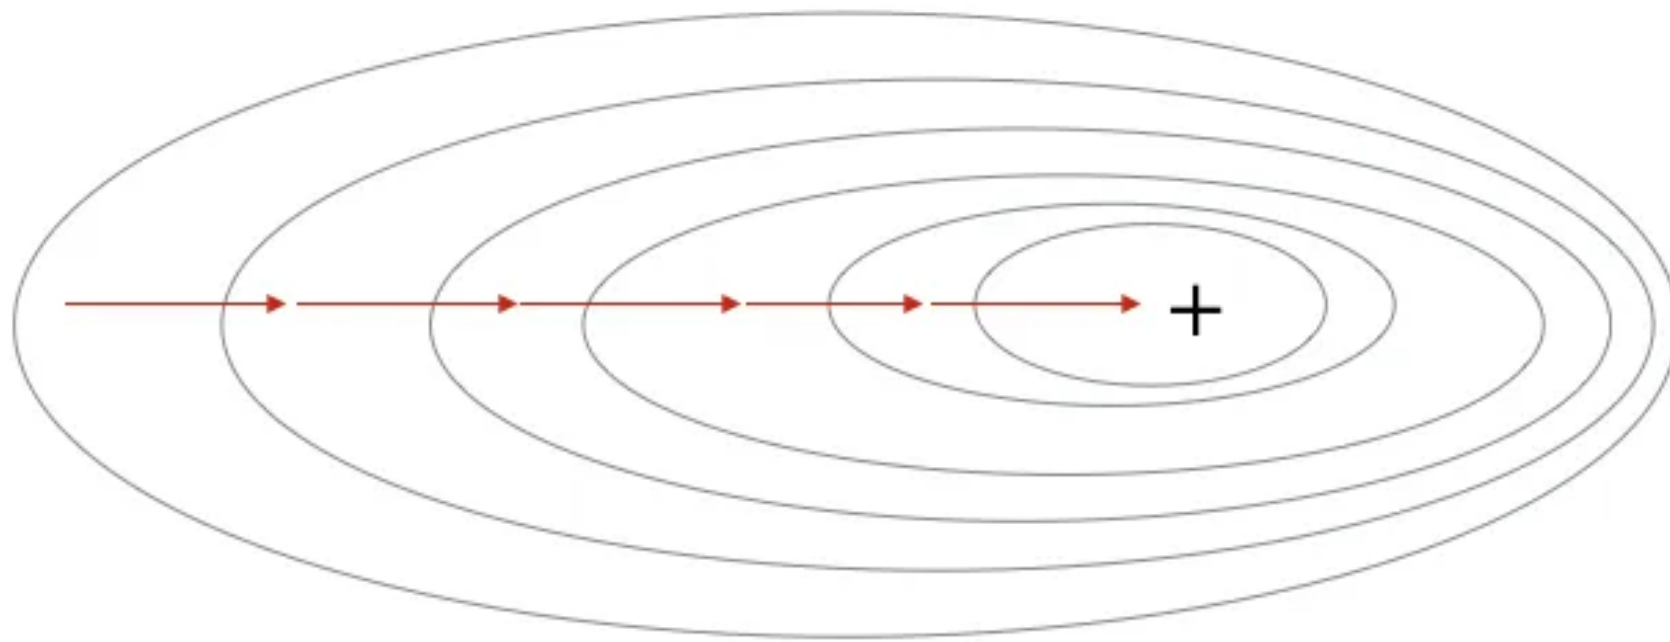


# Gradient descent - Batch, mini batch & stochastic

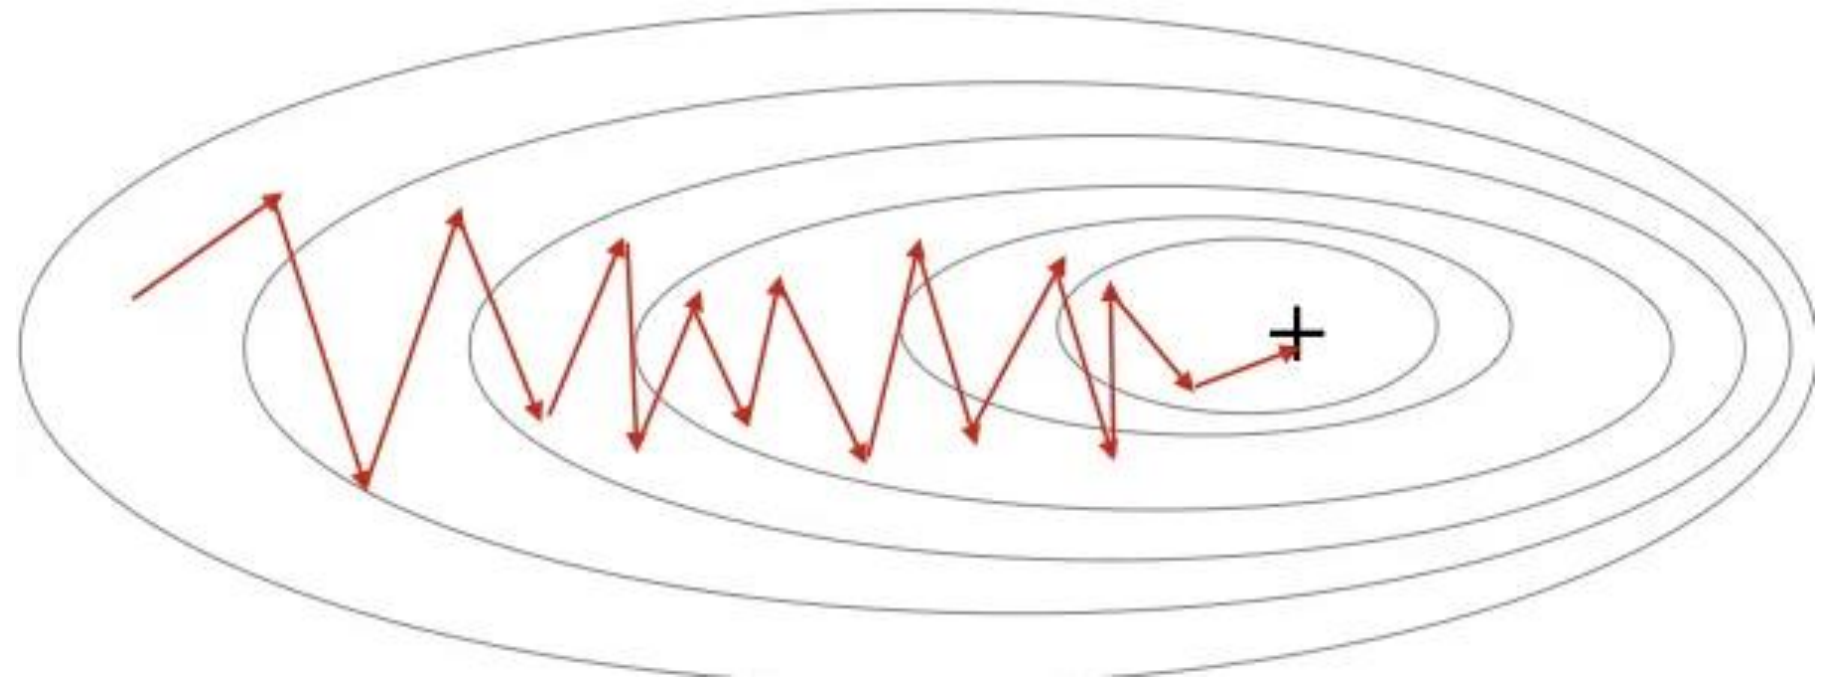
- Batch gradient descent
  - Entire dataset used, Offline training
  - Good for small data set
- Stochastic gradient descent
  - One record used at a time, Online training
  - Good for streaming data
- Mini-batch gradient descent
  - Large dataset is cut into chunks, Calc J on each chunk
  - Iterate over entire dataset many times progressively reducing cost

# Gradient descent comparison

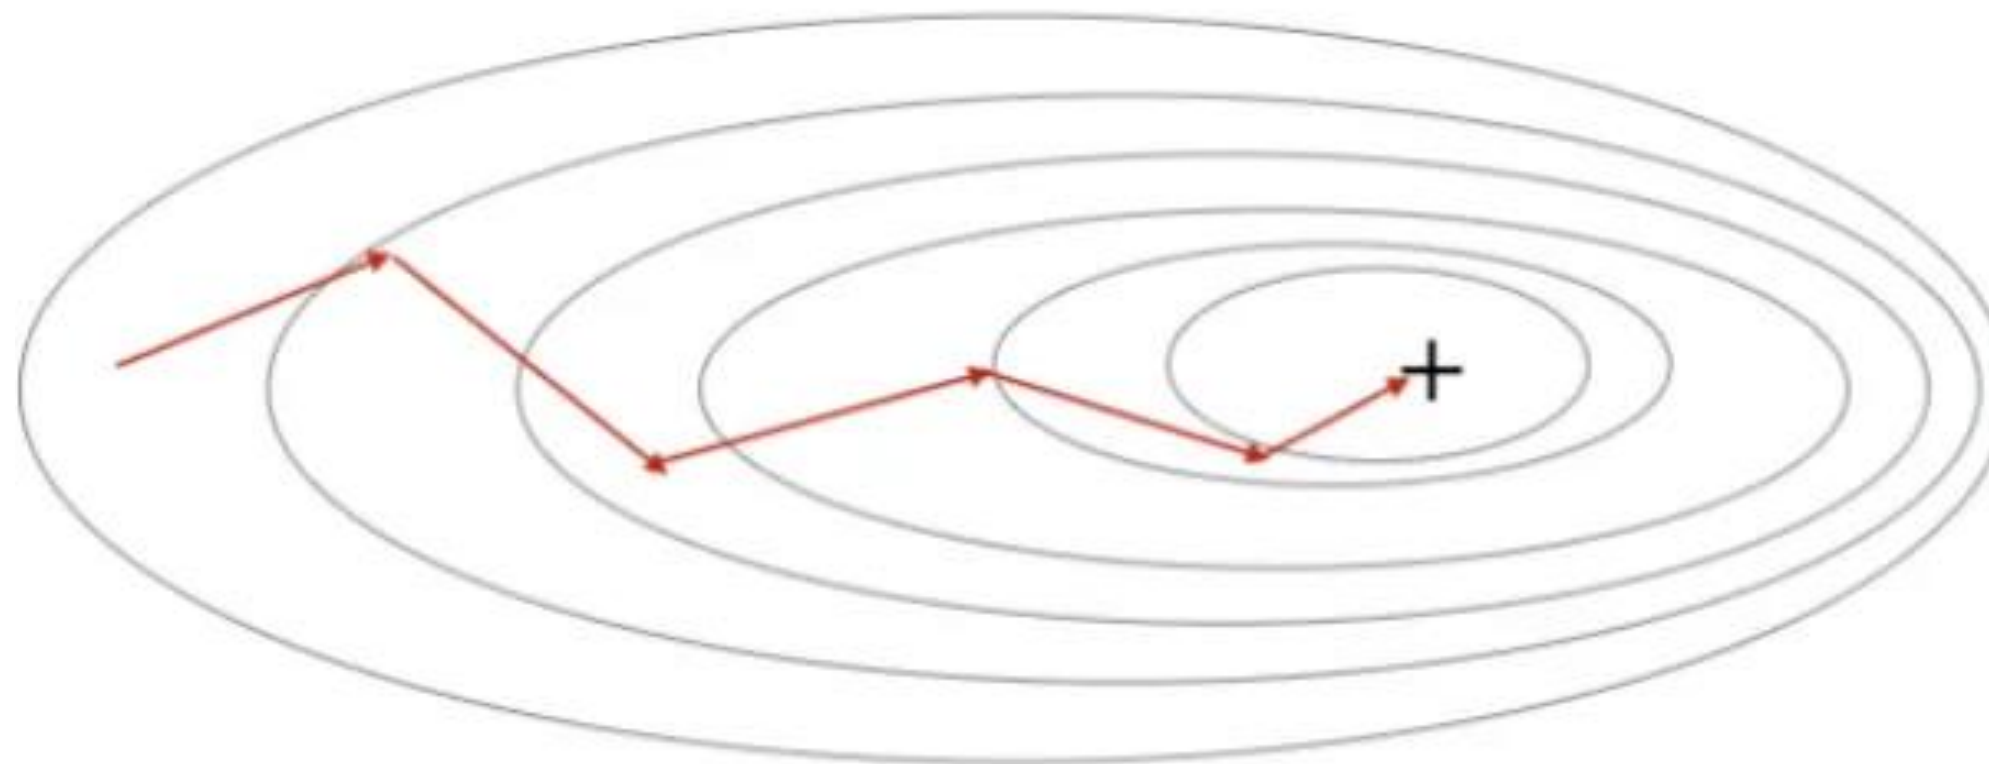
## Batch Gradient Descent



## Stochastic Gradient Descent

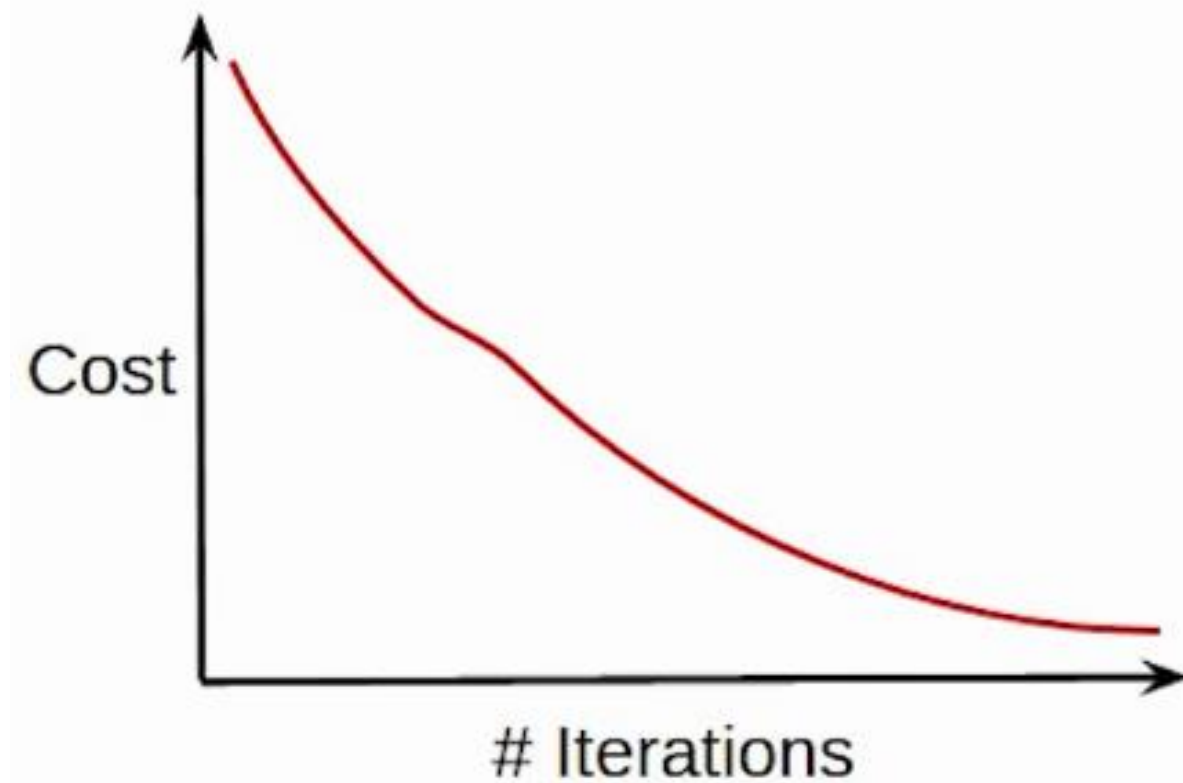


## Mini-batch Gradient Descent

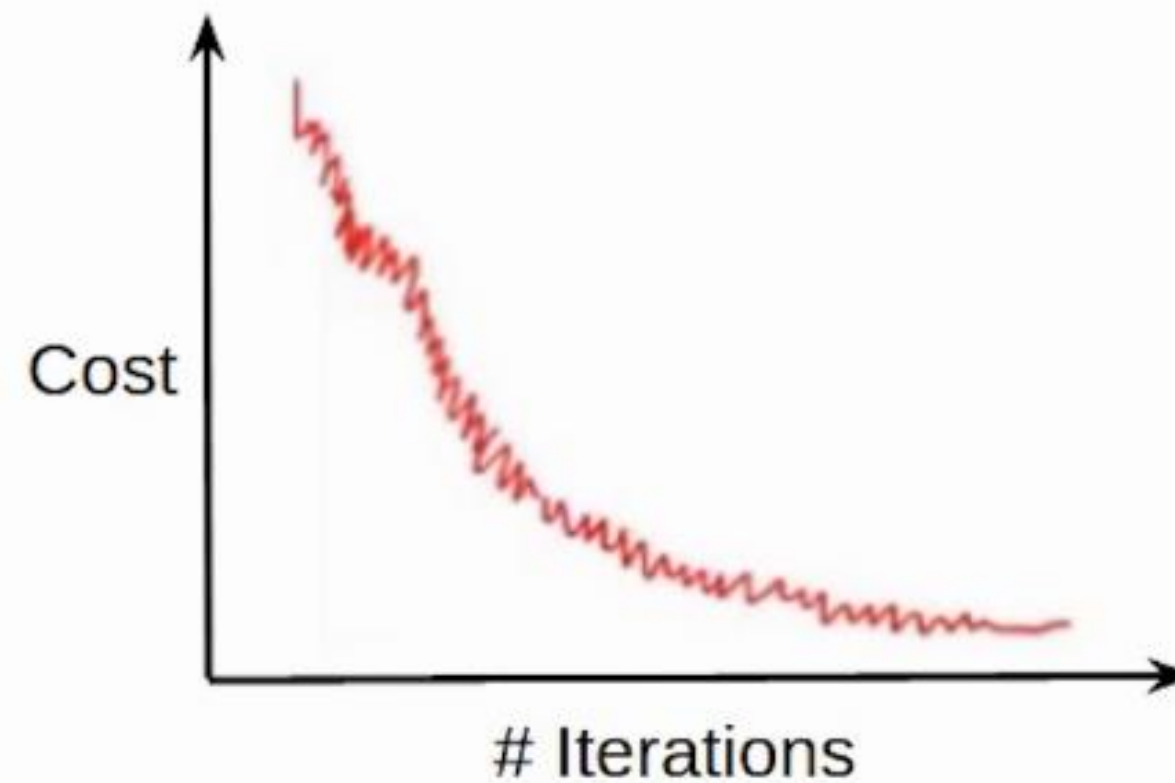


# Cost function comparison

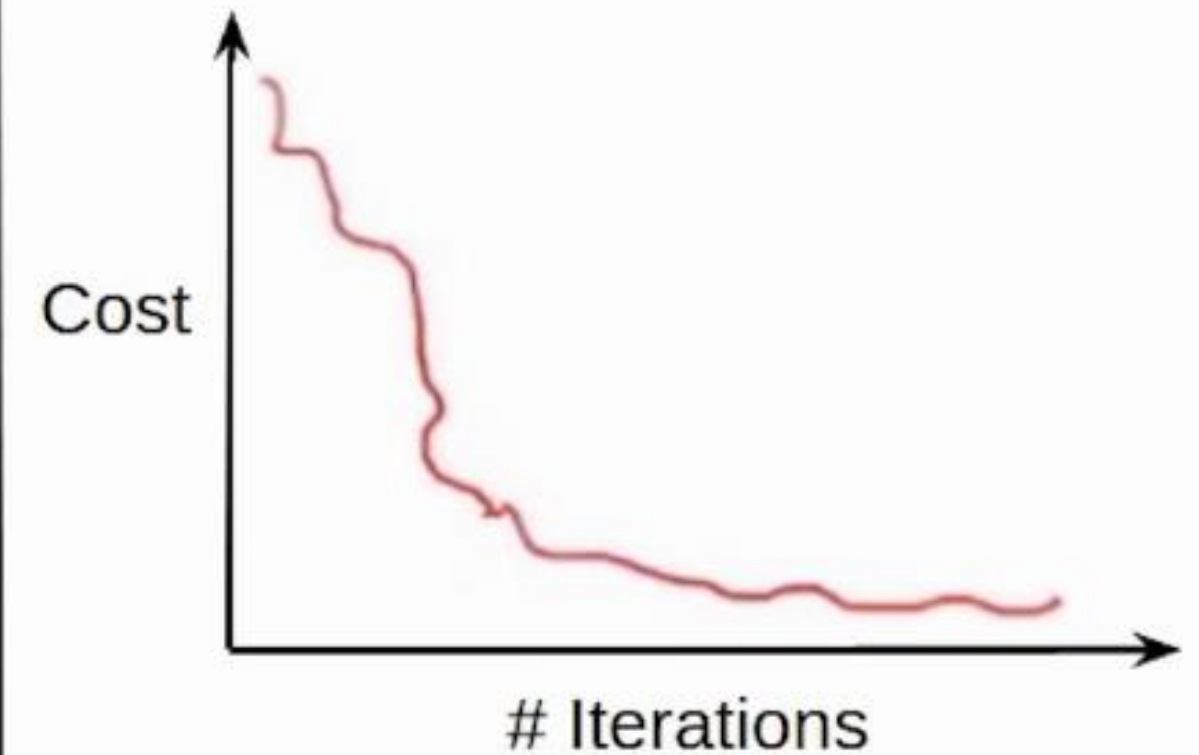
- Cost function reduces smoothly



- Lot of variations in cost function



- Smoother cost function as compared to SGD





# Gradient descent comparison

## Batch Gradient Descent

- Entire dataset for updation
- Cost function reduces smoothly
- Computation cost is very high

## Stochastic Gradient Descent (SGD)

- Single observation for updation
- Lot of variations in cost function
- Computation time is more

## Mini-Batch Gradient Descent

- Subset of data for updation
- Smoother cost function as compared to SGD
- Computation time is lesser than SGD
- Computation cost is lesser than Batch Gradient Descent






# Coding Linear Regression

# Linear Regression Dataset

- Advertising.csv

TV	Radio	Newspaper	Sales
230.1	37.8	69.2	22.1
44.5	39.3	45.1	10.4
17.2	45.9	69.3	9.3
151.5	41.3	58.5	18.5
180.8	10.8	58.4	12.9



# Coding with statsmodels

```
import statsmodels.api as sm
lm = sm.OLS(y, X)
model = lm.fit()
```

```
model.summary()
```

OLS Regression Results						
Dep. Variable:	sales		R-squared (uncentered):		0.982	
Model:	OLS		Adj. R-squared (uncentered):		0.982	
Method:	Least Squares		F-statistic:		3566	
Date:	Sun, 28 Mar 2021		Prob (F-statistic):		2.43e-171	
Time:	13:42:33		Log-Likelihood:		-423.54	
No. Observations:	200		AIC:		853.1	
Df Residuals:	197		BIC:		863.0	
Df Model:	3					
Covariance Type:	nonrobust					
	coef	std err	t	P> t	[0.025	0.975]
TV	0.0538	0.001	40.507	0.000	0.051	0.056
radio	0.2222	0.009	23.595	0.000	0.204	0.241
newspaper	0.0168	0.007	2.517	0.013	0.004	0.030
Omnibus:	5.982	Durbin-Watson:	2.038			
Prob(Omnibus):	0.050	Jarque-Bera (JB):	7.039			
Skew:	-0.232	Prob(JB):	0.0296			
Kurtosis:	3.794	Cond. No.	12.6			



# Coding with sklearn

```
from sklearn.linear_model import LinearRegression  
lm = LinearRegression()  
model = lm.fit(X,y)
```

```
model.intercept_
```

```
array([2.93888937])
```

```
model.coef_
```

```
array([[ 0.04576465,  0.18853002, -0.00103749]])
```

```
model.predict(new_data)
```

```
array([[6.15044172]])
```

# Different types of gradient descent in sklearn

- Batch gradient descent
  - `sklearn.linear_model.LinearRegression`
- Stochastic gradient descent
  - `sklearn.linear_model.SGDRegressor`
- Mini-batch gradient descent
  - `sklearn.linear_model.SGDRegressor`
  - `partial_fit()`
  - Pass each mini batch into `partial_fit()`
  - Cannot use `partial_fit()` in Pipeline!

# Evaluation metrics for Regression


```
model.score(X_train, y_train)
```

```
:  
0.910413637900632
```

```
model.score(X_test, y_test)
```

```
:  
0.8495077592917368
```

- What does score mean in Regression?






# Evaluation metrics for Regression

- Mean Squared Error (MSE) 
$$= \frac{1}{n} \sum_{i=1}^n (\hat{y}^{(i)} - y^{(i)})^2$$
- Root Mean Squared Error (RMSE) 
$$= \sqrt{\frac{1}{n} \sum_{i=1}^n (\hat{y}^{(i)} - y^{(i)})^2}$$
- Mean Absolute Error (MAE) 
$$= \frac{1}{n} \sum_{i=1}^n |\hat{y}^{(i)} - y^{(i)}|$$
- R-Squared 
$$R^2 = 1 - \frac{SS_{reg}}{SS_{avg}} = 1 - \frac{\sum_{i=1}^n (\hat{y}^{(i)} - y^{(i)})^2}{\sum_{i=1}^n (\hat{y}^{(i)} - \bar{y})^2}$$
- Adjusted R-Squared 
$$R_{adj}^2 = 1 - \frac{\sum_{i=1}^n (\hat{y}^{(i)} - y^{(i)})^2}{\sum_{i=1}^n (\hat{y}^{(i)} - \bar{y})^2} \frac{(n-1)}{(n-k-1)}$$


# R-squared intuition



$$R^2 = 1 - \frac{SS_{reg}}{SS_{avg}} = 1 - \frac{\sum_{i=1}^n (\hat{y}^{(i)} - y^{(i)})^2}{\sum_{i=1}^n (\hat{y}^{(i)} - \bar{y})^2}$$

- Denominator is variance w.r.t. mean
- Numerator is variance w.r.t. regression line
- Lesser the variance wrt regression line the better
- How much variance is explained by linear regression?
  - More the merrier (Implies less error is left after regression)
- R-squared between 0 & 1. Higher the better

# Adjusted R-squared intuition



$$R_{adj}^2 = 1 - \frac{\sum_{i=1}^n (\hat{y}^{(i)} - y^{(i)})^2}{\sum_{i=1}^n (\hat{y}^{(i)} - \bar{y})^2} \frac{(n-1)}{(n-k-1)}$$

- If additional feature is added R squared increases
- But if the feature less useful in explaining variance, then adjusted R-squared decreases
- Penalized for using more features that do not add value

# Recap

- Population and Sample Regression
- Simple Linear Regression Intuition
- Linear Regression Algorithm
- Gradient Descent
- Impact of Scaling in Gradient Descent
- Closed form analytical solution
- Types of Gradient Descent
- Regression Evaluation Metrics





QUESTIONS